

Operating System Concepts

Lecture 3: System Calls, Linking & Loading

Omid Ardakanian
oardakan@ualberta.ca
University of Alberta

Recap

- Definition: a trap or exception is a **software-generated interrupt** caused by an error (e.g., division by zero) or a request for OS services (**system call**)

0: 0x00080000	Illegal address
1: 0x00100000	Memory violation
2: 0x00100480	Illegal instruction
3: 0x00123010	System call

Trap Vector

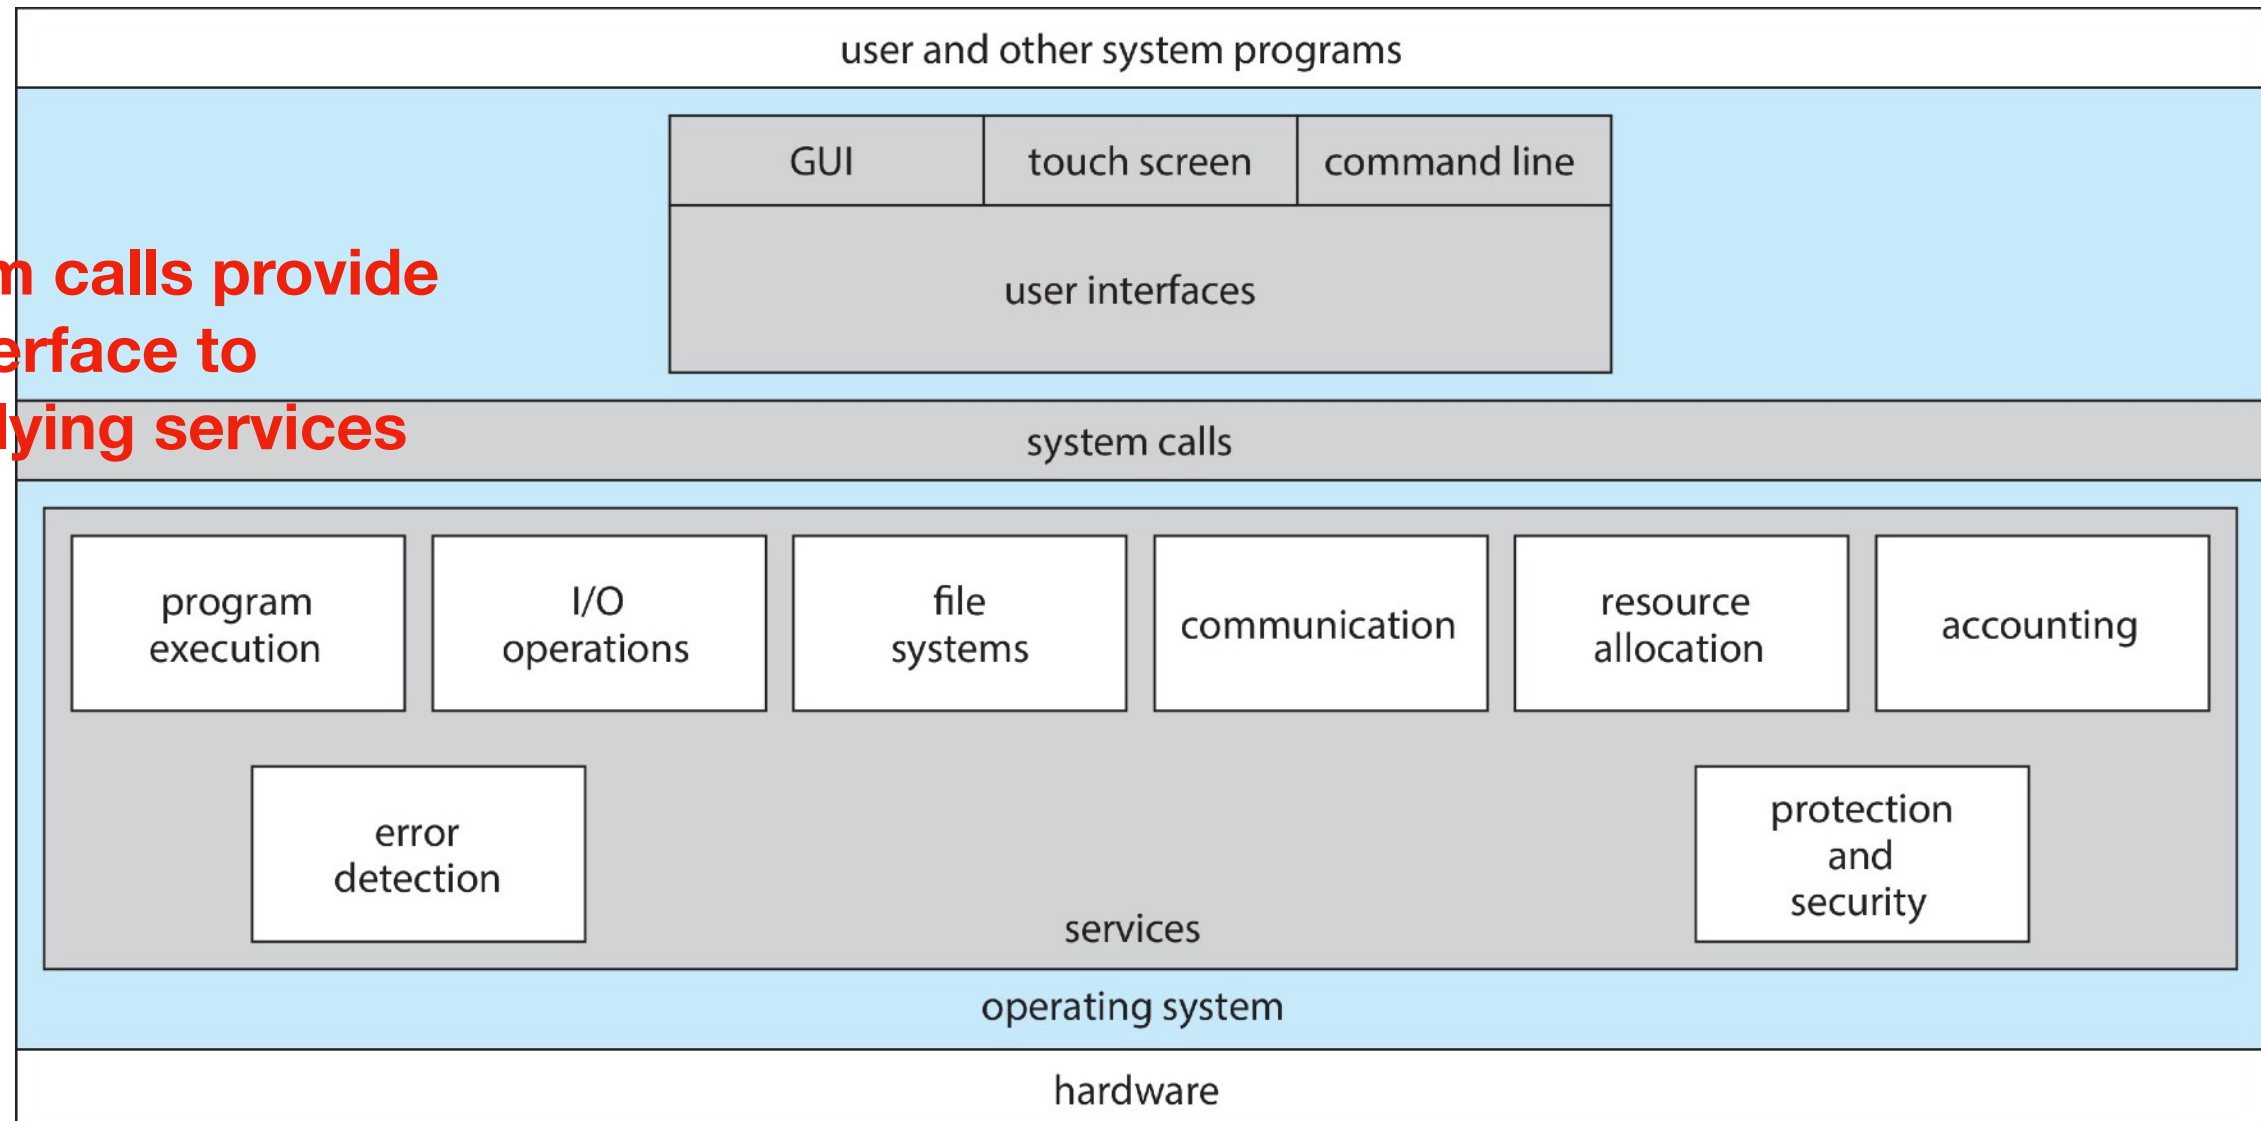
Today's class

- OS services
 - User interface
 - System calls: interface to OS services
 - Protection
- Basics of compiling, linking, and loading
- OS structure (time permitting)
 - Examples

OS Services

Operating System Services

system calls provide
an interface to
underlying services



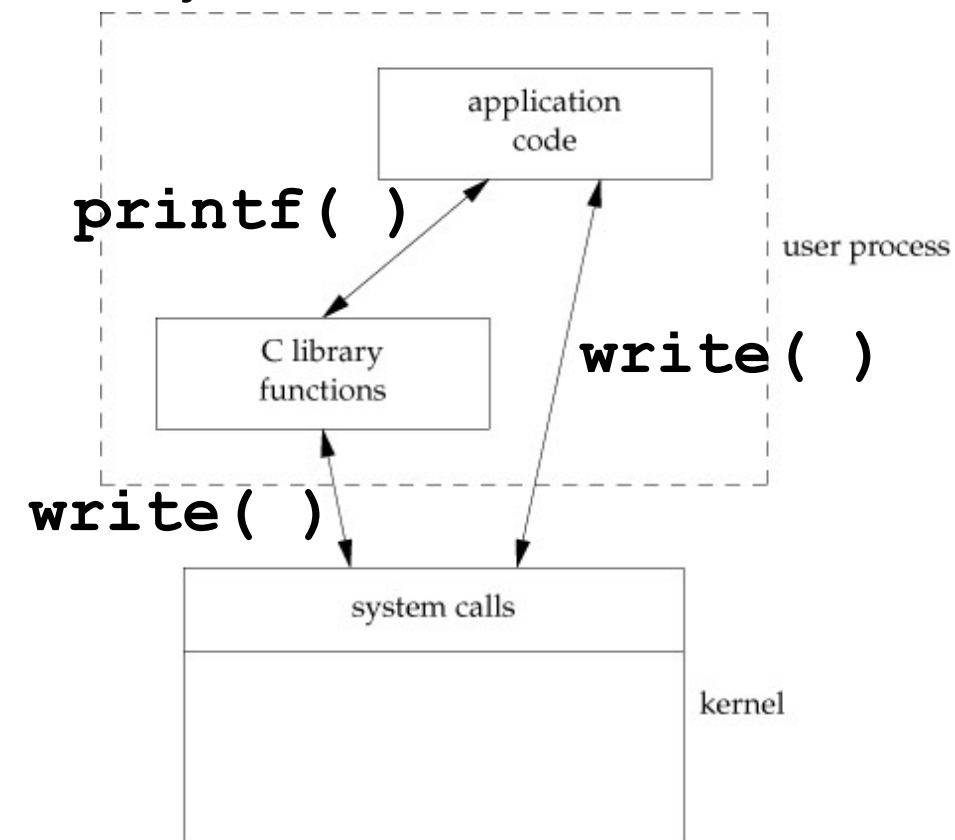
- **Services:** user interface; program execution (loading and error handling); I/O, file system, and interprocess communication services, resource allocation (scheduling), logging, protection and security, etc.
- **Goals:**
 - making the programming task easier and increasing user convenience
 - ensuring the efficient operation of the system (i.e., resource allocation, logging, protection and security)
 - ensuring the efficient operation of the system (i.e., resource allocation, logging, protection and security)
 - making the programming task easier and increasing user convenience

System calls... what are they?

- System calls are services provided by the kernel to application programs
 - a typical OS exports a few hundred (~300-400) system calls
 - see Section 2 of the Linux manual page for the list of Linux system calls:
<https://man7.org/linux/man-pages/man2/syscalls.2.html>
- System calls are necessary to access devices and files, request memory, set access permissions, stop and start processes, communicate with processes, set a timer, etc.

System calls... what are they?

- Each system call is typically defined as a function in the C library (e.g. glibc)
 - in UNIX-based systems, most system call wrapper functions are declared in `unistd.h`
- System calls can be invoked
 - via `syscall()` by providing their assembly language interface number, e.g. `syscall(SYS_write, ...)`; this is called a **raw system call**
 - indirectly via wrapper functions, e.g. `write(...)`, which may perform some pre/post-processing and error handling
- Many higher-level functions in the standard C library invoke system calls too
 - e.g. `printf(...)` from `stdio.h` calls `write(...)` from `unistd.h` after performing string formatting and type conversion



UNIX system calls

- Process control
 - `fork()`, `exec()`, `wait()`, `exit()`, ...
- Memory management
 - `brk()`, `sbrk()`, ...
- File/device management
 - `open()`, `close()`, `read()`, `write()`, `stat()`, `lseek()`, `link()`, `ioctl()`, ...
- Information maintenance
 - `getpid()`, `sleep()`, `time()`, ...
- Protection
 - `chmod()`, `chown()`, ...
- Communications
 - `pipe()`, `shm_open()`, `mmap()`, `socket()`, `accept()`, `send()`, `recv()`, ...

Linux system calls can be found at:

<http://man7.org/linux/man-pages/man2/syscalls.2.html>

Tracing system calls executed when you run an application

Use the `strace` command in Linux (see the man page of `strace`)

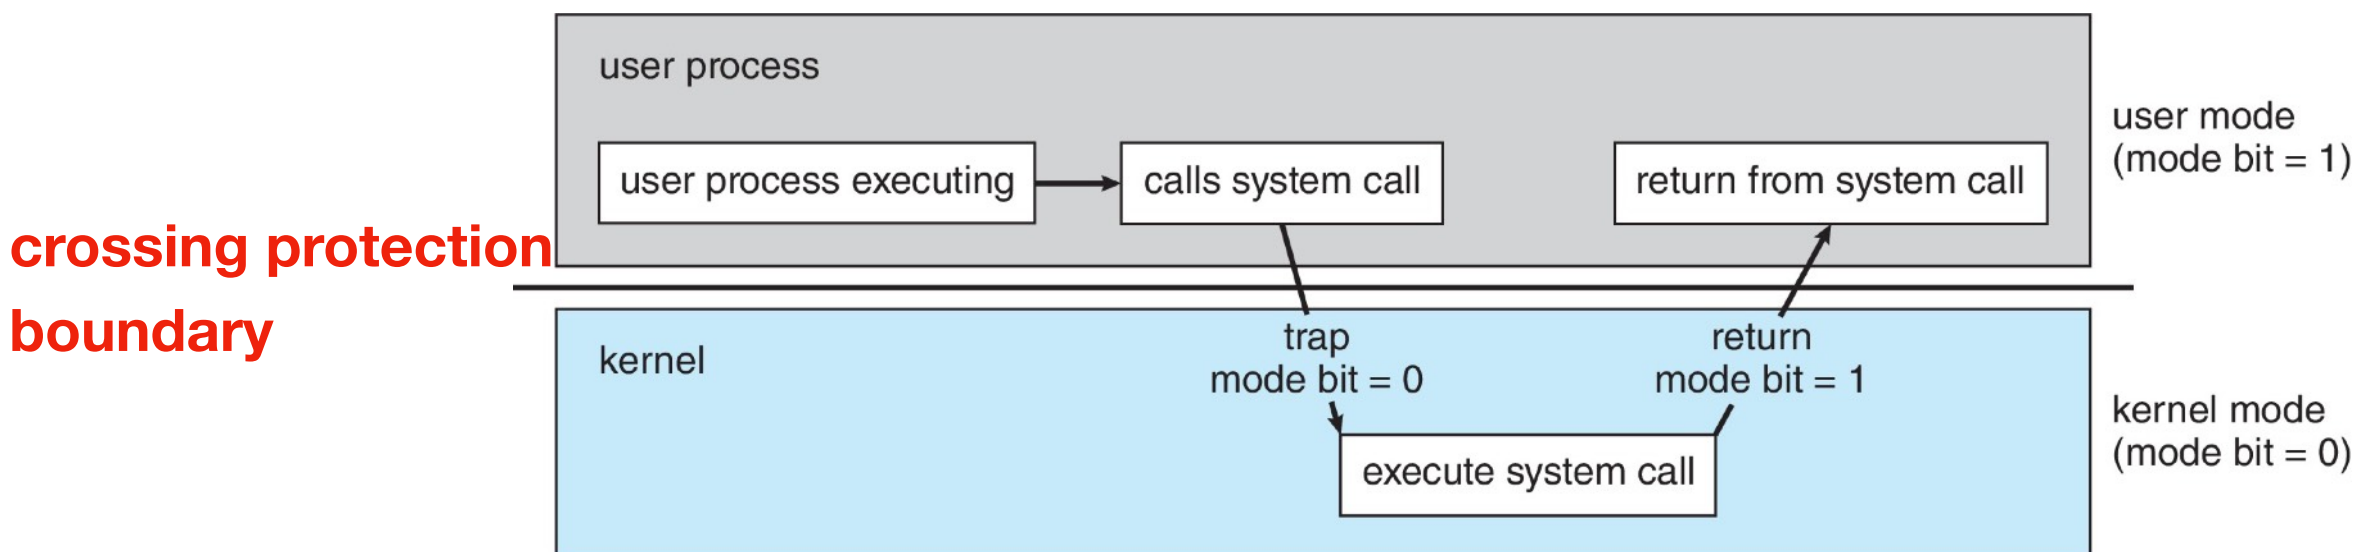
[illegible]

POSIX Application Programming Interface (API)

- Application programmers usually use functions from an API rather than invoking system calls directly
 - **Potability**: UNIX-based operating systems could have different declarations and implementations for their system calls, but they all support the same **POSIX API**; hence if you use the POSIX API you can expect that your program compiles and runs on any system that supports it
 - **Ease of use**: system calls are often more detailed and more difficult to use than their corresponding API functions
- **POSIX** (Portable Operating System Interface) is a family of standards implemented primarily for UNIX-based operating systems
 - POSIX compliant systems like Linux and macOS must implement the POSIX core standard (POSIX.1)
 - Linux man page says if a system call is specified in POSIX standards
 - POSIX thread library (aka pthreads) is defined in an extension of this API known as POSIX.1c
- **Win32 API** is the standard API for Windows operating systems

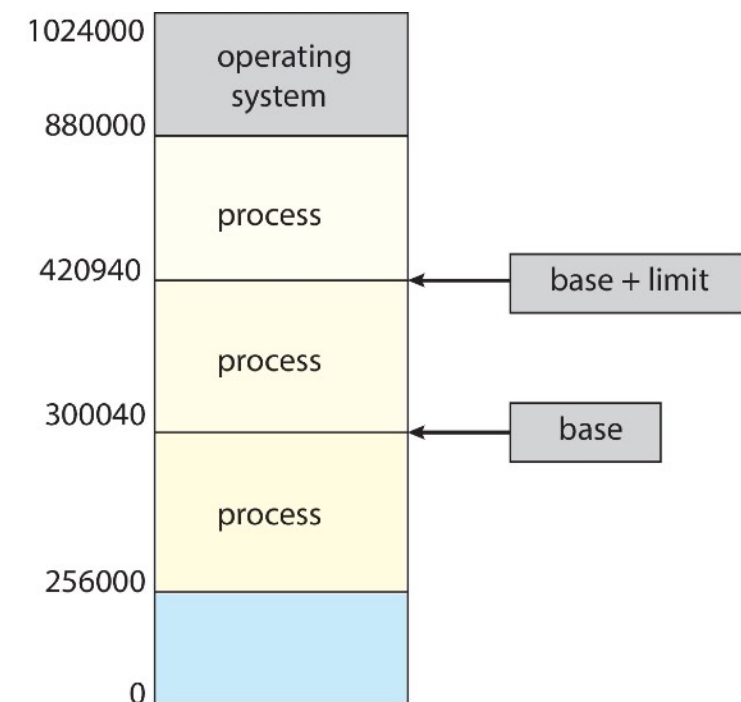
Protection bit — recap

- Hardware has a status **bit** that indicates the current mode (user or kernel)
 - there could be more than two operation modes
e.g., ARMv8 systems have seven modes
 - user programs run in the user mode
 - kernel code runs in the kernel mode with **full privileges** of hardware
- Operation modes provide the means for protecting OS from errant users
 - code could be buggy or malicious
- Examples of **privileged instructions** are I/O control, timer management, interrupt management; they can only run in the kernel mode
- Invoking a system call allows user program to run privileged instructions



Memory protection

- Hardware must provide support so that the OS can
 - protect user programs from each other
 - protect the OS from user programs
- The simplest technique is to use **base and limit registers**
 - base (**relocation**) and limit registers are loaded by the OS before context switching
 - base register holds the smallest legal physical memory address of the process
 - limit register holds the size of the memory allocated to a process
 - CPU checks each reference in user mode (instruction and data addresses) to ensure that it falls between base and base+limit



Timer interrupt - Another form of protection

What if the process running on the CPU does not wait for I/O or signals?

- Kernel protects CPU from being hogged using timer interrupts that occur at regular intervals (e.g., every 100 microseconds)
 - frequency is set by the kernel
 - it's yet another protection mechanism
- At each timer interrupt, the kernel chooses a new process to execute on CPU (scheduling)
- Interrupts can be temporarily deferred (it is crucial for implementing **mutual exclusion**) but user programs do not run with enough privilege to defer timer interrupts

0: 0x2ff080000	keyboard
1: 0x2ff100000	mouse
2: 0x2ff100480	timer
3: 0x2ff123010	Disk 1

Interrupt Vector

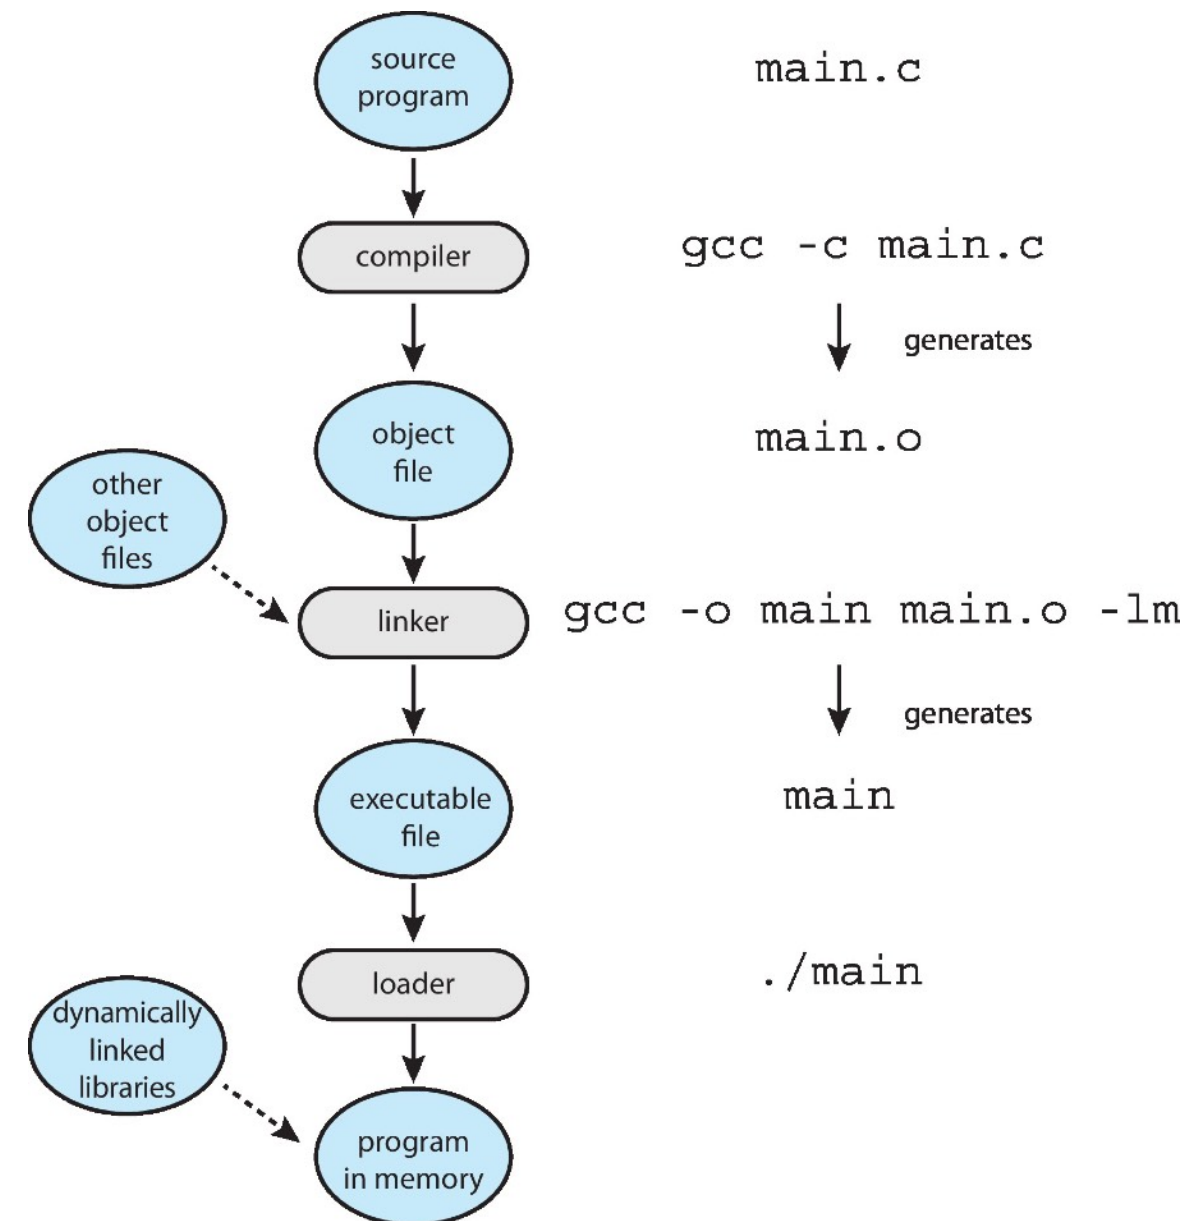
Summary

OS Services	Hardware Support
Interrupts	interrupt controller, interrupt request lines, interrupt vector
System calls	trap vector
I/O	interrupt and memory mapping
Protection	operation modes, privileged instructions, base and limit registers, timer interrupt
Scheduling & accounting	timer
Synchronization	atomic operations
Virtual memory	memory management unit (MMU), translation look-aside buffer (TLB)

Running a user program

From Program to Process

- Compiling is the process of converting a source file (ASCII code) into an object file
 - object files miss certain information, such as functions declared in other files and libraries
- Linking is the process of combining **relocatable object files** and specific libraries into a single binary executable file
- Loading is the process of bringing this executable file into memory
 - the loader is executed when you enter the name of the executable file on the command line; this is done using the `execve()` system call
 - the loader sets up the process memory to contain code and data from executable
 - a library can be conditionally linked and loaded if it is required during the run time; this can be done using dynamically linked libraries (DLLs)



Executable and Linkable Format (ELF)

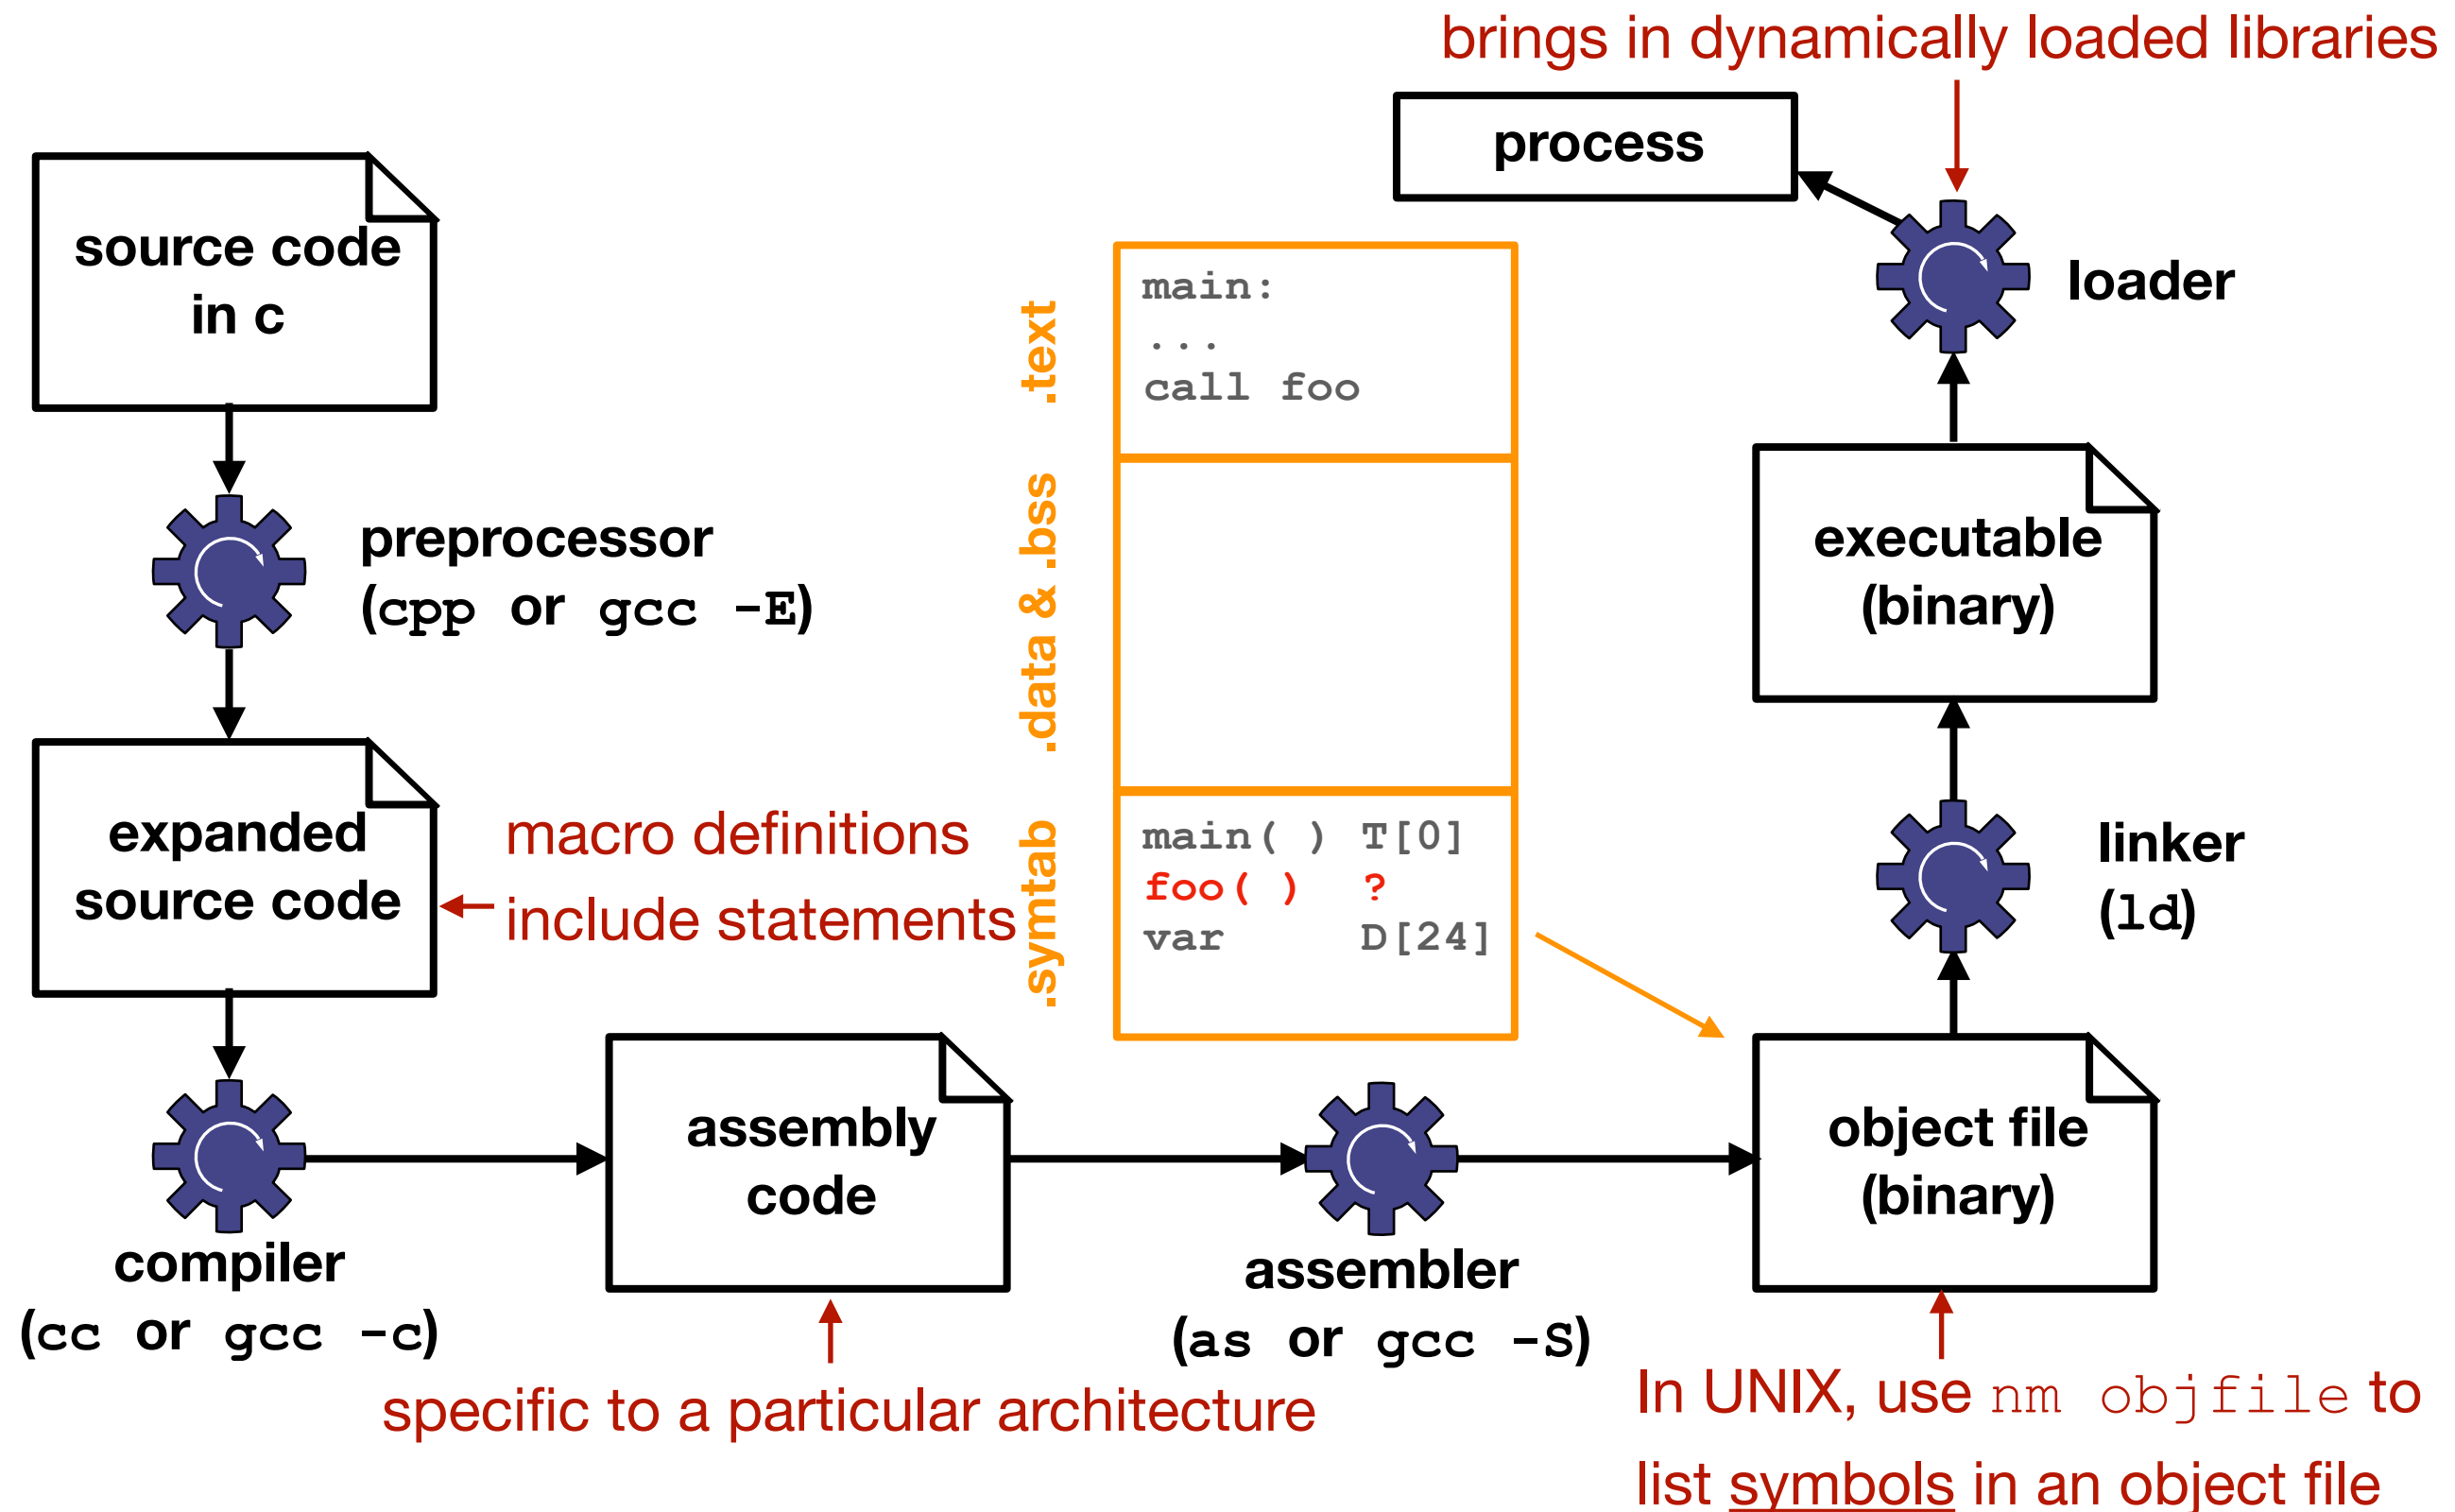
- Executable files must have a standard format, so that OS knows how to load and start them
- ELF is a common standard across modern UNIX operating systems, such as Linux
 - ELF header starts with a 4-byte magic string: 0x7F followed by 0x45 0x4C 0x46 (ELF in ASCII)
- **ELF relocatable file** contains the compiled machine code and a symbol table containing metadata about functions and variables referenced in the program
 - the symbol table is necessary for linking the relocatable object file with other object files
- **ELF executable file** contains the address of the first instruction of the program (program's **entry point**: `_start` function)
 - this file can be loaded into memory

Executable and Linkable Format (ELF)

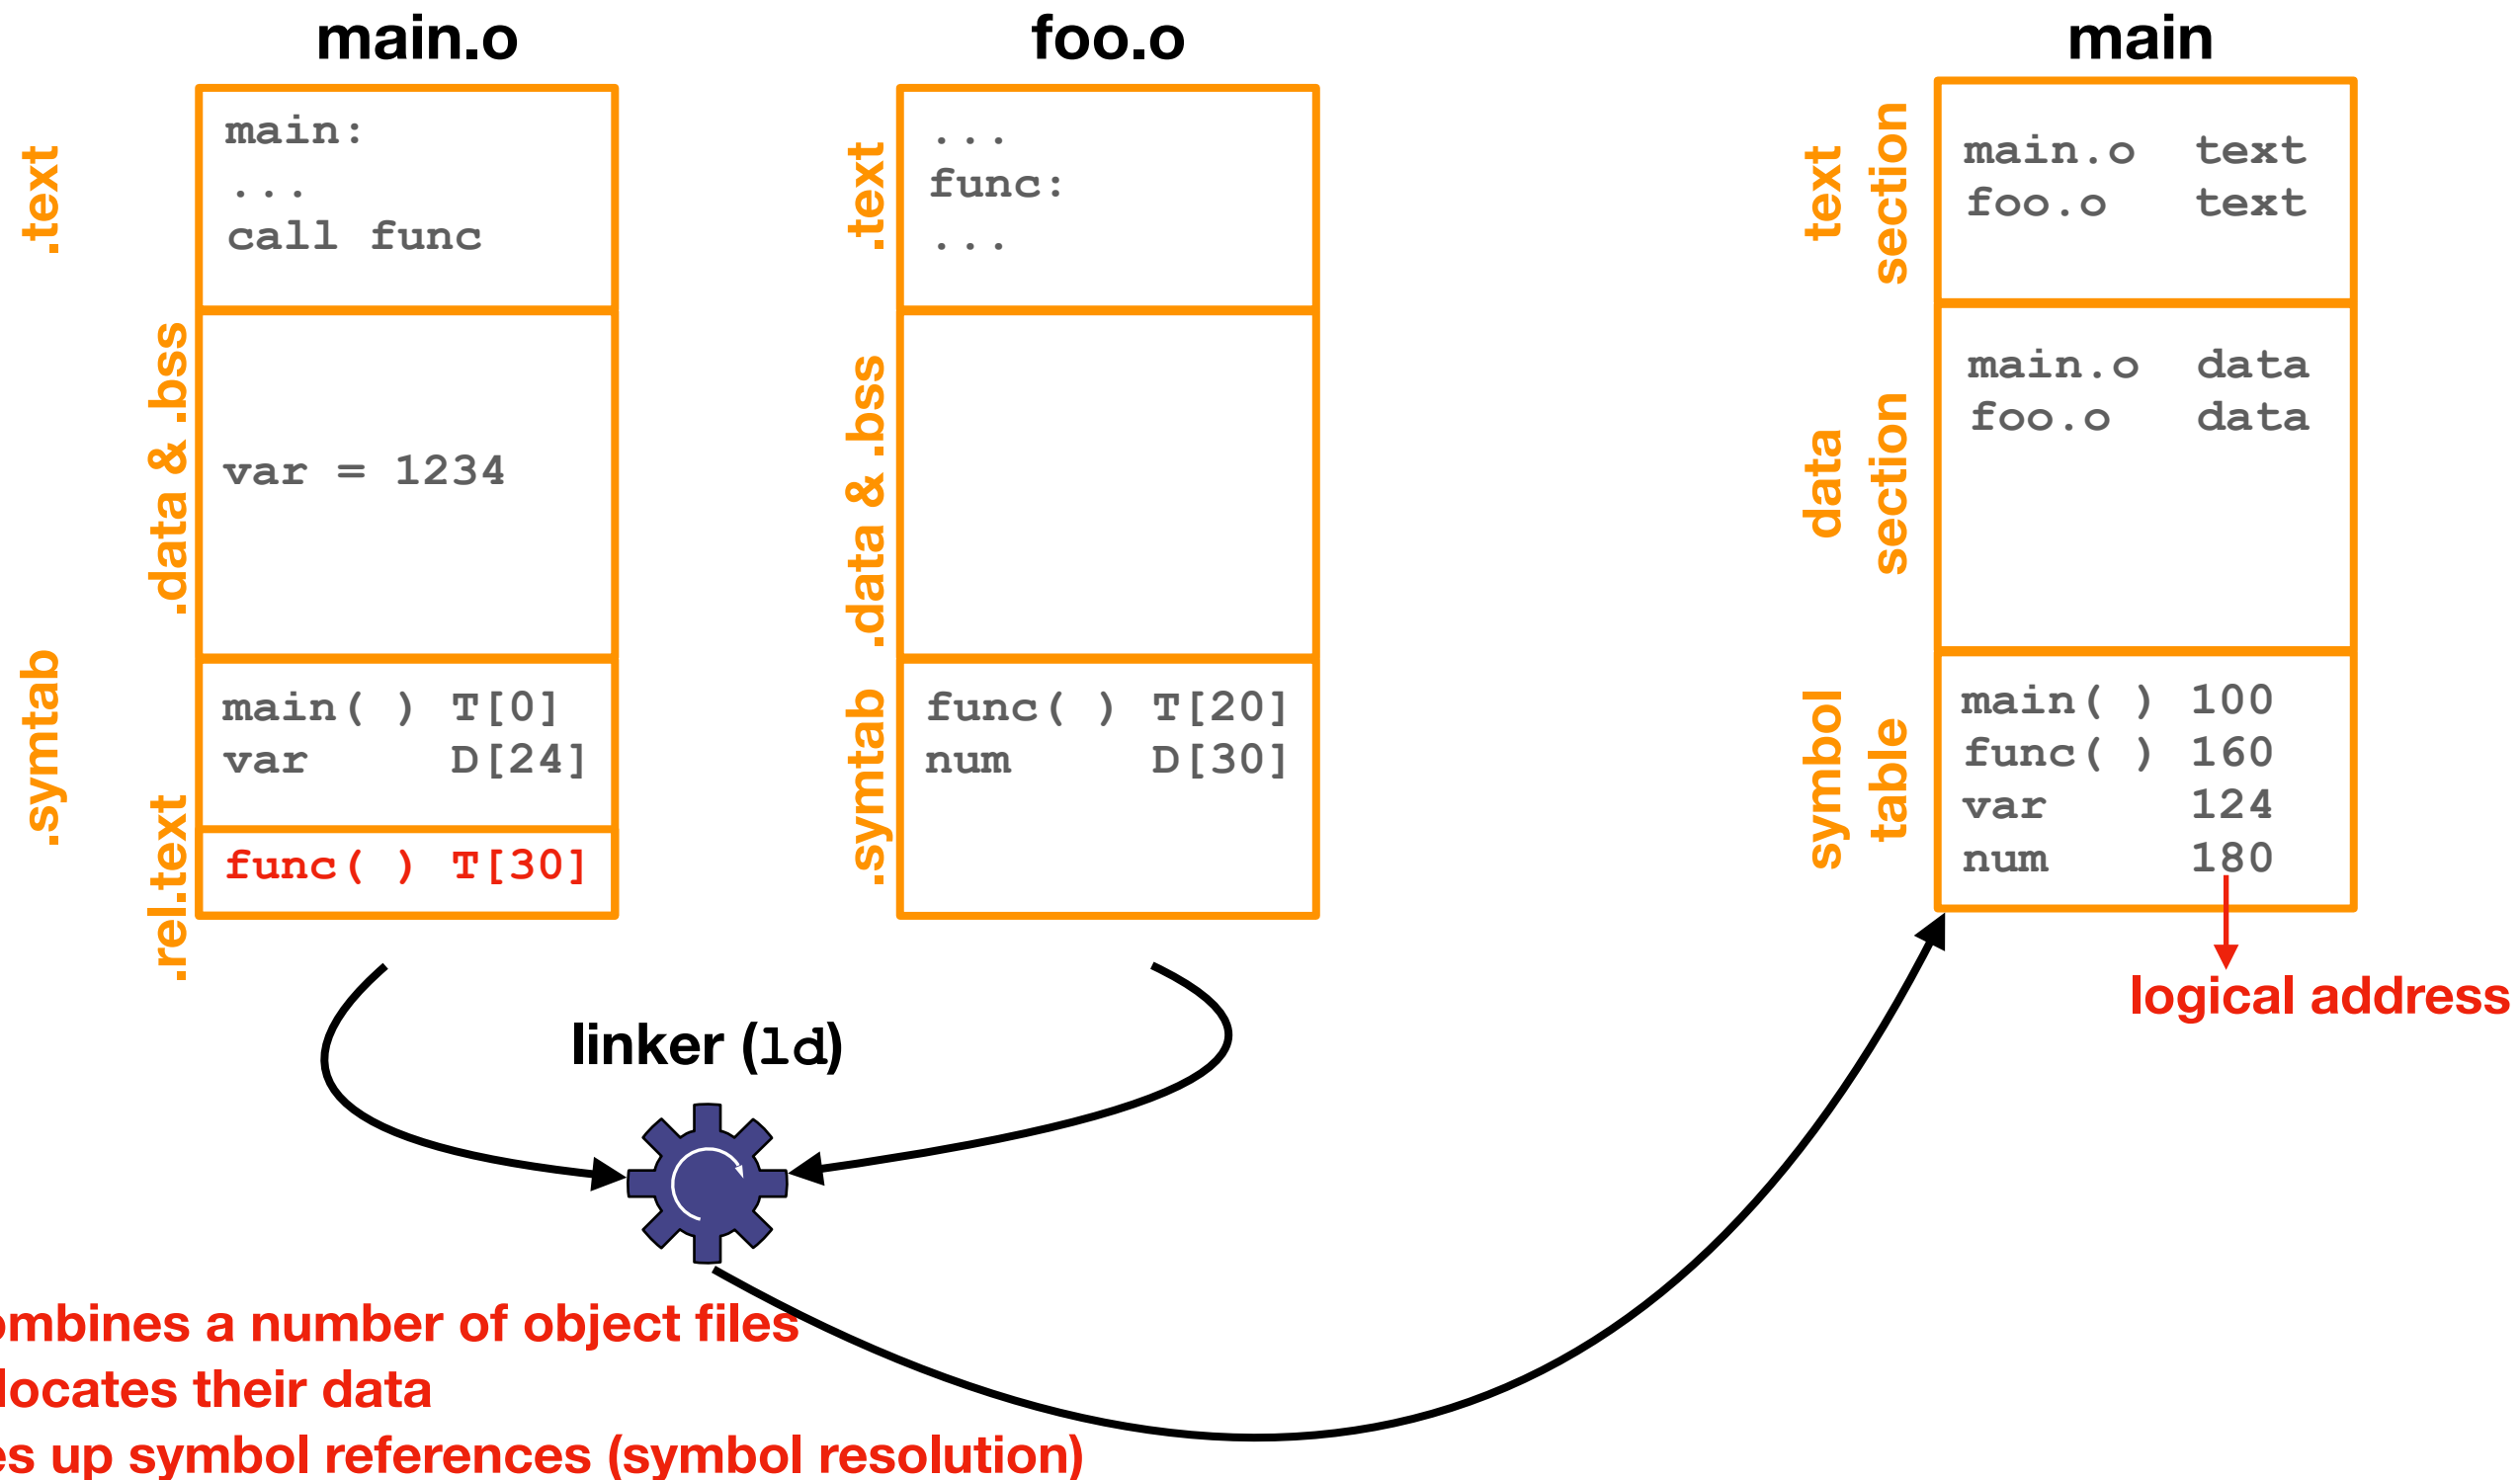
ELF file is divided into multiple sections

- **.text** contains the machine code
 - In UNIX-based systems, see the content of this section using `objdump -drS objfile`
- **.data** contains **initialized** global variables
- **.bss** (block storage start) contains **uninitialized** global variables
 - occupies no space in the object file actually
 - there is no point storing more zeroes on your disk
- **.rodata** contains read-only data such as constant strings
 - e.g., the format strings in `printf` statements
- **.symtab** contains the symbol table
 - information about functions and global variables defined and referenced in the program
- **.rel.text** and **.rel.data** contain relocation information for functions and global variables that are referenced but not defined (external references)
 - linker modifies this section when combining the object files, resolving external references

Compiler, Linker and Loader in action

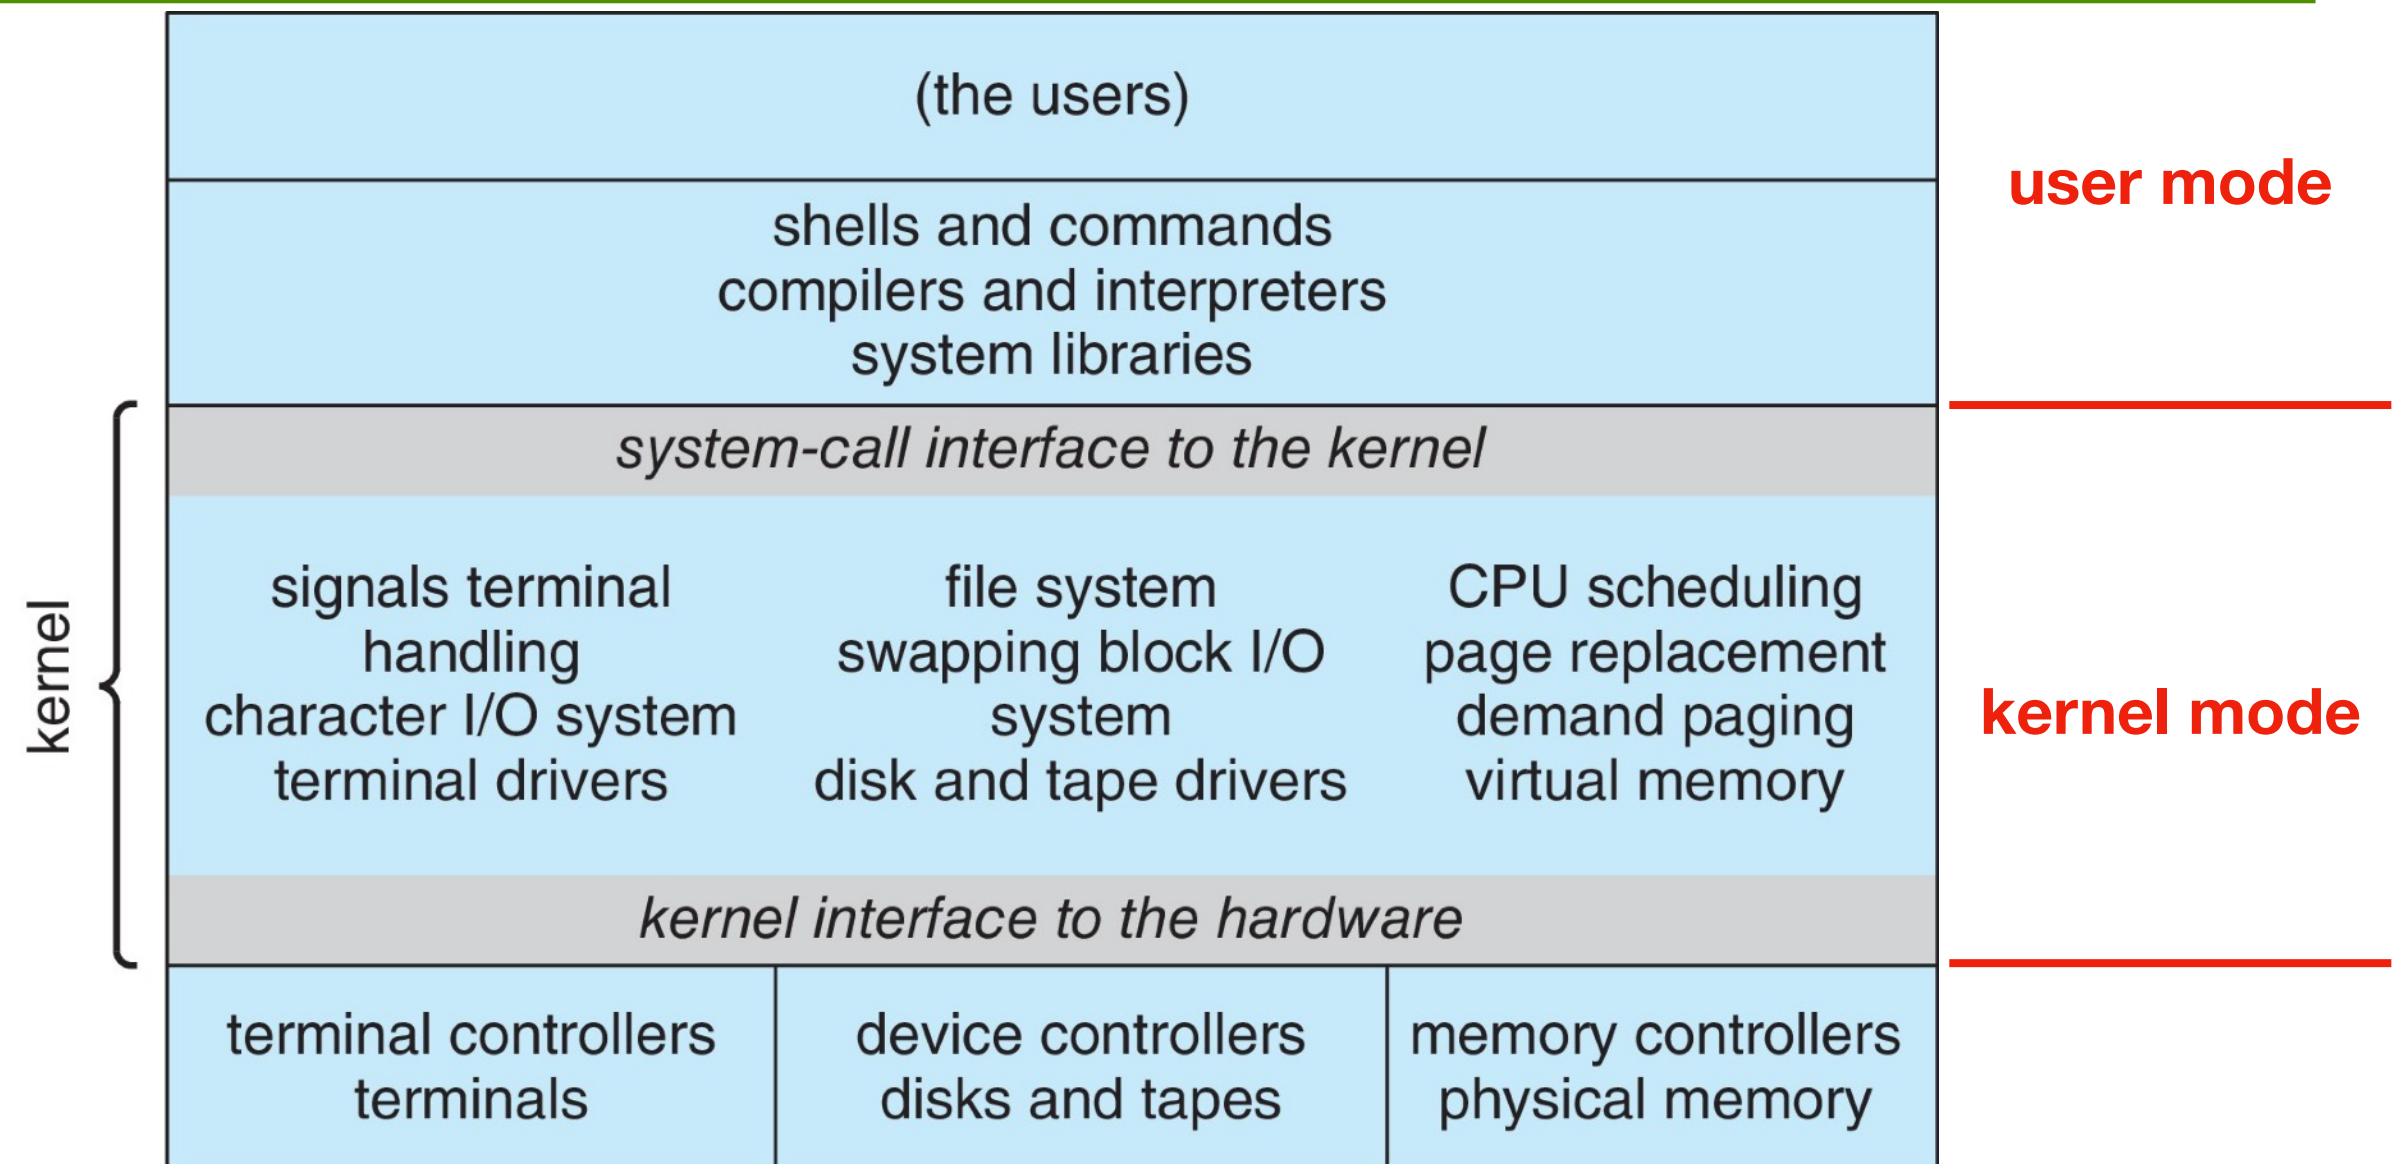


Linking



OS Structure

UNIX system has a *monolithic* structure

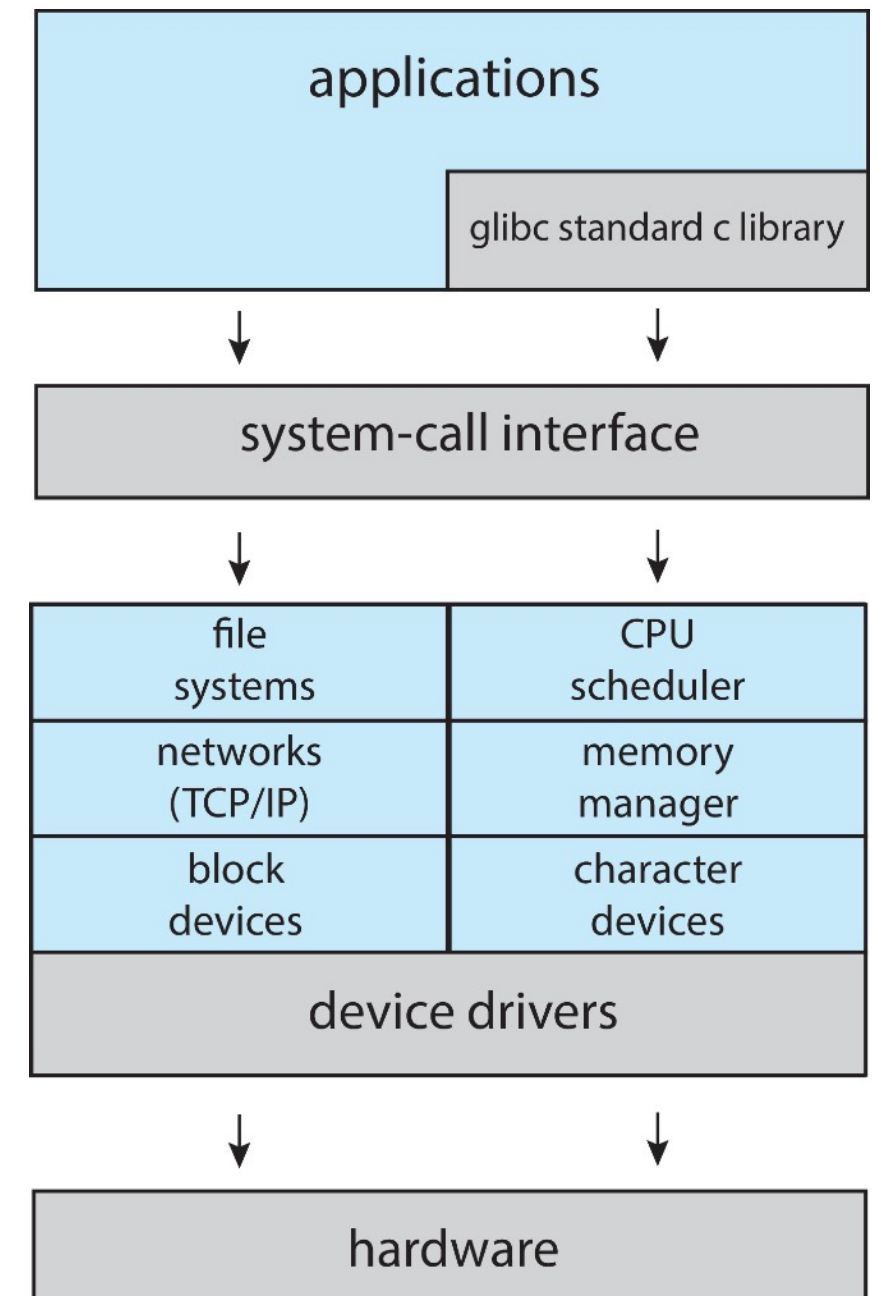


All functionality of the kernel placed into a single static binary file that runs in a single address space

- faster communication with the kernel
- little overhead in the system call interface

Linux has a monolithic structure too


- core (~ 30% of Linux source code) + dynamically loaded kernel modules
 - examples are device drivers, file systems, network protocols
- modules can be loaded at boot time or during run time
 - adding new modules does not require recompiling the kernel
 - each module talks to other modules through known interfaces



Linux has a monolithic structure too

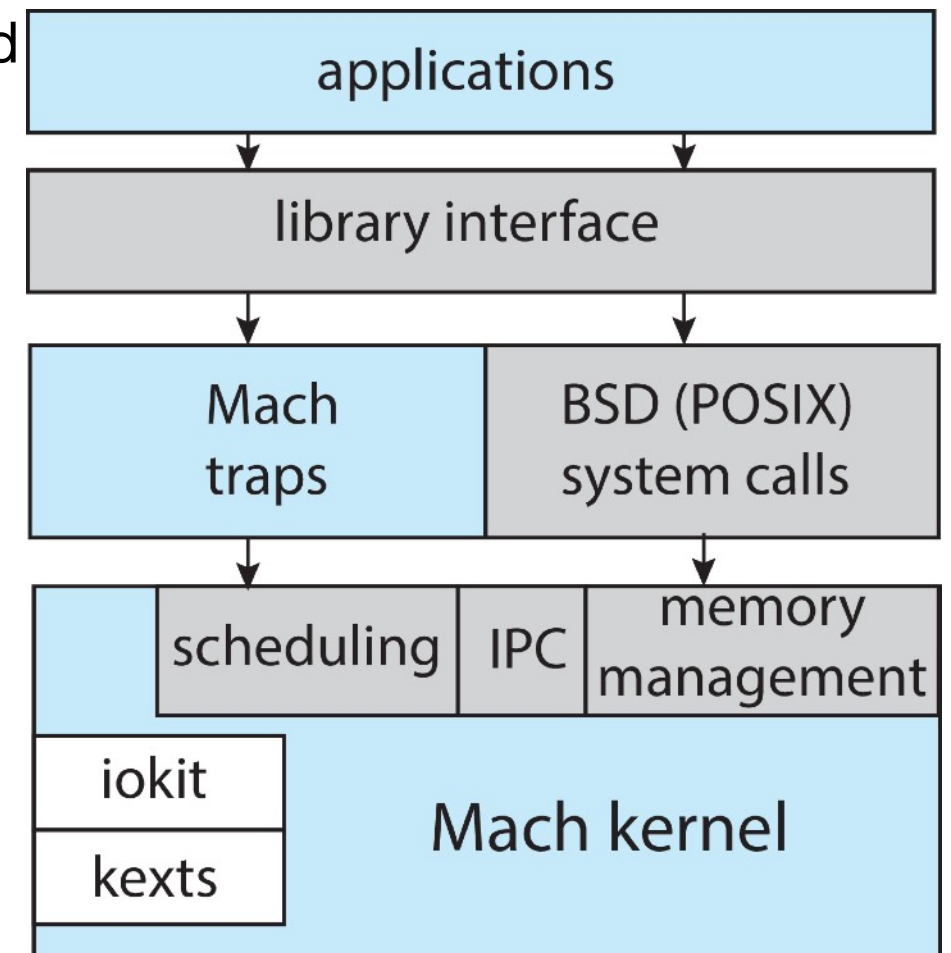
- Applications use the GNU version of the standard C library (**glibc**) which provides the functionality required by POSIX
 - glibc provides a wrapper around system calls
 - glibc is the system-call interface to the kernel
- Linux kernel source: <https://www.kernel.org/>
 - Directory structure:

```
include: public headers
kernel: core kernel components (e.g., scheduler)
arch: hardware-dependent code
fs: file systems
mm: memory management
ipc: interprocess communication
drivers: device drivers
usr: user-space code
lib: common libraries
```



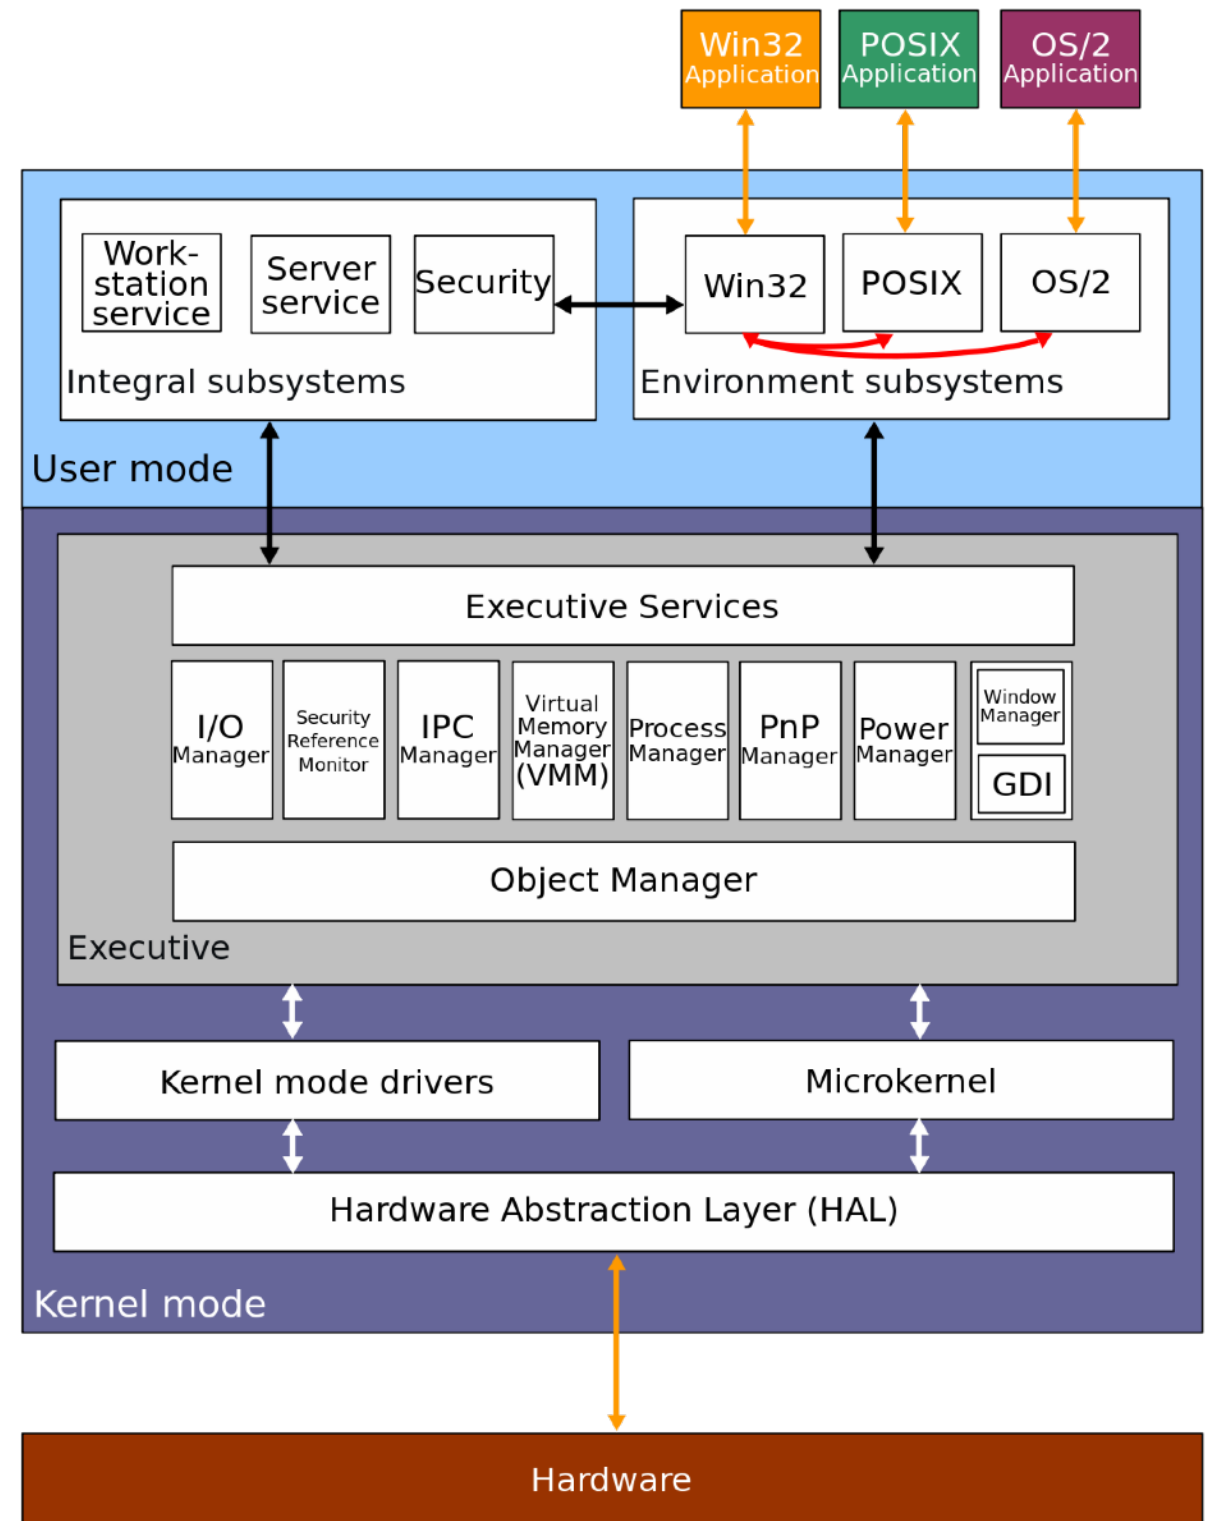
Darwin, macOS kernel, has a hybrid structure

- Darwin is Mach **microkernel** + BSD (threads, command line interface, networking, file system) + user-level services (GUI)
- Layered structure
 - advantages: modularity, simplicity, portability, ease of design/debugging
 - disadvantage: communication overhead between layers, extra copying, book-keeping
- Microkernel structure
 - a small kernel providing
 - interprocess communication (message passing usually through **ports**)
 - basic functionality (e.g., scheduling and virtual memory management)
 - other OS functionality (device drivers, file system, networking, user interface, etc.) implemented as user-level programs



Windows NT

- Layered architecture with several modules
- Windows executive provides core OS services (Ntoskrnl.exe)
- Kernel itself resides between HAL and the executive layer
 - responsible for thread scheduling and interrupt handling
- Windows native APIs are undocumented
 - POSIX, OS/2 and Win32 APIs are documented and can be used by applications
- The hardware abstraction layer (HAL.dll) provides portability across a variety of hardware platforms
 - device drivers call functions in HAL to interface with the hardware



Homework

- Familiarize yourself with Linux system calls
- Read the Linux Journal article about linking and loading