

Operating System Concepts

Lecture 4: Process Abstraction

Omid Ardakanian
oardakan@ualberta.ca
University of Alberta

Today's class

- Process abstraction
 - How does OS create this abstraction?
 - Why is it useful?
 - What happens during a context switch? What are the roles of the **dispatcher** and **scheduler**?

Process abstraction


- A process is a program during execution. It is an execution environment with restricted rights
 - has its own resources (CPU registers, memory to contain program code and data, file descriptors, etc.)
 - encapsulates one or more **threads** that share process resources
 - is characterized by a **unique identifier** (PID)
- Different processes may run different instances of the same program (process \neq program)
 - e.g., you can run several instances of a web browser
- Why do we need this abstraction?
 - necessary for concurrent execution and protection

Memory layout for process

- Multiple sections of a process
 - text section containing the program code
 - data section containing global variables (initialized and uninitialized)
 - stack containing temporary data, function parameters, return addresses, and local variables
 - heap containing memory that is dynamically allocated at runtime using `malloc` library function or the `brk/sbrk` system call

Which variables are allocated on the stack?

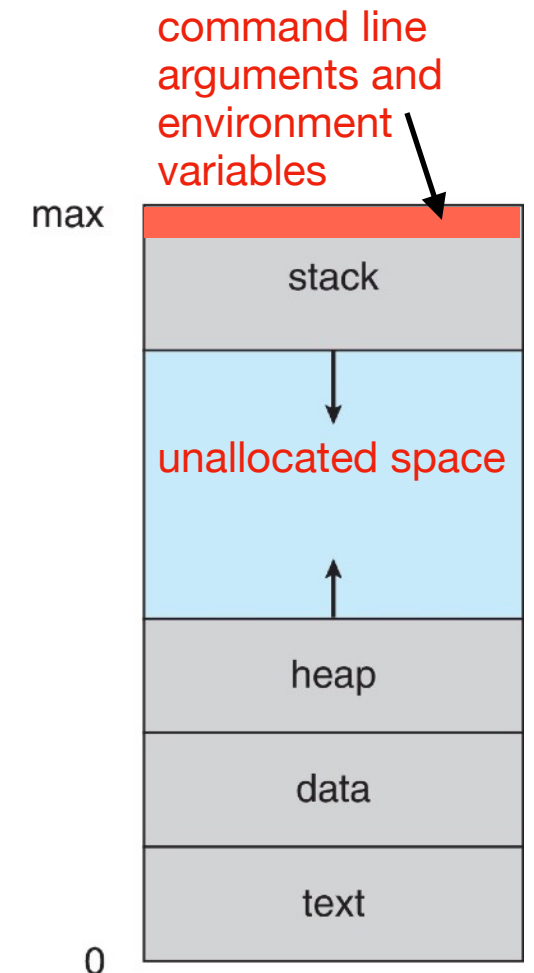
```
#include <stdio.h>
void foo (int n) {
    int i, a[5], *b;
    if (n == 0) return;
    b = new int[n];
    printf ("foo(%d): %p,%p,%p,%p \n", n, &i, a, &b, b);
    foo(n-1);
}
int main ( ) { foo(10); }
```



Programmer's view of memory:
single space containing
this one process only

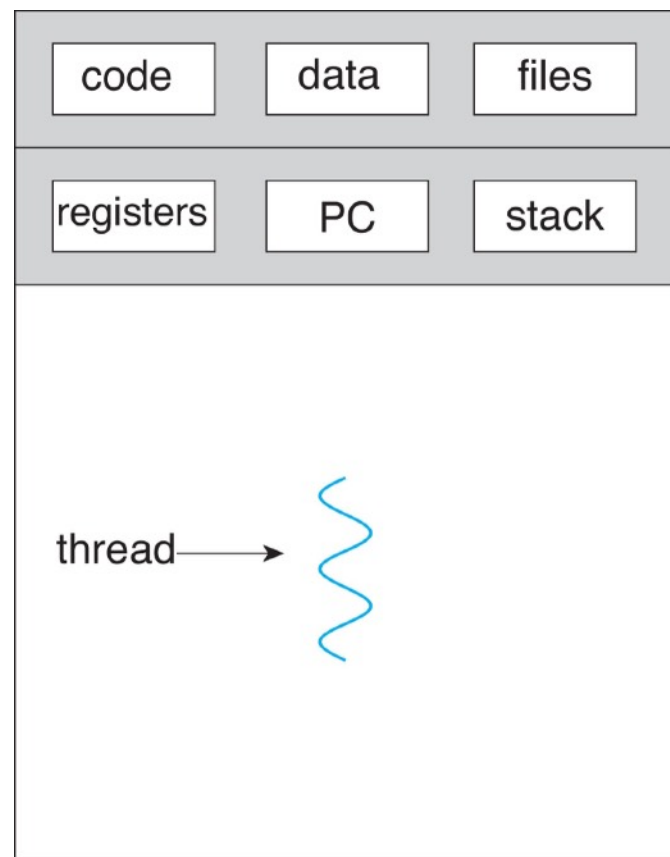
Memory layout of a process

- Multiple sections of a process
 - text section (fixed size) containing the program code
 - data section (fixed size) containing global variables (initialized/uninitialized)
 - stack containing temporary data, function parameters, return addresses, and local variables
 - heap containing memory that is dynamically allocated at runtime using `malloc` library function or the `brk/sbrk` system call
- Each process has a distinct and isolated **address space** (i.e., set of addresses that can be accessed by its code)
 - addresses in the executable file are as if it is loaded at memory address 00000000
 - these addresses need to be adjusted when the program is **relocated** to somewhere else
 - no process can read or write memory of another process
- Hardware translates virtual addresses to physical addresses

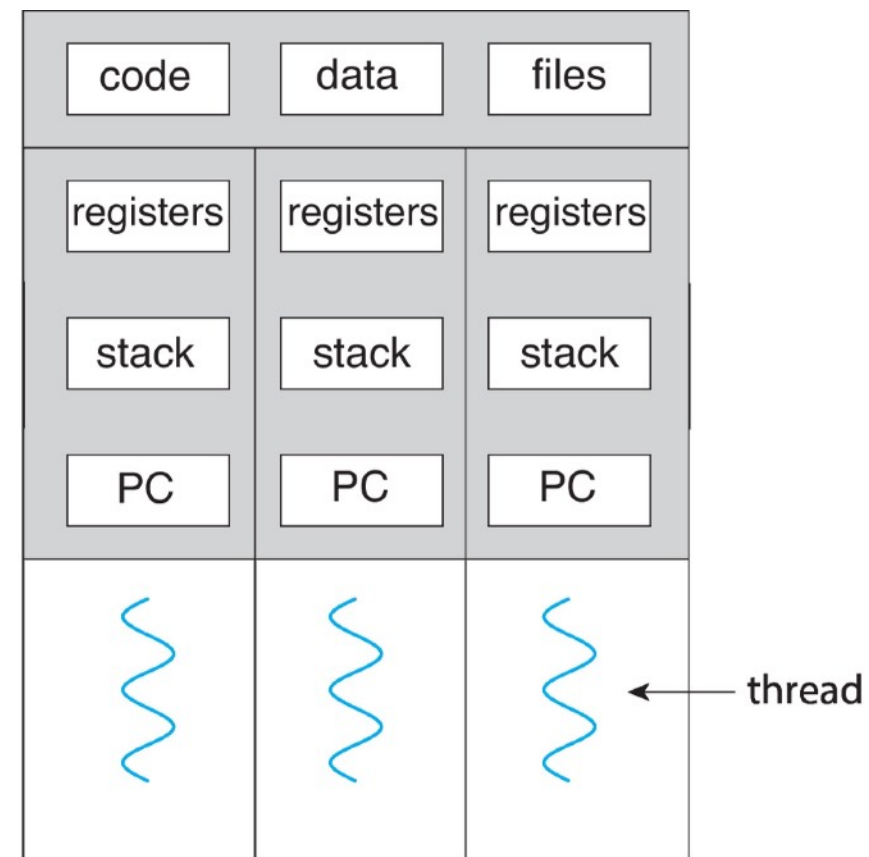


Single vs. multi-threaded process

- A thread is a sequential execution stream of instructions
- The address space of the process is shared among its threads
 - sharing heap, text, static data sections in addition to file descriptors
- Threads can execute simultaneously on different cores of a multicore system
 - e.g. in a word processor you can simultaneously type a character and run the spell checker!



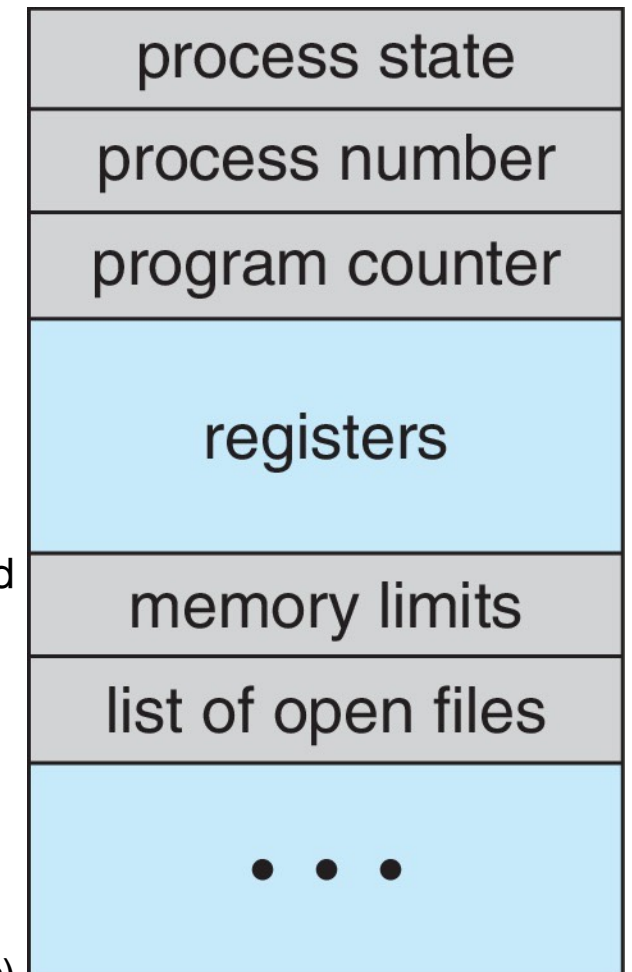
single-threaded process



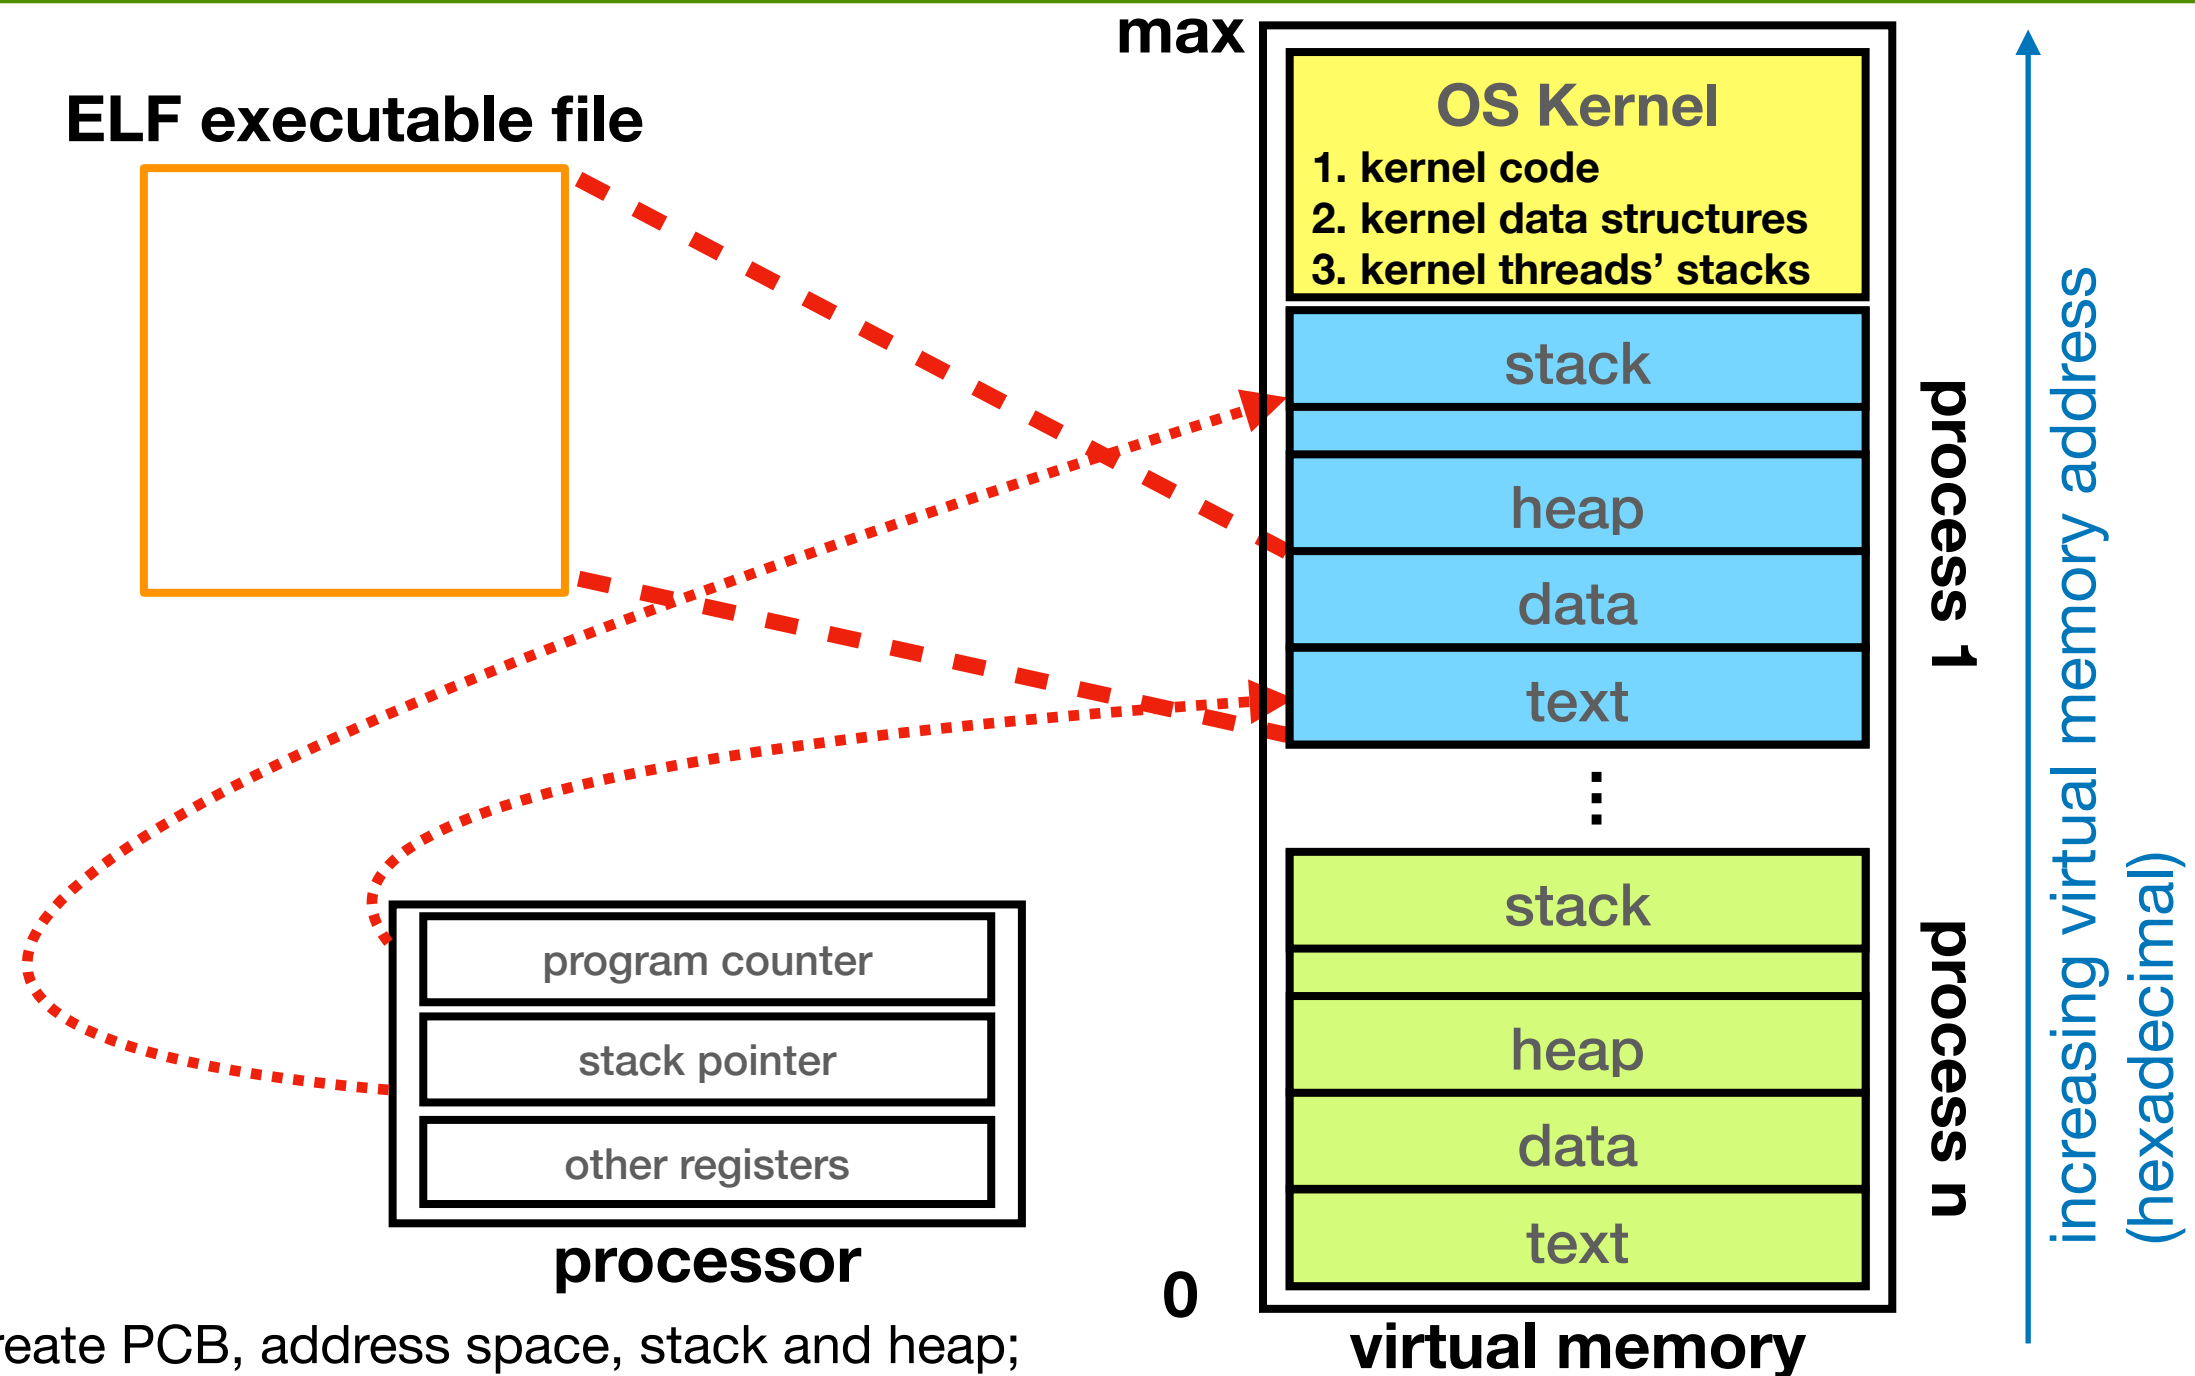
multithreaded process

Process control block


- For every process, OS keeps the state of process execution (metadata) in a **process control block** (PCB)
- PCB is a kernel data structure in memory; it represents run-time information about the process, defining its **context**
 - Process status (running, ready, blocked/waiting)
 - Process ID (PID) and its children's PIDs
 - CPU registers, including program counter (PC), stack pointer (SP), base/relocation and limit registers, page-table base register (PTBR), and general-purpose registers
 - Thread control block(s)
 - Accounting information (e.g., execution time, time elapsed since start)
 - Scheduling information (e.g., priorities, queue pointers for state queues)
 - Set of OS resources in use (e.g., list of open files, I/O devices allocated to the process)
 - Current working directory
 - Username of owner
 - ...
- PCB in Linux is represented by the C structure called `task_struct`
It is defined in `<linux/sched.h>`
 - `task_struct` contains `mm_struct` which represents the address space of a given process



Loading — revisited



1. create PCB, address space, stack and heap; pushes argc and argv on the stack
2. initialize registers and program counter
3. load code and data sections of executable file into memory



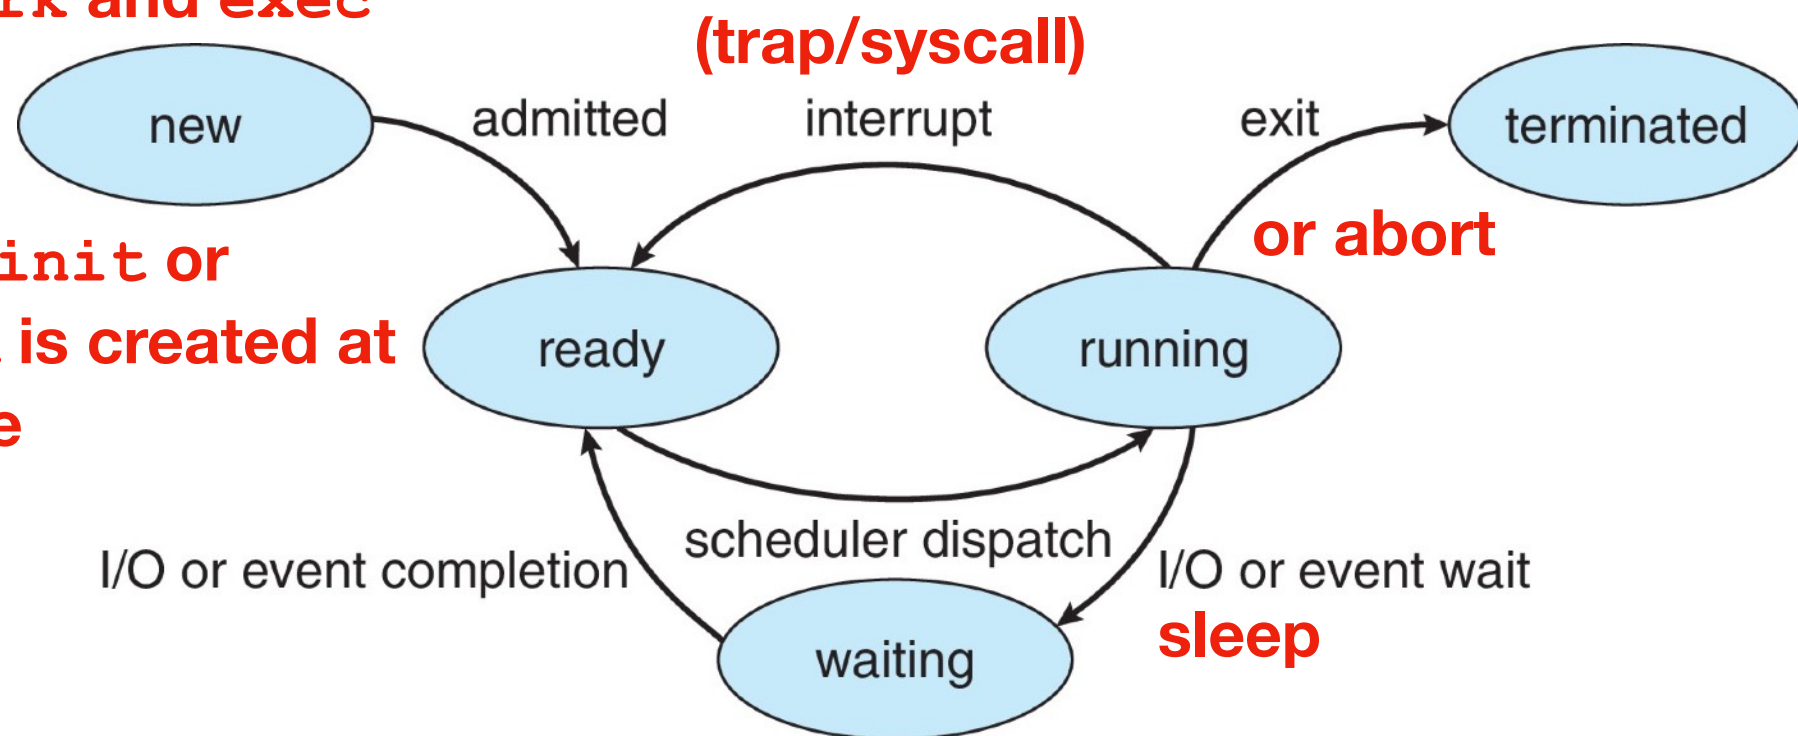
Keeping track of processes

- OS juggles many processes at a time
 - only one process can be running per core (the kernel maintains a pointer to this process)
 - but many processes can be in ready and waiting states
- OS puts PCBs of the active processes in appropriate queues
 - ready queue (organized by the process-scheduling priority, the arrival time, etc.)
 - wait queue for each device
 - **zombie** queue (terminated processes that have an entry in the process table are kept in that queue until they are reaped by their parent)
- state change happens as a result of the process actions (e.g., termination or invoking system calls), OS actions (scheduling), and external actions (hardware interrupts)

Process/thread lifecycle

**new process created
using fork and exec**

**process init or
systemd is created at
boot time**

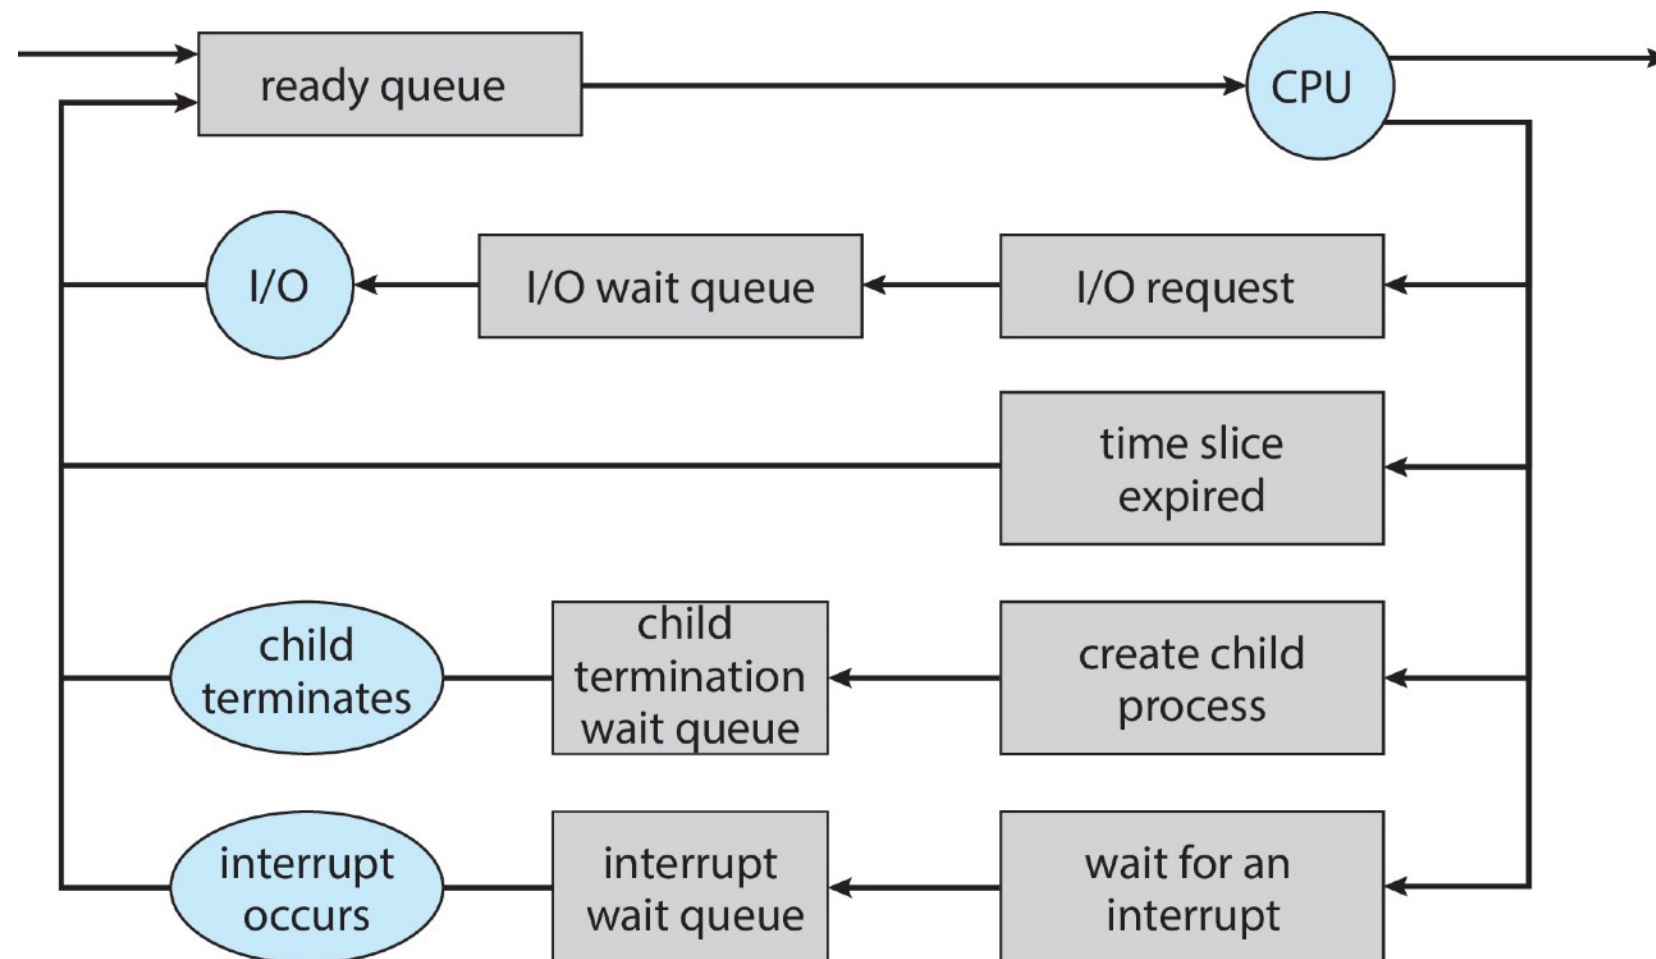


Zombie and orphan processes

- a process that has terminated, but its parent has not yet read its exit status becomes **zombie**
- a process becomes **orphan** when its parent terminates while it is still running

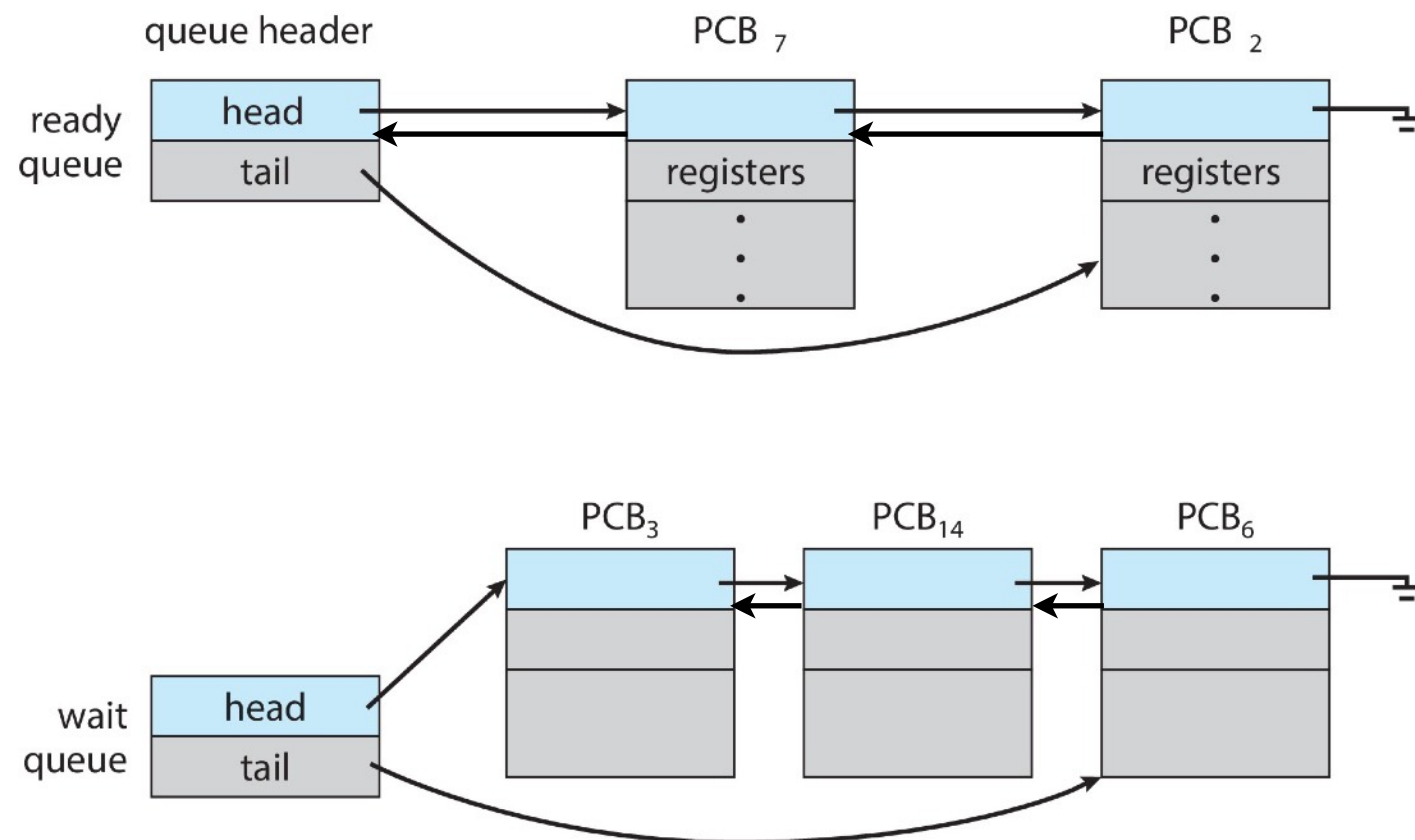
Multiprogramming

- Only one process is active at a time (per CPU core)
- OS gives out CPU to different processes in the ready queue
 - the number of processes currently in memory is the **degree of multiprogramming**



Scheduling

- The scheduler maintains a data structure (e.g., a **doubly linked list**) of PCBs
 - PCBs are moved from one queue to another queue
 - job queue: every process in the system
 - ready queue: processes residing in the main memory, waiting to run on CPU
 - device queue: processes waiting for a particular I/O device, each device has its own queue



Scheduling

- The scheduler selects a process among processes in the ready queue to run
 - selected process runs on CPU
 - if no process left in the ready queue, the CPU runs an idle process
 - scheduling can be performed for fairness, minimum latency, providing real-time guarantees, etc.
- The (short-term) scheduler makes a decision very fast (<10 ms) and is called very often (e.g., every 100ms)

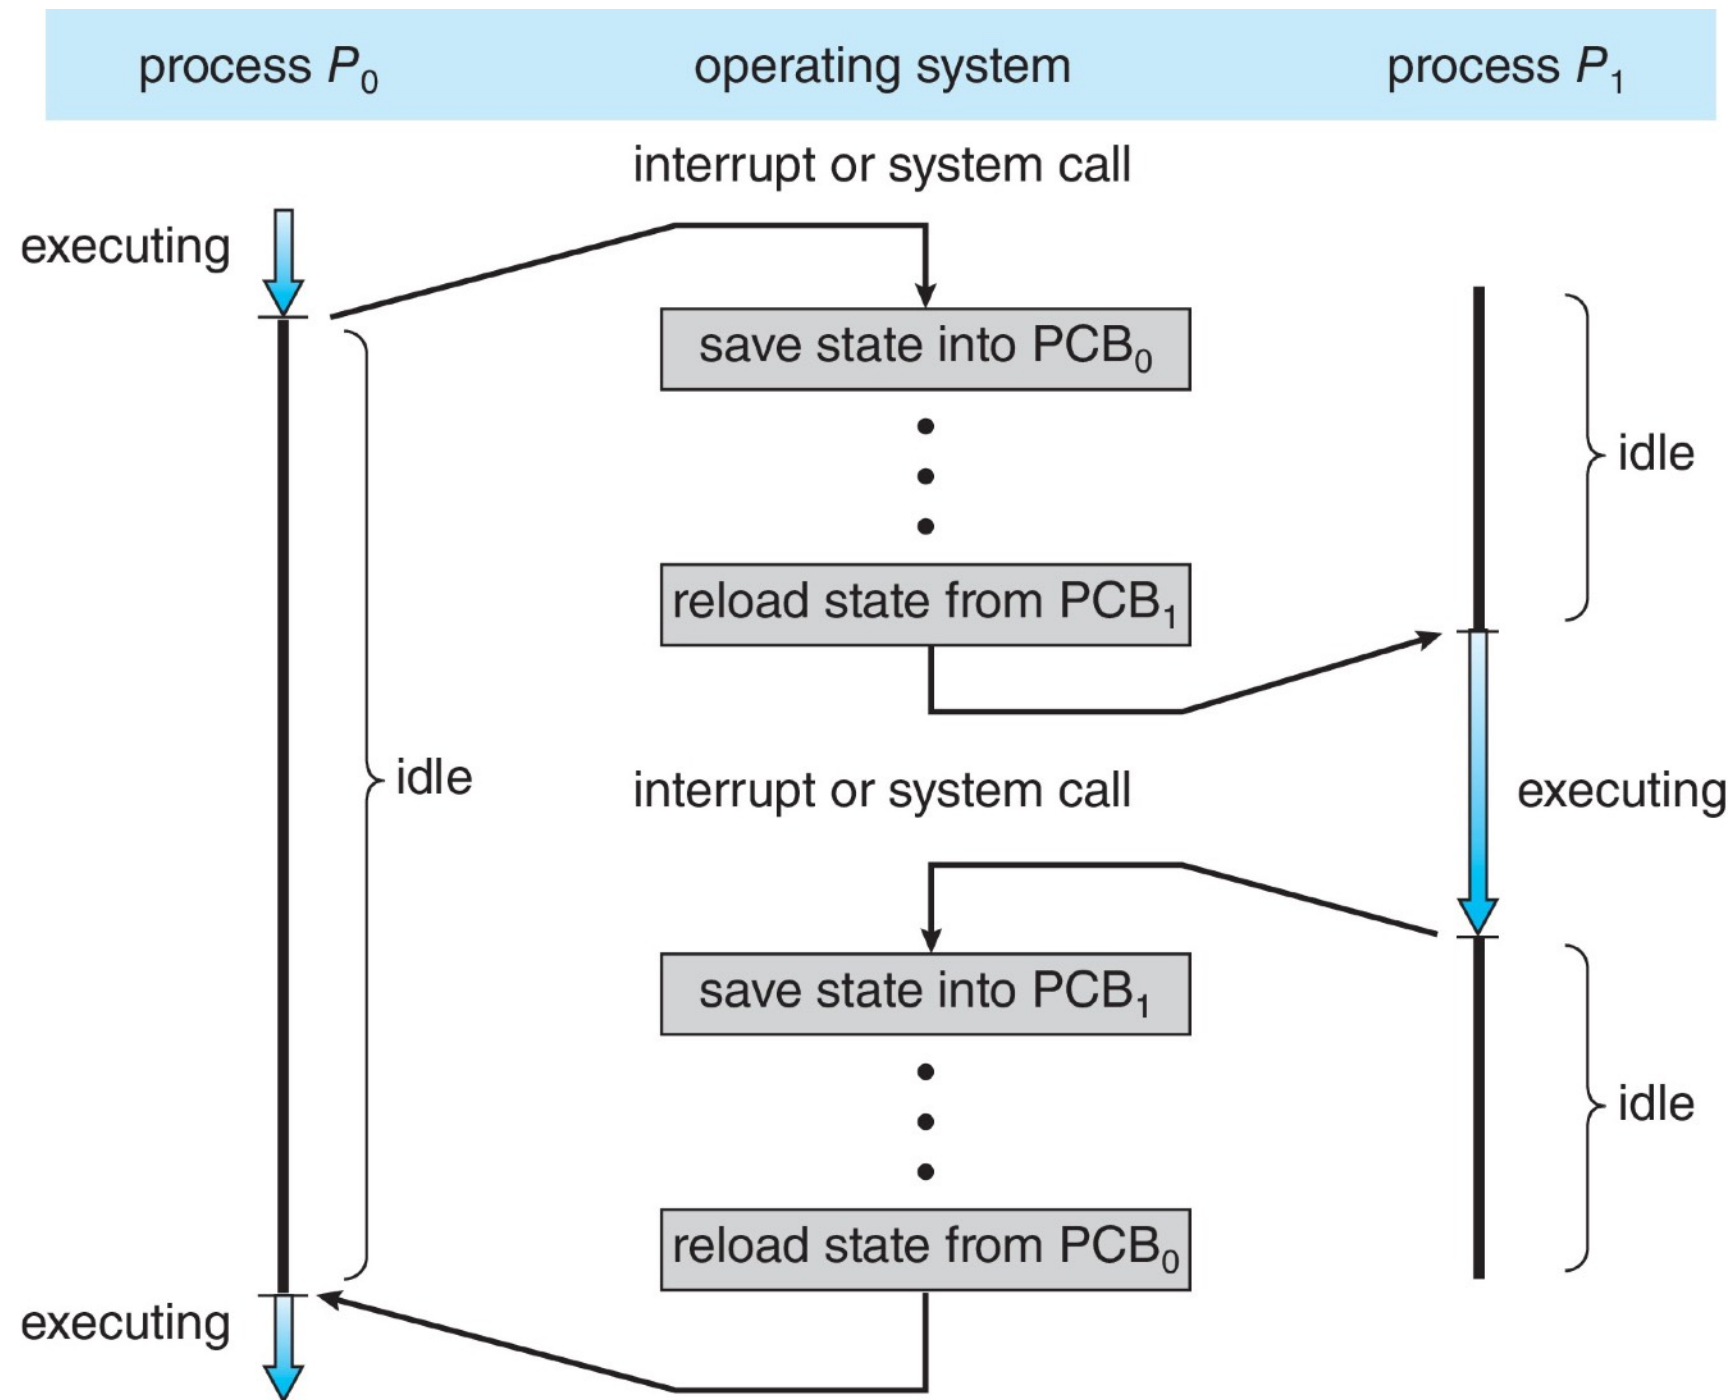
Context switching

- Definition: stopping a process and starting another one
 - it is a relatively expensive operation
 - time-sharing systems may do hundreds of context switches per second
 - the **context** includes the value of CPU registers, the process state, and memory management information
- OS starts executing a process in the ready state by loading hardware registers (PC, SP, etc) from its PCB
 - while a process is running, the CPU modifies the Program Counter (PC), Stack Pointer (SP), registers, etc.
- When OS stops executing a process, it saves the values of the registers (PC, SP, etc.) into its PCB
 - so that they can be restored the next time the process is selected for running on the CPU

Tradeoff

- Context switching is pure **overhead**, i.e., the system does no useful work
 - the cost of a context switch and the time between switches are closely related
 - but fast context switching is necessary for responsiveness
 - OS must balance the context switch frequency with the scheduling requirement (response time, fairness, etc.)

Context switching



Scheduler and dispatcher

- The scheduler selects a process to run next
- The dispatcher makes it happen
 - performs context switching
 - switches to user mode
 - jumps to the proper location in the user program