

Operating System Concepts

Lecture 5: Process Control

Omid Ardakanian
oardakan@ualberta.ca
University of Alberta

Today's class

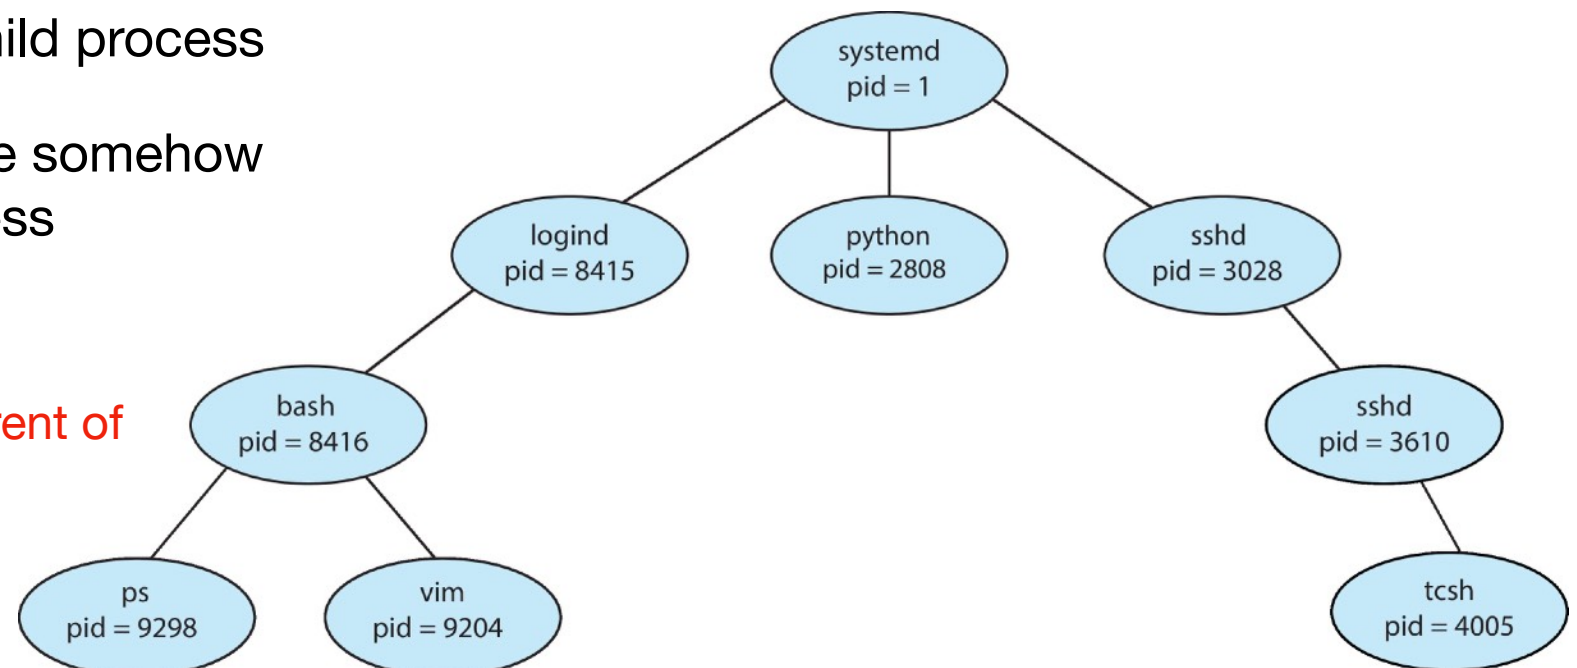
- Process Control
 - How to create a process?
 - How to terminate a process?
- Examples

Process management system calls in UNIX

- `getpid()` returns the current PID
- `fork()` copies the current process (a new PID is assigned to the child process)
- `execve()` loads a new binary file into memory (without changing its PID)
- `wait()` waits until one of its child processes terminates
- `waitpid()` waits until the specified child process terminates
- `_exit()` terminates a process
- `pause()` causes the calling process to sleep until a signal is delivered that either terminates the process or causes the invocation of a signal-catching function
- `nanosleep()` suspends execution of a process for at least the specified time or the delivery of a signal
- `kill()` sends a signal (interrupt-like notification) to another process (if permitted)
- `sigaction()` sets handlers for signals except `SIGKILL` and `SIGSTOP`
 - these signals cannot be caught, blocked, or ignored

Creating a process with the fork system call

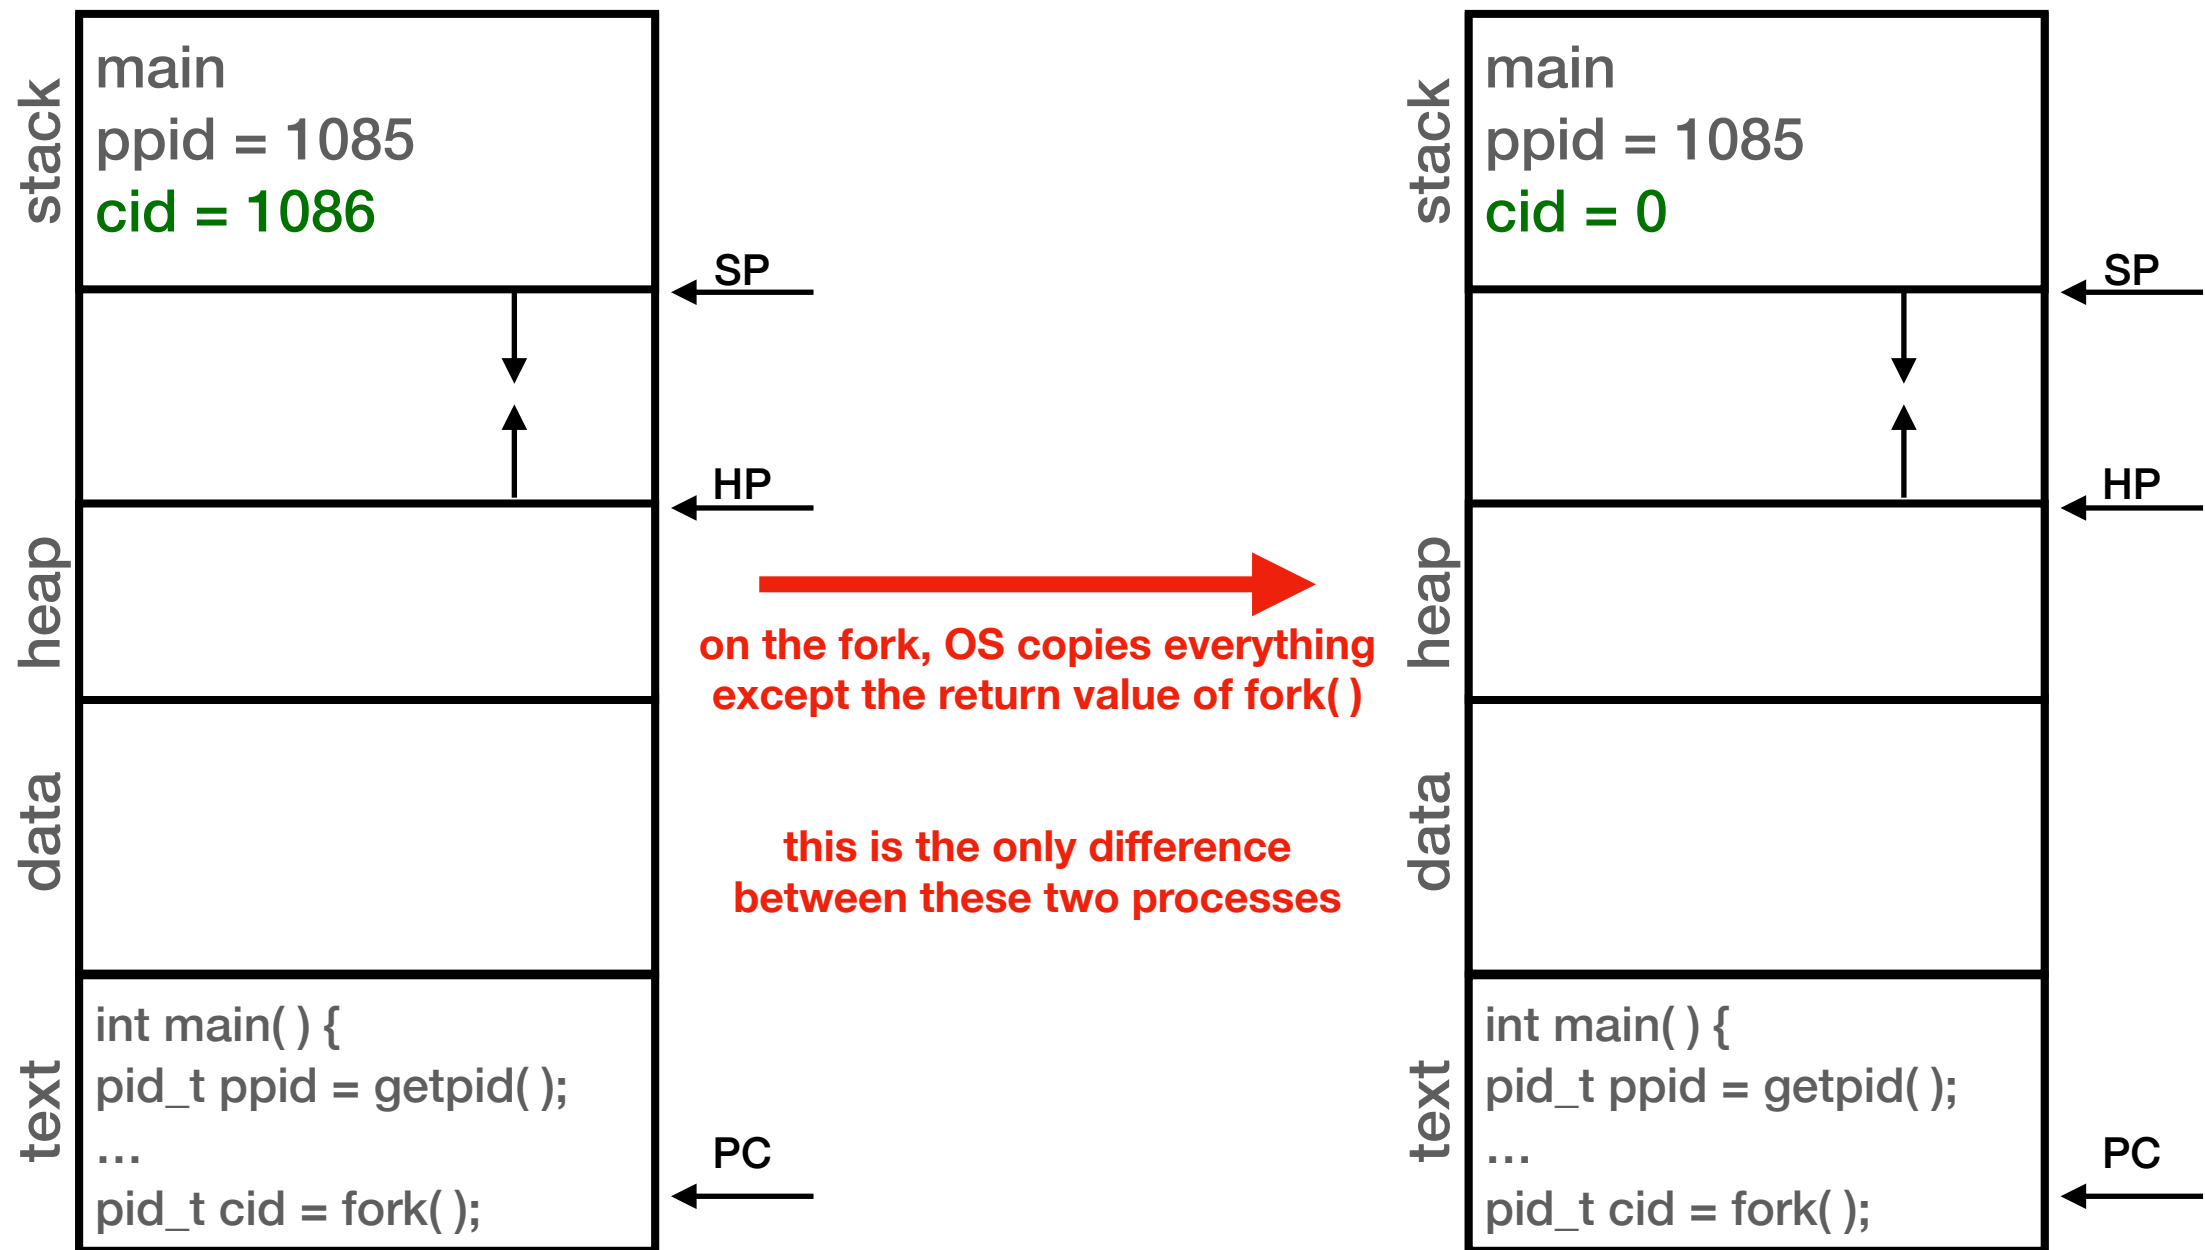
- The `fork()` system call creates a **child** process that inherits a copy of its **parent**'s memory, open file descriptors, CPU registers, etc.
- Both parent and child processes execute from the instruction following `fork()`
 - does the child or parent process run first? **We don't know!**
- The return value from `fork()` is of type `pid_t` (like an integer)
 - when > 0 : running in (original) parent process and the return value is pid of new child
 - when $= 0$: running in new child process
 - when < 0 : error! must handle somehow still running in original process



the `systemd` process is the root of this tree (parent of all user processes created when the system has booted); it has PID = 1

What happens on a fork?

```
pid_t cid = fork( );
```



Common problems with fork

The `fork()` system call is

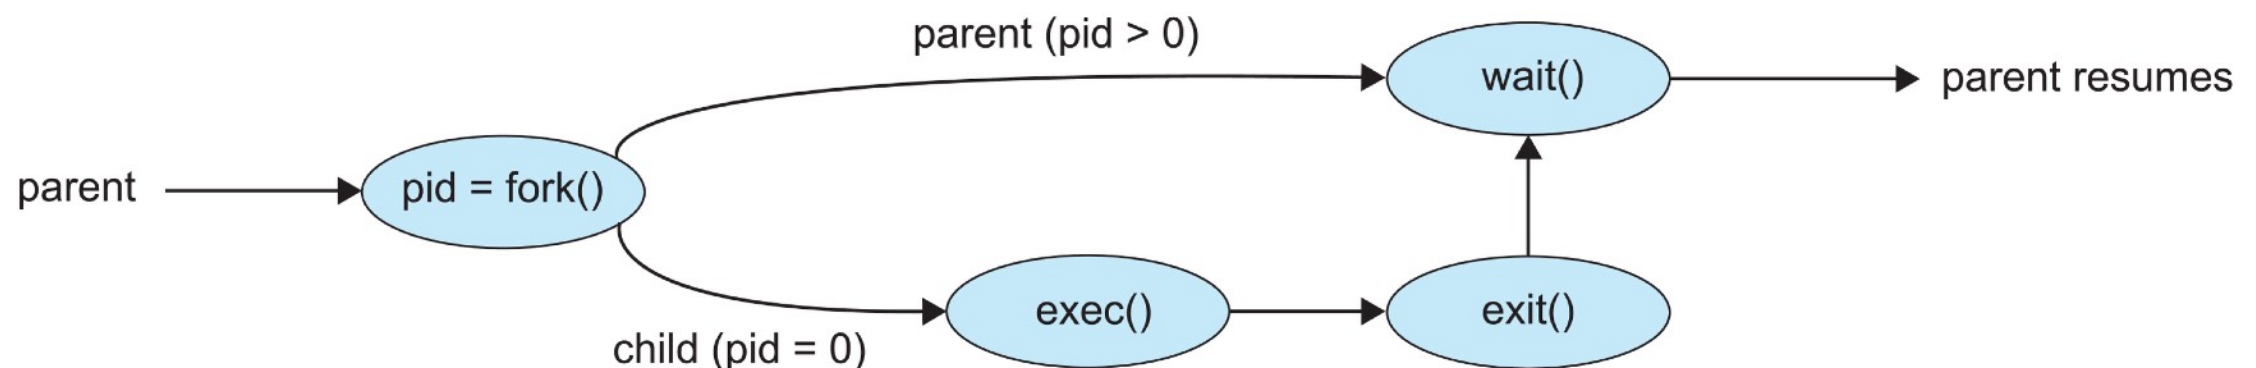
- inefficient and slow
 - the cost of copying the entire address space of a process is high
- insecure
 - the parent process must explicitly remove states that the child process does not need (scrubbing secrets from memory)
- not thread-safe
 - the child process created by the fork system call will have a single thread only (a copy of the calling thread)
 - **Problem:** one thread doing memory allocation and holding a **heap lock**, while another thread forks. Any attempt to allocate memory in the child (and thus acquire the same lock) will immediately **deadlock** waiting for an unlock operation that will never happen
 - **Solution:** not using fork in a multithreaded process, or calling `exec` immediately afterwards

Program loading with the exec system call

- The `execve ()` system call allows a process to load a different program and start execution from its main function
 - allows a process to specify the number of arguments (`argc`) and the string argument array (`argv`) that must be sent to the new process
- If the call is successful, the same process runs a different program
 - code, data, **stack and heap** sections are overwritten
- We normally call `execve ()` right after calling `fork ()`
 - hence, the memory copied during `fork ()` is useless
 - in performance-sensitive applications, the `vfork ()` system call allows creating a process without creating an identical memory image; in this case child process must call `execve ()` immediately — undefined behaviour if it doesn't

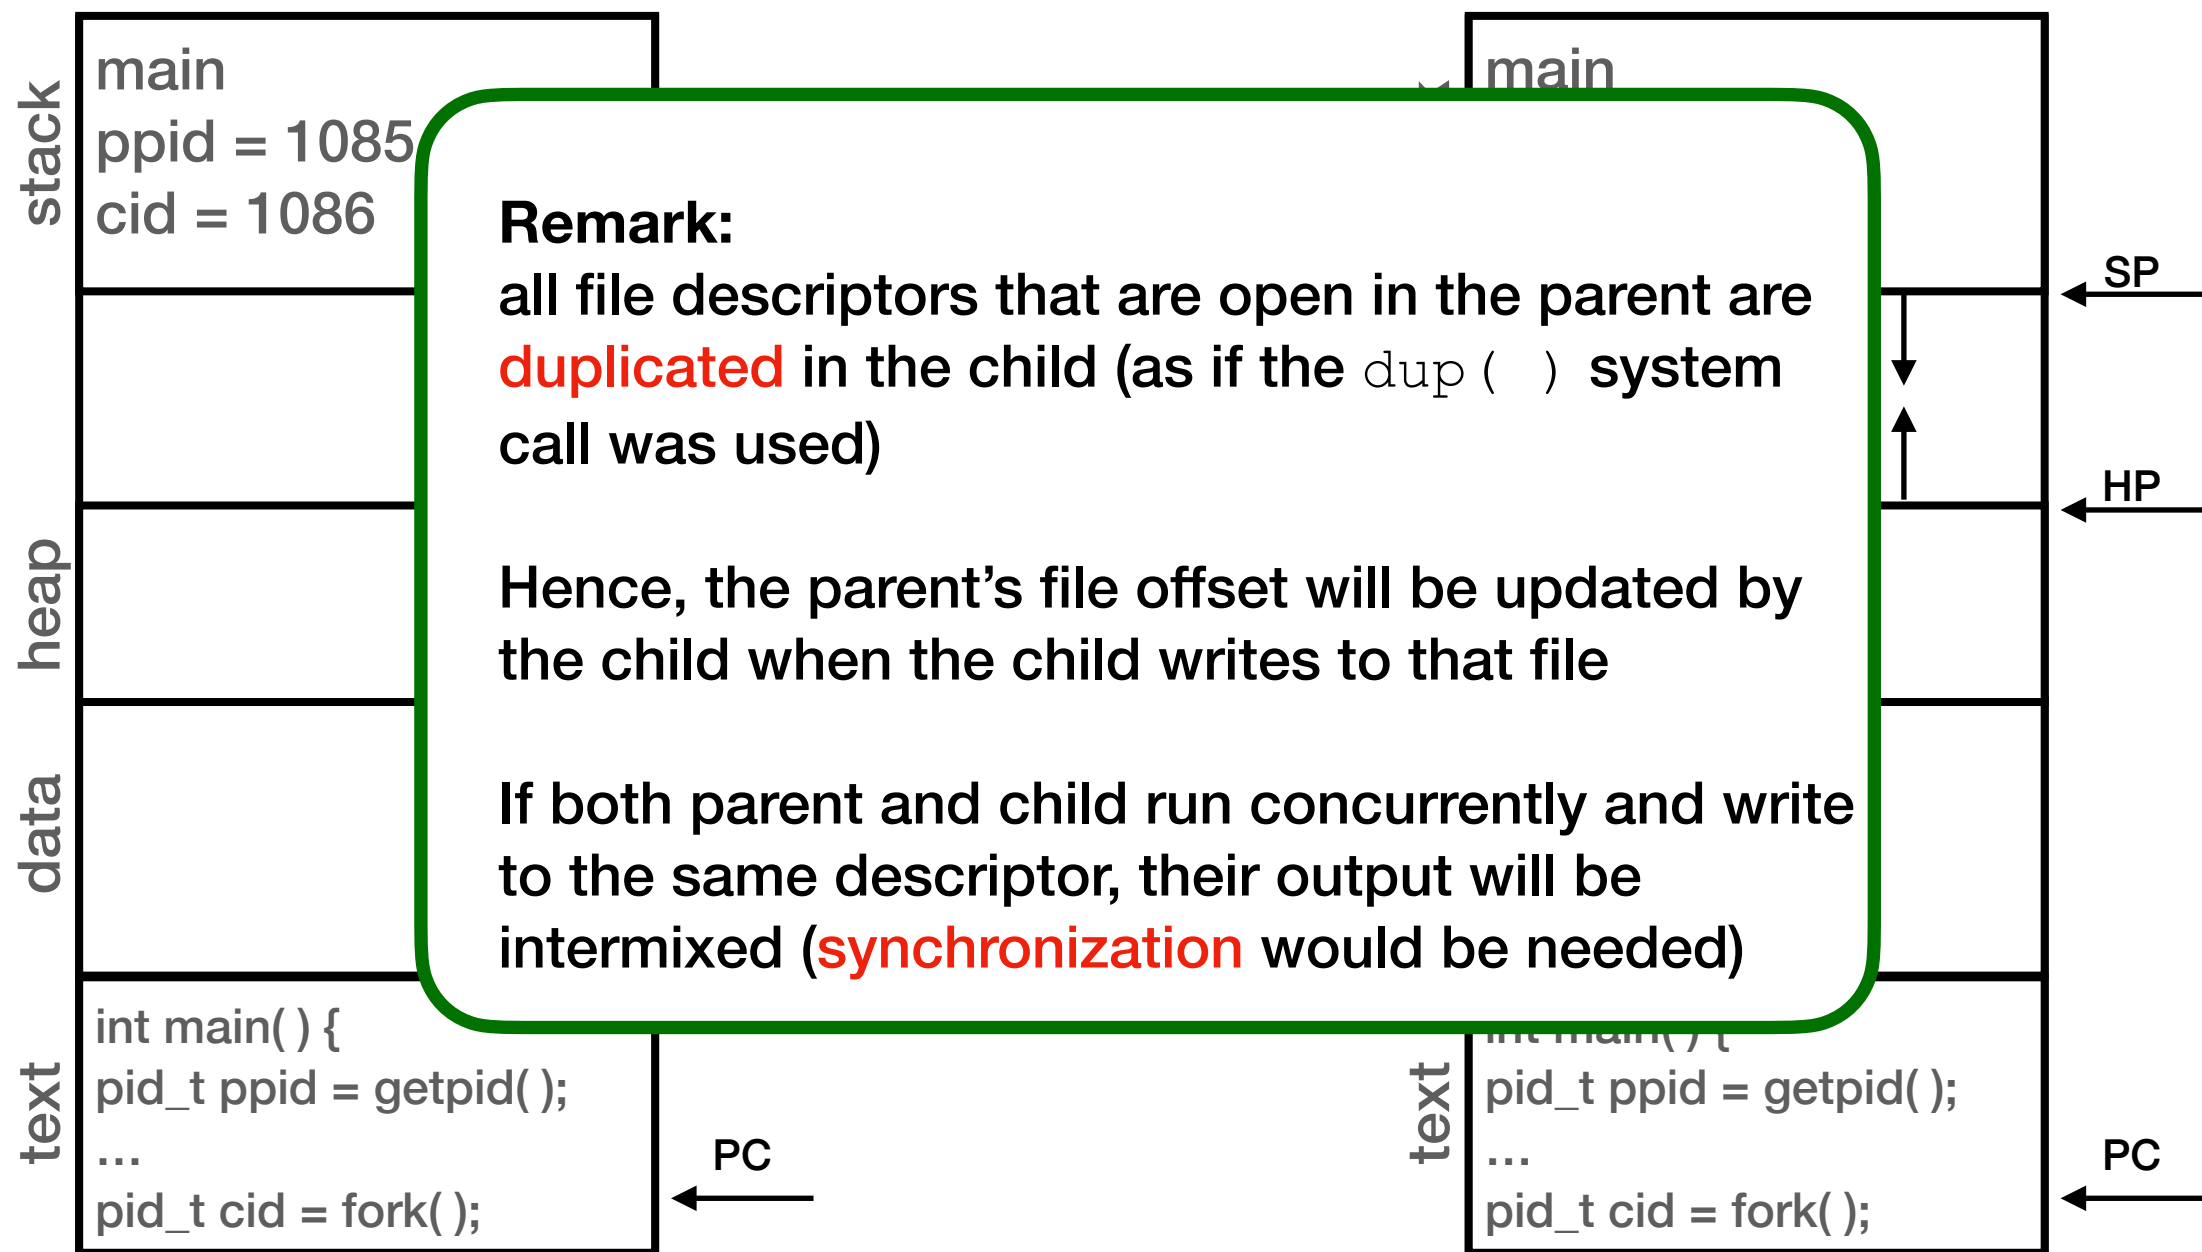
Waiting for the child process to terminate

- The parent can execute concurrently with its children or can wait until some or all of them terminate
- The `wait()` system call enables the parent process to wait for a child process to change state (e.g. terminate)
 - puts the parent to sleep waiting for a child's result
 - when a child calls `_exit()`, the kernel notifies the parent by sending the `SIGCHLD` signal to the parent; this unblocks the parent and returns the child's return value along with the child's PID
 - if there are no children alive, `wait()` returns immediately
 - also, if there are zombies waiting for their parents, `wait()` returns one of the return values immediately (and deallocates the zombie)



What happens on a fork?

```
pid_t cid = fork( );
```



Terminating a process

- Process termination is the ultimate resource reclamation by the OS
 - closes all open files, connections, etc.
 - deallocates memory and most of the OS data structures supporting the process
 - checks if parent is alive
 - if so, holds the **exit status** until parent requests it; in this case, process does not really die, but it enters the zombie state (**Why?**)
 - if not, it deallocates all data structures; the process is dead at this point
 - cleans up all waiting zombies

Normal and abnormal termination

- A process can terminate normally by returning from `main`, or directly calling the standard C library function `exit()` or the system call `_exit()`;
 - open file descriptors are closed; children are inherited by the `init` process
- When the `main` function returns, `exit()` is called indirectly
 - `exit()` calls all exit handlers that have been registered using `atexit()` — a glibc function
 - `_exit()` does not call exit handlers
- For abnormal termination of a process, call `abort()` which generates `SIGABRT`
 - functions registered using `atexit()` are not called
 - it may not close open files or flush stream buffers!
- A process can terminate a child using the `kill()` system call
 - `kill(cid, SIGKILL)`

Zombie and orphan processes

- A process that has terminated, but its parent has not (yet) called `wait()` becomes a **zombie**
 - the `ps` command prints the state of a zombie process as `Z`
- A process becomes **orphan** when its parent terminates while it is still running
 - hence the parent terminates without invoking the `wait()` system call
 - UNIX and Linux systems assign the `init` process (PID=1) as the new parent of the orphan process (aka **reparenting**)
 - you can check this `if (getppid() == 1)`
 - the `init` process periodically calls `wait()` allowing the exit status of any orphaned process to be collected and their process table entries be deleted

Fork example

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>

int main() {
    ...
    pid_t ppid = getpid();           // store parent's pid
    pid_t pid = fork();              // create a child
    if(pid == 0){                    // child continues
        printf("Child pid: [%d]\n", getpid());
        ...
    } else if (pid > 0) {             // parent continues
        printf("Parent pid: [%d] Child pid: [%d]\n", ppid, pid);
        ...
    } else {
        perror("fork failed!");
        exit(1);
    }
    ...
}
```

Run the `ps` command to
check the processes' IDs

Combining fork and wait

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/wait.h>

int main() {
    ...
    pid_t ppid = getpid();           // store parent's pid
    pid_t pid = fork();              // create a child
    if(pid == 0){                    // child continues
        ...
    } else if (pid > 0) {            // parent continues
        ...
        pid_t cpid = wait(&child_status); // how was it stopped or terminated
    } else {
        perror("fork failed!");
    }
    ...
}
```

Combining fork and exec

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/wait.h>

main() {
    ...
    pid_t ppid = getpid();           // store parent's pid
    pid_t pid = fork();              // create a child
    if(pid == 0){                    // child continues
        execve("/bin/ls", arg0, arg1, ...); // mark the end with a null pointer
        // exec doesn't return on success! so if we got here, it must have failed!
        perror("exec failed!");
    } else if (pid > 0) {            // parent continues
        ...
        cpid = wait(&status);        // pass NULL if not interested in exit status
        if (WIFEXITED(status))       // true if the child terminated normally
            printf("child exit status was %d\n", WEXITSTATUS(status));
    } else {
        perror("fork failed!");
    }
    ...
}
```

Parent can kill its child!

```
#include <signal.h>
#include <unistd.h>
#include <stdio.h>

main() {
    ...
    int ppid = getpid();           // store parent's pid
    int pid = fork();              // create a child
    if(pid == 0){                  // child continues here
        sleep(10);                // child sleeps for 10 seconds
        ...
        exit(0);
    }
    else {                         // parent continues here
        ...
        printf( "Type any character to kill the child.\n" );
        char answer[10];
        gets(answer);
        if ( !kill(pid, SIGKILL) ) {
            printf("Killed the child.\n");
        }
    }
}
```


Other system calls for process control

- OS must include calls to enable special control of a process:
 - Priority manipulation:
 - the `nice(incr)` system call adjusts the process priority by adding `incr` to its nice value
 - lower nice values have higher scheduling priority
 - a process could be “**nice**” and reduce its share of the CPU by adjusting its nice value
 - Debugging support:
 - the `ptrace()` system call allows a process to be put under control of another process by having its system calls intercepted; very useful for breakpoint debugging
 - the other process can check the arguments of the system call made by the process being traced, set breakpoints, examine registers, etc.
 - Alarms and time:
 - the `sleep()` system call puts a process on a timer queue waiting for some number of seconds, supporting alarm functionality

Process termination in UNIX systems

- the `kill` **system call** sends a signal to a process or process group based on the specified PID
- the `kill` **command** sends a `SIGTERM` signal by default
- the `killall` **command** sends an arbitrary signal to processes based on process name

Process monitoring in UNIX systems

- `ps` displays information about a selection of the active processes
 - `ps -el` lists complete information about all processes that are currently active in the system
 - `ps -u [username]` lists all processes created by a specific user
- `top` provides a dynamic real-time view of a running system (repetitive update on active processes)
- `pstree` displays a tree of processes

Shell

- Acts as a process control system
 - allowing programmers to create and manage a set of processes to do some tasks
 - Windows, Linux, MacOS have their own shells
- When you log in to a machine running UNIX, you create a shell process
- Every command launched in the shell is a child process of the shell process (an implicit `fork()` and `execve()` pair)
- The separation of `fork()` and `execve()` enables features like input/output redirection, pipes, etc.
 - the shell runs code after the call to `fork()` and before the call to `execve()`

Summary

- OS creates, deletes, suspends, and resumes processes
- OS allocates resources to active processes
 - memory, I/O devices, files
- OS schedules processes
 - context switches between them
- OS supports **interprocess communication** and provides **synchronization** mechanisms