

Operating System Concepts

Lecture 6: Signals

Omid Ardakanian
oardakan@ualberta.ca
University of Alberta

Other system calls for process control

- OS must include calls to enable special control of a process:
 - Priority manipulation:
 - the `nice(incr)` system call adjusts the process priority by adding `incr` to its nice value
 - lower nice values have higher scheduling priority
 - a process could be “**nice**” and reduce its share of the CPU by adjusting its nice value
 - Debugging support:
 - the `ptrace()` system call allows a process to be put under control of another process by having its system calls intercepted; very useful for breakpoint debugging
 - the other process can check the arguments of the system call made by the process being traced, set breakpoints, examine registers, etc.
 - Alarms and time:
 - the `sleep()` system call puts a process on a timer queue waiting for some number of seconds, supporting alarm functionality

Process termination in UNIX systems

- the `kill` **system call** sends a signal to a process or process group based on the specified PID
- the `kill` **command** sends a `SIGTERM` signal by default
- the `killall` **command** sends an arbitrary signal to processes based on process name

Process monitoring in UNIX systems

- `ps` displays information about a selection of the active processes
 - `ps -el` lists complete information about all processes that are currently active in the system
 - `ps -u [username]` lists all processes created by a specific user
- `top` provides a dynamic real-time view of a running system (repetitive update on active processes)
- `ps tree` displays a tree of processes

Shell

- Acts as a process control system
 - allowing programmers to create and manage a set of processes to do some tasks
 - Windows, Linux, MacOS have their own shells
- When you log in to a machine running UNIX, you create a shell process
- Every command launched in the shell is a child process of the shell process (an implicit `fork()` and `execve()` pair)
- The separation of `fork()` and `execve()` enables features like input/output redirection, pipes, etc.
 - the shell runs code after the call to `fork()` and before the call to `execve()`

Today's class

- Signals
 - Generation
 - Disposition
 - Blocking
- Interprocess communication (IPC)

Signals are software interrupts

- Signals provide a way of handling **asynchronous events** and sometimes **synchronous events** (e.g. divide by zero, segmentation fault)
 - they are small messages sent to a process, hence can be viewed as a form of IPC
 - often sent by the kernel, but can be sent from other processes too
- Mac OS X 10.6.8 and Linux 3.2.0 each support 31 different signals
- Every signal has a name that begins with SIG
 - names are defined by positive integer constants in `<signal.h>`
 - e.g. 9 is `SIGKILL`: kill; 11 is `SIGSEGV`: segmentation fault (e.g. dereferencing a null pointer), 8 is `SIGFPE`: erroneous arithmetic operation (e.g. divide by zero)
 - use `kill -l` to get a list of signals (this is architecture dependent)
- every signal has a default action associated with it; this action is performed once the signal is delivered unless a custom handler was installed
 - default actions: (a) terminate, (b) terminate with a **core dump**, (c) ignore, and (d) stop

Terminology

- We say that a signal is
 - posted (or generated or sent) if the event that causes it has occurred
 - delivered (or caught) if the action associated with it is taken
 - this action is referred to as **signal disposition**
 - pending if it was posted but not yet delivered
 - intermediate state between generated and delivered
 - signals will be pending if the target process blocks them
 - blocked if the target process does not want it delivered
 - the target process asked (using **signal mask**) the kernel to block that signal
- A signal can be process-directed or thread-directed (such as `SIGSEGV` and `SIGFPE`)
 - a process-directed signal may be delivered to any one of the threads that does not currently have the signal blocked

Signal generation

- Using a keystroke combination in shell
 - Ctrl-C causes `SIGINT` (interrupt) to be generated and sent to the foreground process
 - Ctrl-Z causes `SIGTSTP` (stop) to be generated and sent to the foreground process (note `SIGSTOP` \neq `SIGTSTP`)
 - Ctrl-\ causes `SIGQUIT` (quit) to be generated
- OS kernel wants to notify a process that its execution has led to a hardware exception (e.g., divide by zero, floating-point overflow, segmentation fault)
- The `kill` command or the `kill(pid_t pid, int sig)` system call is used
- A software condition occurs, e.g.
 - `SIGURG` (generated when out-of-band data arrives over a network connection)
 - `SIGPIPE` (generated when a process writes to a pipe that has no reader)
 - `SIGALRM` (generated when a timer set by the `alarm` function expires)

POSIX.1 reliable-signal

- A process can send a signal to another process or a group of processes
 - using the `kill()` system call
 - **only if it has permission:** the real or effective user ID of the receiver is the same as that of the sender
- A process can send a signal with accompanying **data** to another process (just like IPC)
 - using the `sigqueue()` function
- A process can send itself a signal using the `raise()` function
 - similar to `kill(getpid(), sig)`
- Waiting for a signal
 - the `pause()` system call puts a process to sleep until any signal is caught; this signal can be generated by the `alarm()` function
 - it returns only when a signal was caught and the signal-catching function returned

POSIX.1 reliable-signal

- Signal dispositions
 1. `SIG_DFL`: let the default action happens (in most cases terminate process or terminate process with a core dump; in some cases ignore)
 2. `SIG_IGN`: ignore the signal (except `SIGKILL` and `SIGSTOP` which can never be ignored as they are surefire ways of terminating/stopping a process)
 3. Address of a user-defined signal handler (a function that takes a single integer argument and returns void); kernel catches it by invoking this function
 - `SIGKILL` and `SIGSTOP` cannot be caught
 - a signal handler can return or call `exit`; if it returns, the normal sequence of instructions are executed
- System calls can be interrupted by a signal
 - some of them are automatically restarted, e.g., `ioctl`, `read`, `readv`, `write`, `writv`, `wait`, and `waitpid`

POSIX signal environment

- Examining or modifying the action associated with a particular signal (except for `SIGKILL` and `SIGSTOP`)
 - `sigaction()` supersedes `signal()` from earlier releases of the UNIX System
- Blocking signals in a signal set from delivery to a process
 - `sigprocmask()`
- Manipulating the signal set (a bit vector)
 - `sigemptyset()`, `sigfillset()`, `sigaddset()`, `sigdelset()`, `sigismember()`
 - when a signal is caught and the handler is entered, the current signal is automatically added to the signal mask of the process; this is to prevent subsequent occurrences of that signal from interrupting the signal handler
- Returning the set of signals that are blocked from delivery and currently pending for the calling process
 - `sigpending()`

Example of the sigaction() system call

```
#include <stdlib.h>
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
#include <signal.h>
```

```
struct sigaction {
    void (*sa_handler)(int);
    void (*sa_sigaction)(int, siginfo_t *, void *);
    sigset_t sa_mask;
    int sa_flags;
    void (*sa_restorer)(void);
};
```

```
void signal_callback_handler(int signum) {
    printf("Caught the signal!\n");
```

```
    // uncomment the next line to break the loop when signal is received
    // exit(1);
}
```

```
int main() {
    struct sigaction sa;
    sa.sa_flags = 0;
    sigemptyset(&sa.sa_mask);
    sa.sa_handler = signal_callback_handler;
    sigaction(SIGINT, &sa, NULL); // we are not interested in the old disposition

    // sigaction(SIGTSTP, &sa, NULL);
    while (1) {}
}
```

Process creation and signals

- When a process calls `fork`
 - child inherits parent's signal dispositions and signal mask
 - child starts off with a copy of parent's memory image, so signal-handlers are accessible
- When a process calls `exec`
 - the disposition of any signal being caught (not ignored) changes to its default action
 - the status of all other signals is left alone
 - the signal mask is preserved across the `exec` system call

Process state changes and signals


- When a process changes state (terminates, stops, or resumes), the kernel sends `SIGCHLD` to its parent
 - by default, this signal is ignored unless the parent installs a handler for it
 - the handler can call `waitpid(-1, &status, options)` with appropriate options (i.e. flags bitwise-or'ed together) to understand which process has changed state and record its new state
 - see the man page of `waitpid` for the list of options
 - **note:** if multiple signals come in at the same time, the signal handler is invoked only once, so `waitpid` should be called in a loop inside the handler as multiple processes might have changed state

Interprocess Communication

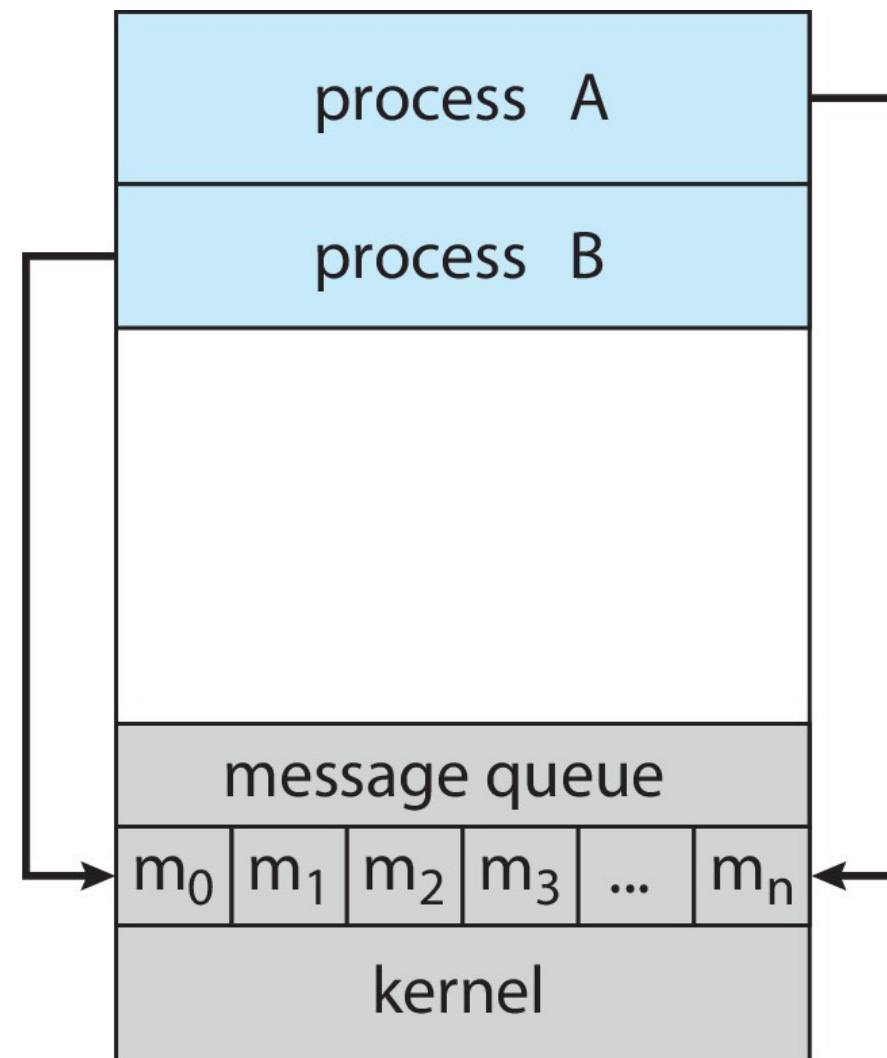
Cooperating processes

- Any two processes are either independent or cooperating
- Cooperating processes work with each other to accomplish a single task
 - improve performance
 - overlapping activities or performing work in parallel
 - improve program structure
 - each cooperating process is smaller than a single monolithic program
- They may need to share information
 - OS makes it possible!
 - shared memory vs. message passing approaches

Models of interprocess communication



(a)



(b)

Homework

- See examples posted on Canvas