

Operating System Concepts

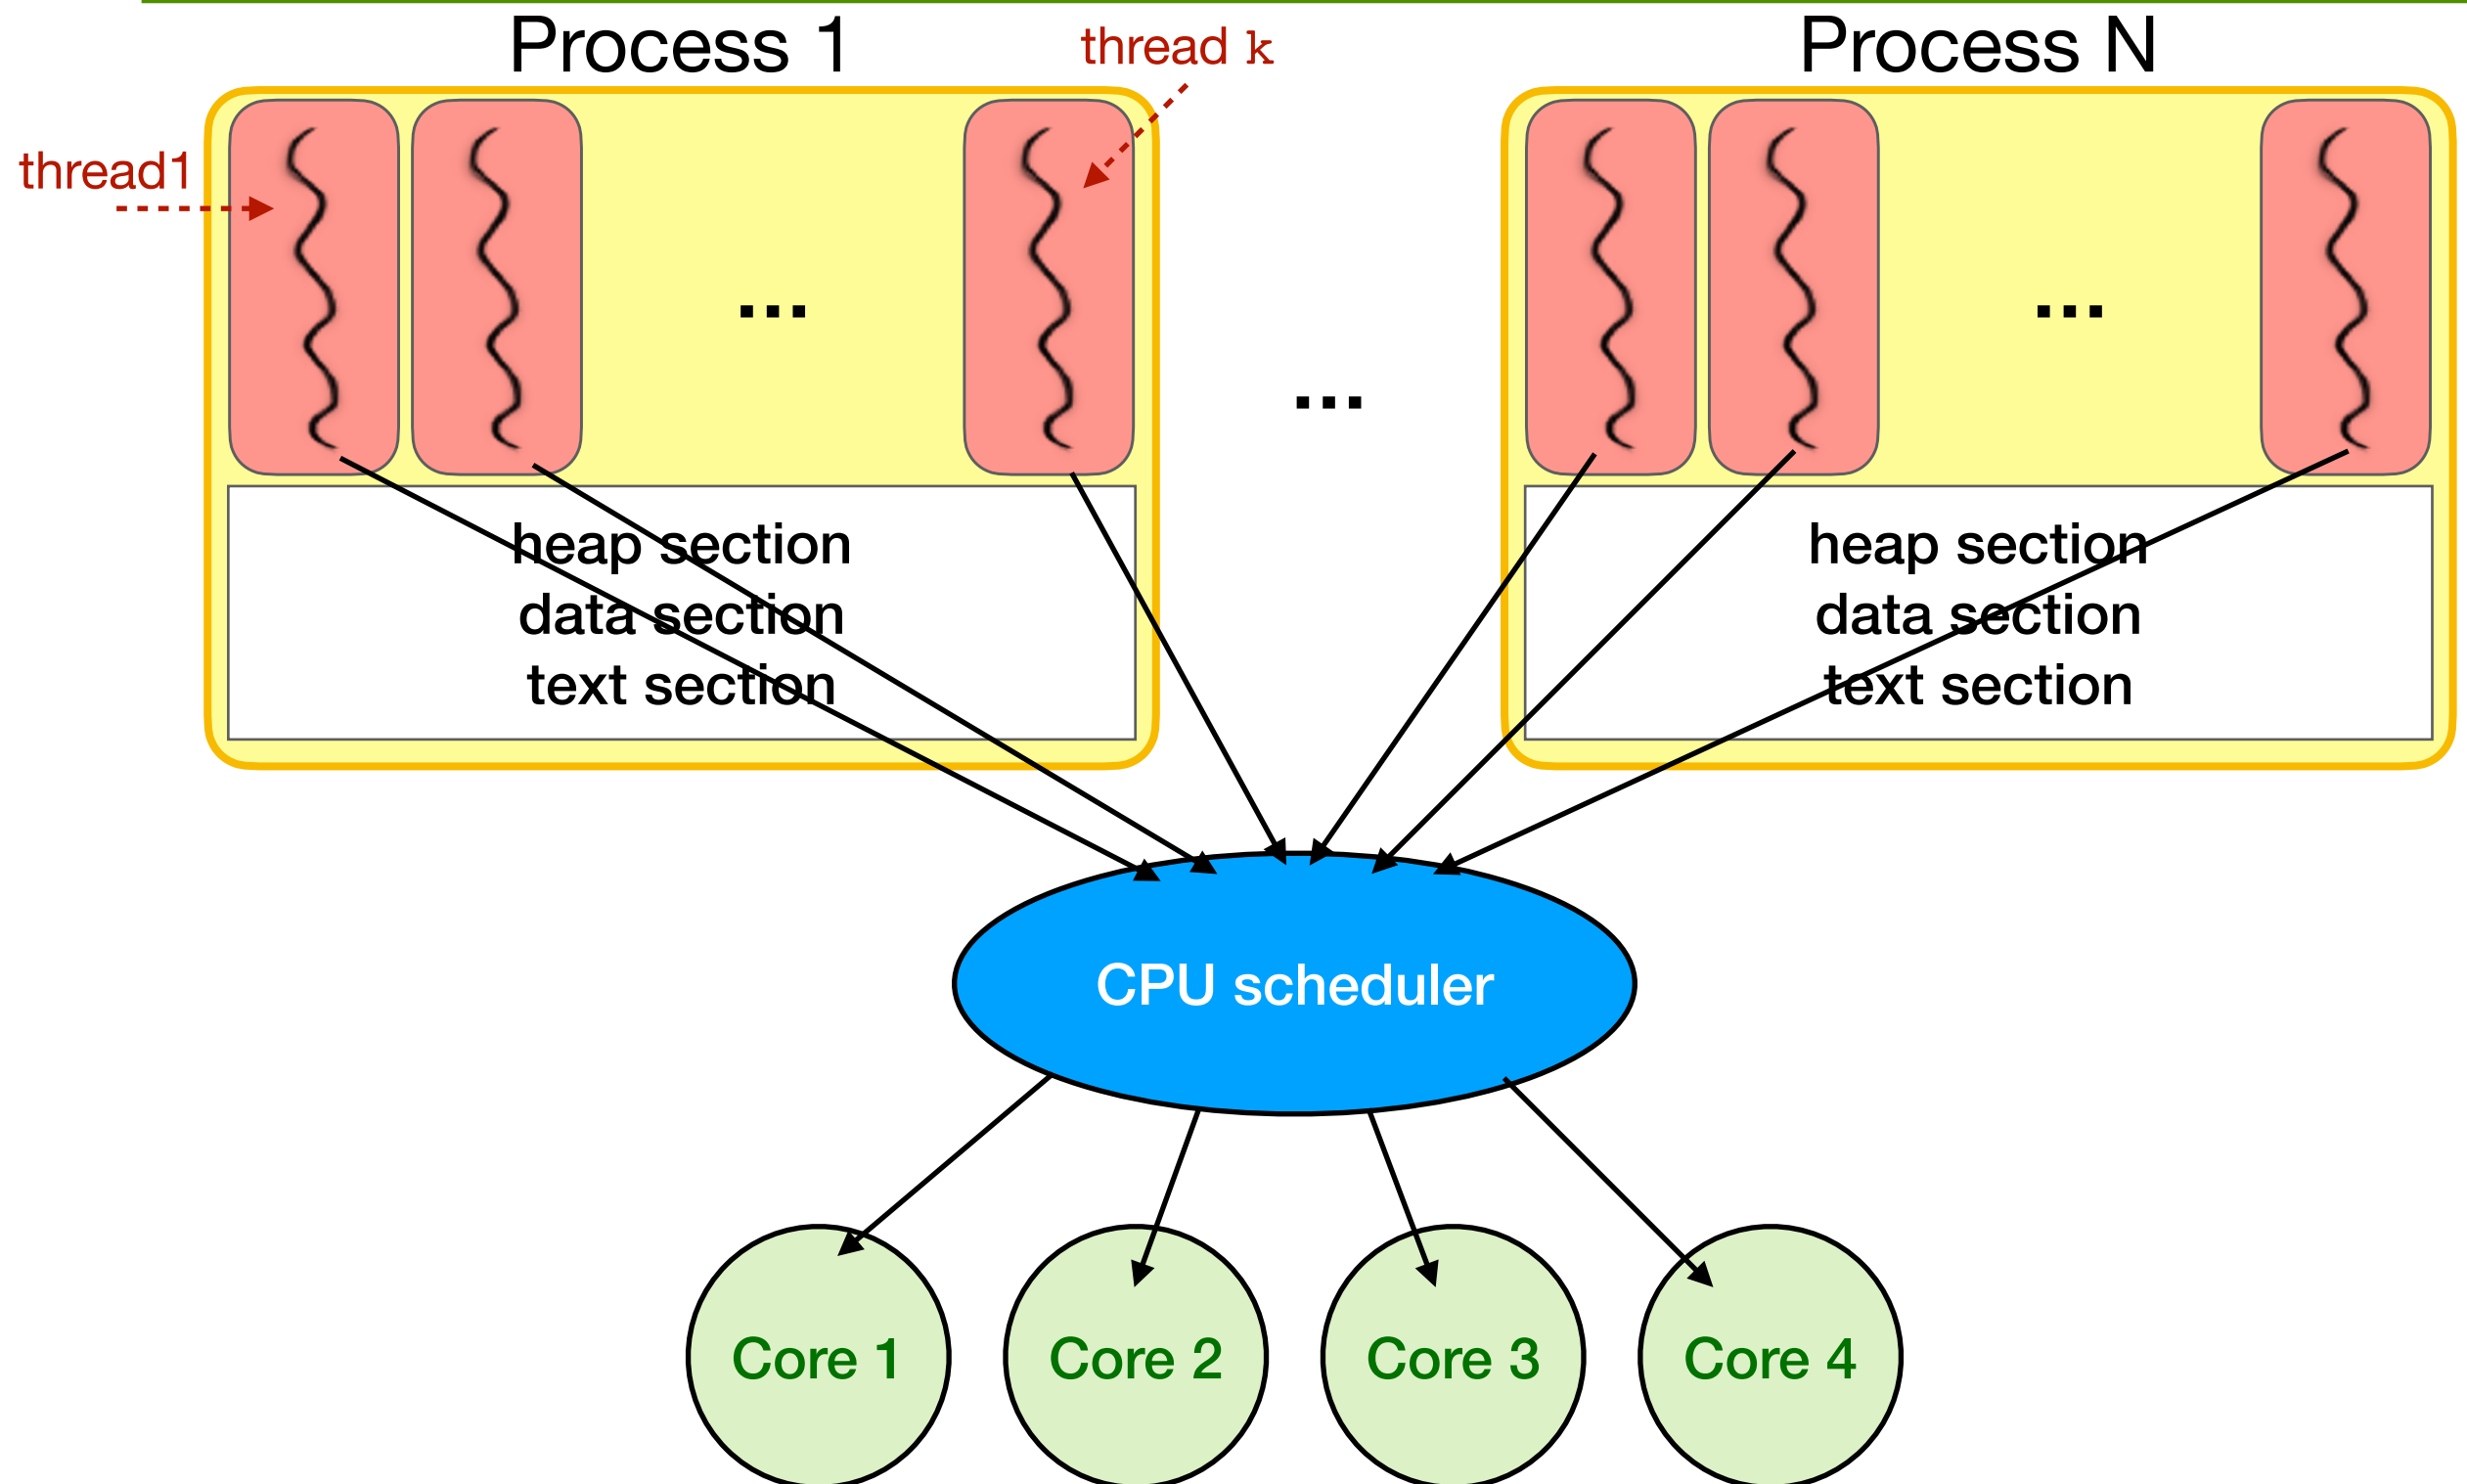
Lecture 12: CPU Scheduling - Part 2

Omid Ardakanian
oardakan@ualberta.ca
University of Alberta

Recap

- Turnaround time: time elapsed from task submission until task completion
- Response time: time elapsed from task submission until the first response is produced
- Waiting time: total time spent by a task waiting in the ready queue
- Throughput: rate of task completion
- CPU utilization: percentage of time CPU is busy (i.e. the ready queue is not empty)

Process versus thread scheduling




Assumptions

- We make a couple of simplifying assumptions today
 - one process per user
 - one thread per process
 - independent processes
 - just one processing core
- Scheduling algorithms were developed in the 70's when these assumptions were realistic

Today's class

- Scheduling algorithms
 - FCFS: First-Come, First-Served
 - RR: Round Robin
 - SJF: Shortest Job First

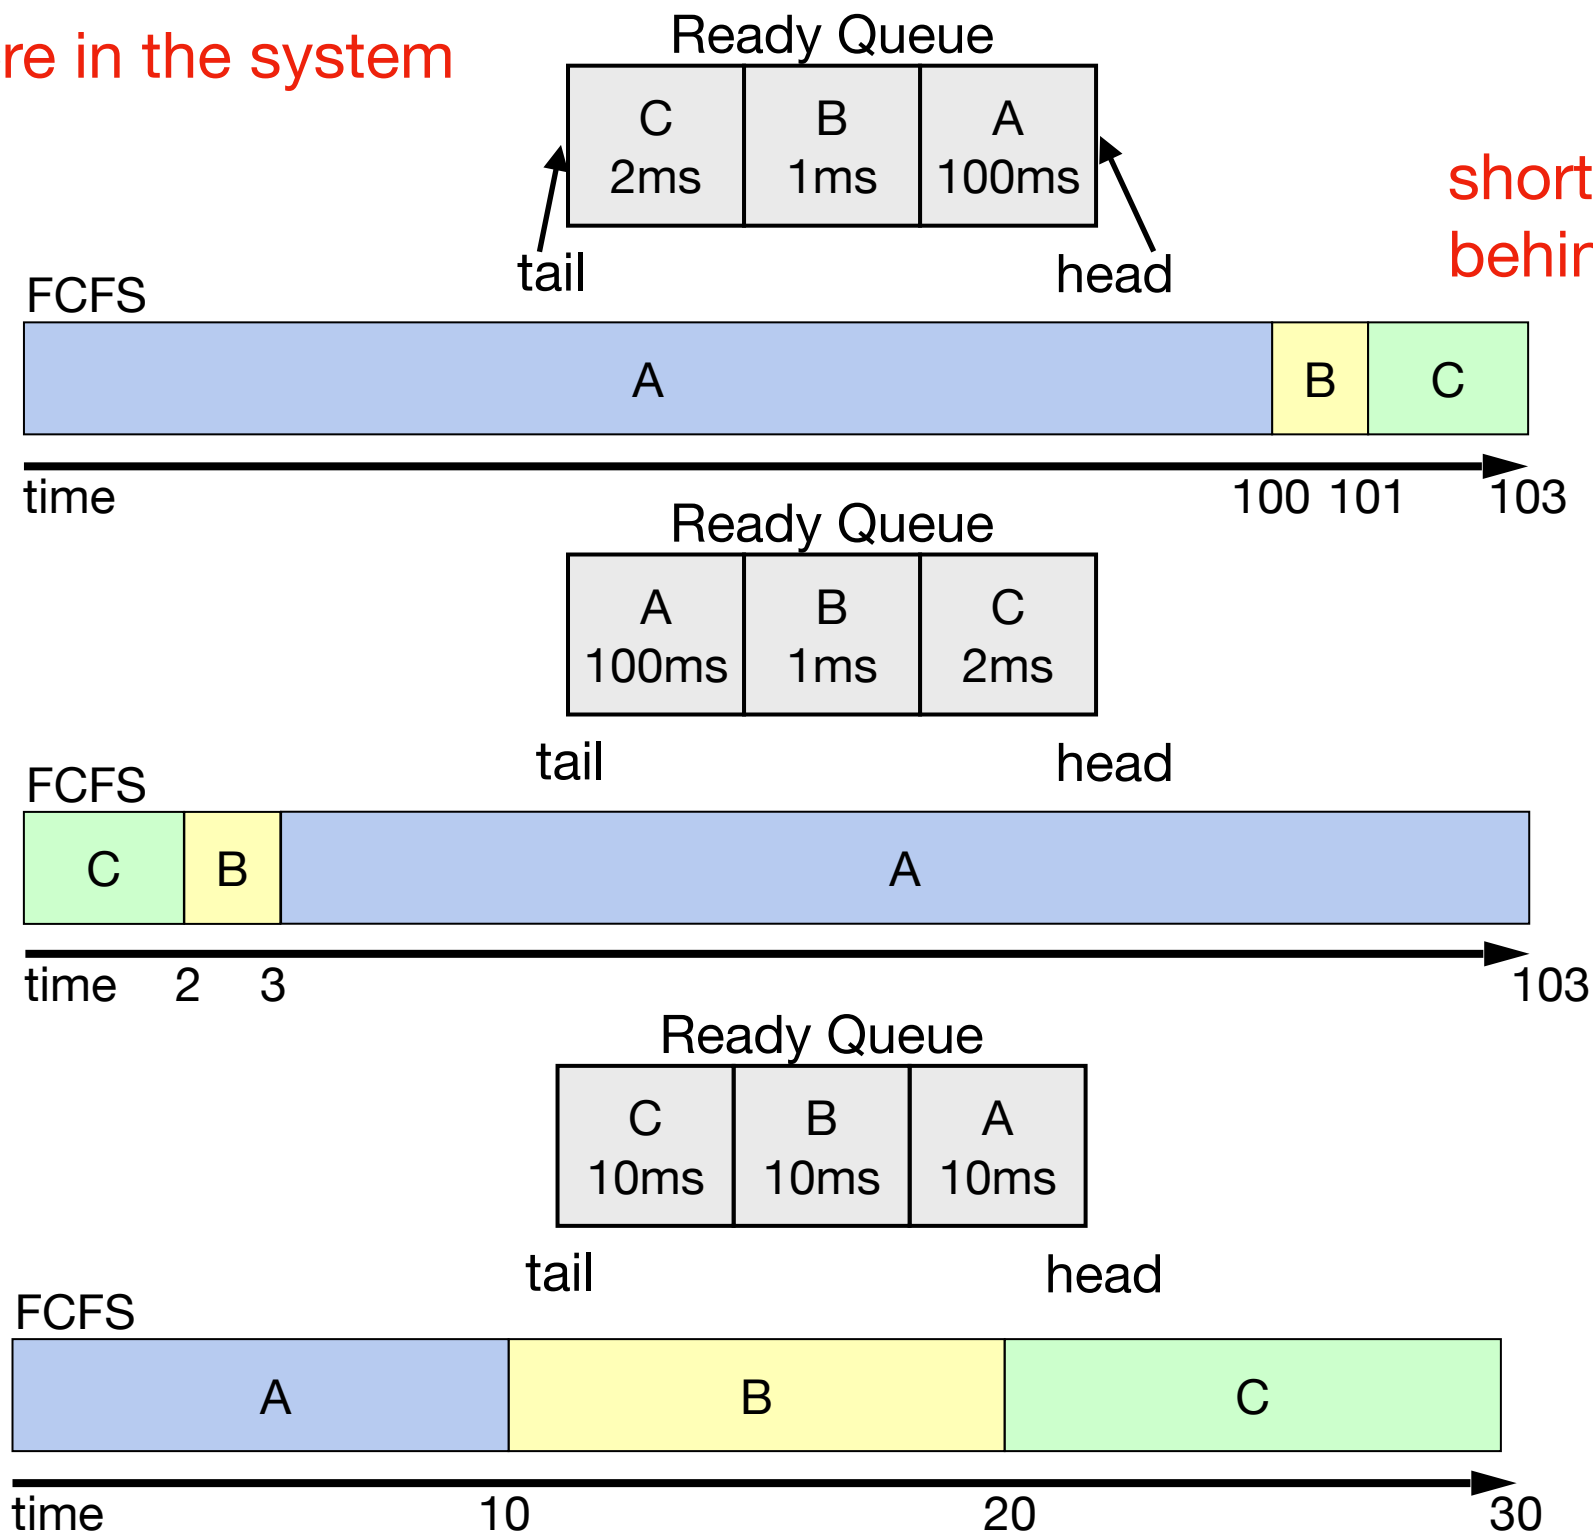


FCFS scheduling

- FCFS: First-Come-First-Served (or FIFO)
 - the scheduler executes jobs **to completion** in the order they arrive
 - in early FCFS schedulers, the job did not relinquish the CPU even when it was doing I/O
 - we will assume a FCFS scheduler that runs when processes are blocked waiting for I/O, but that is **non-preemptive**, i.e., the job keeps the CPU until it blocks (say on an I/O device)

FCFS example

A, B, C were in the system at time 0



short processes stuck behind long processes

Avg. turnaround time:
 $(100+101+103)/3=101.3$

Avg. waiting time:
 $(0+100+101)/3=67$

Avg. turnaround time:
 36

Avg. waiting time:
 1.6

Avg. turnaround time:
 20

Avg. waiting time:
 10

FCFS example

Process	Burst length	Arrival time
A	100	0
B	1	10
C	2	50

FCFS



Avg. turnaround time:
 $[100 + (101 - 10) + (103 - 50)] / 3 = 81.3$

Avg. waiting time:
 $[0 + (100 - 10) + (101 - 50)] / 3 = 47$

Pros and cons of FCFS scheduling


- Advantage: simplicity and low overhead
- Disadvantage:
 - average wait time is highly variable as short jobs may wait behind long jobs
 - if tasks are variable in size, FCFS can cause **poor response time** on average
 - If tasks are equal in size, FCFS is **optimal** in terms of **average response time**
 - not fair
 - may lead to poor overlap of I/O and CPU since CPU bound processes will force I/O bound processes to wait for the CPU, leaving the I/O devices idle

Round Robin scheduling

- Each task gets resources for a fixed period of time (**time quantum**)
 - if it does not finish its execution, it goes back in line (inserted at the end of the ready queue)
 - how to implement? add a timer and use a **preemptive** policy
- With quantum length Q ms, process waits **at most** $(N-1)*Q$ ms to run again if there is a total of N processes

Round Robin scheduling

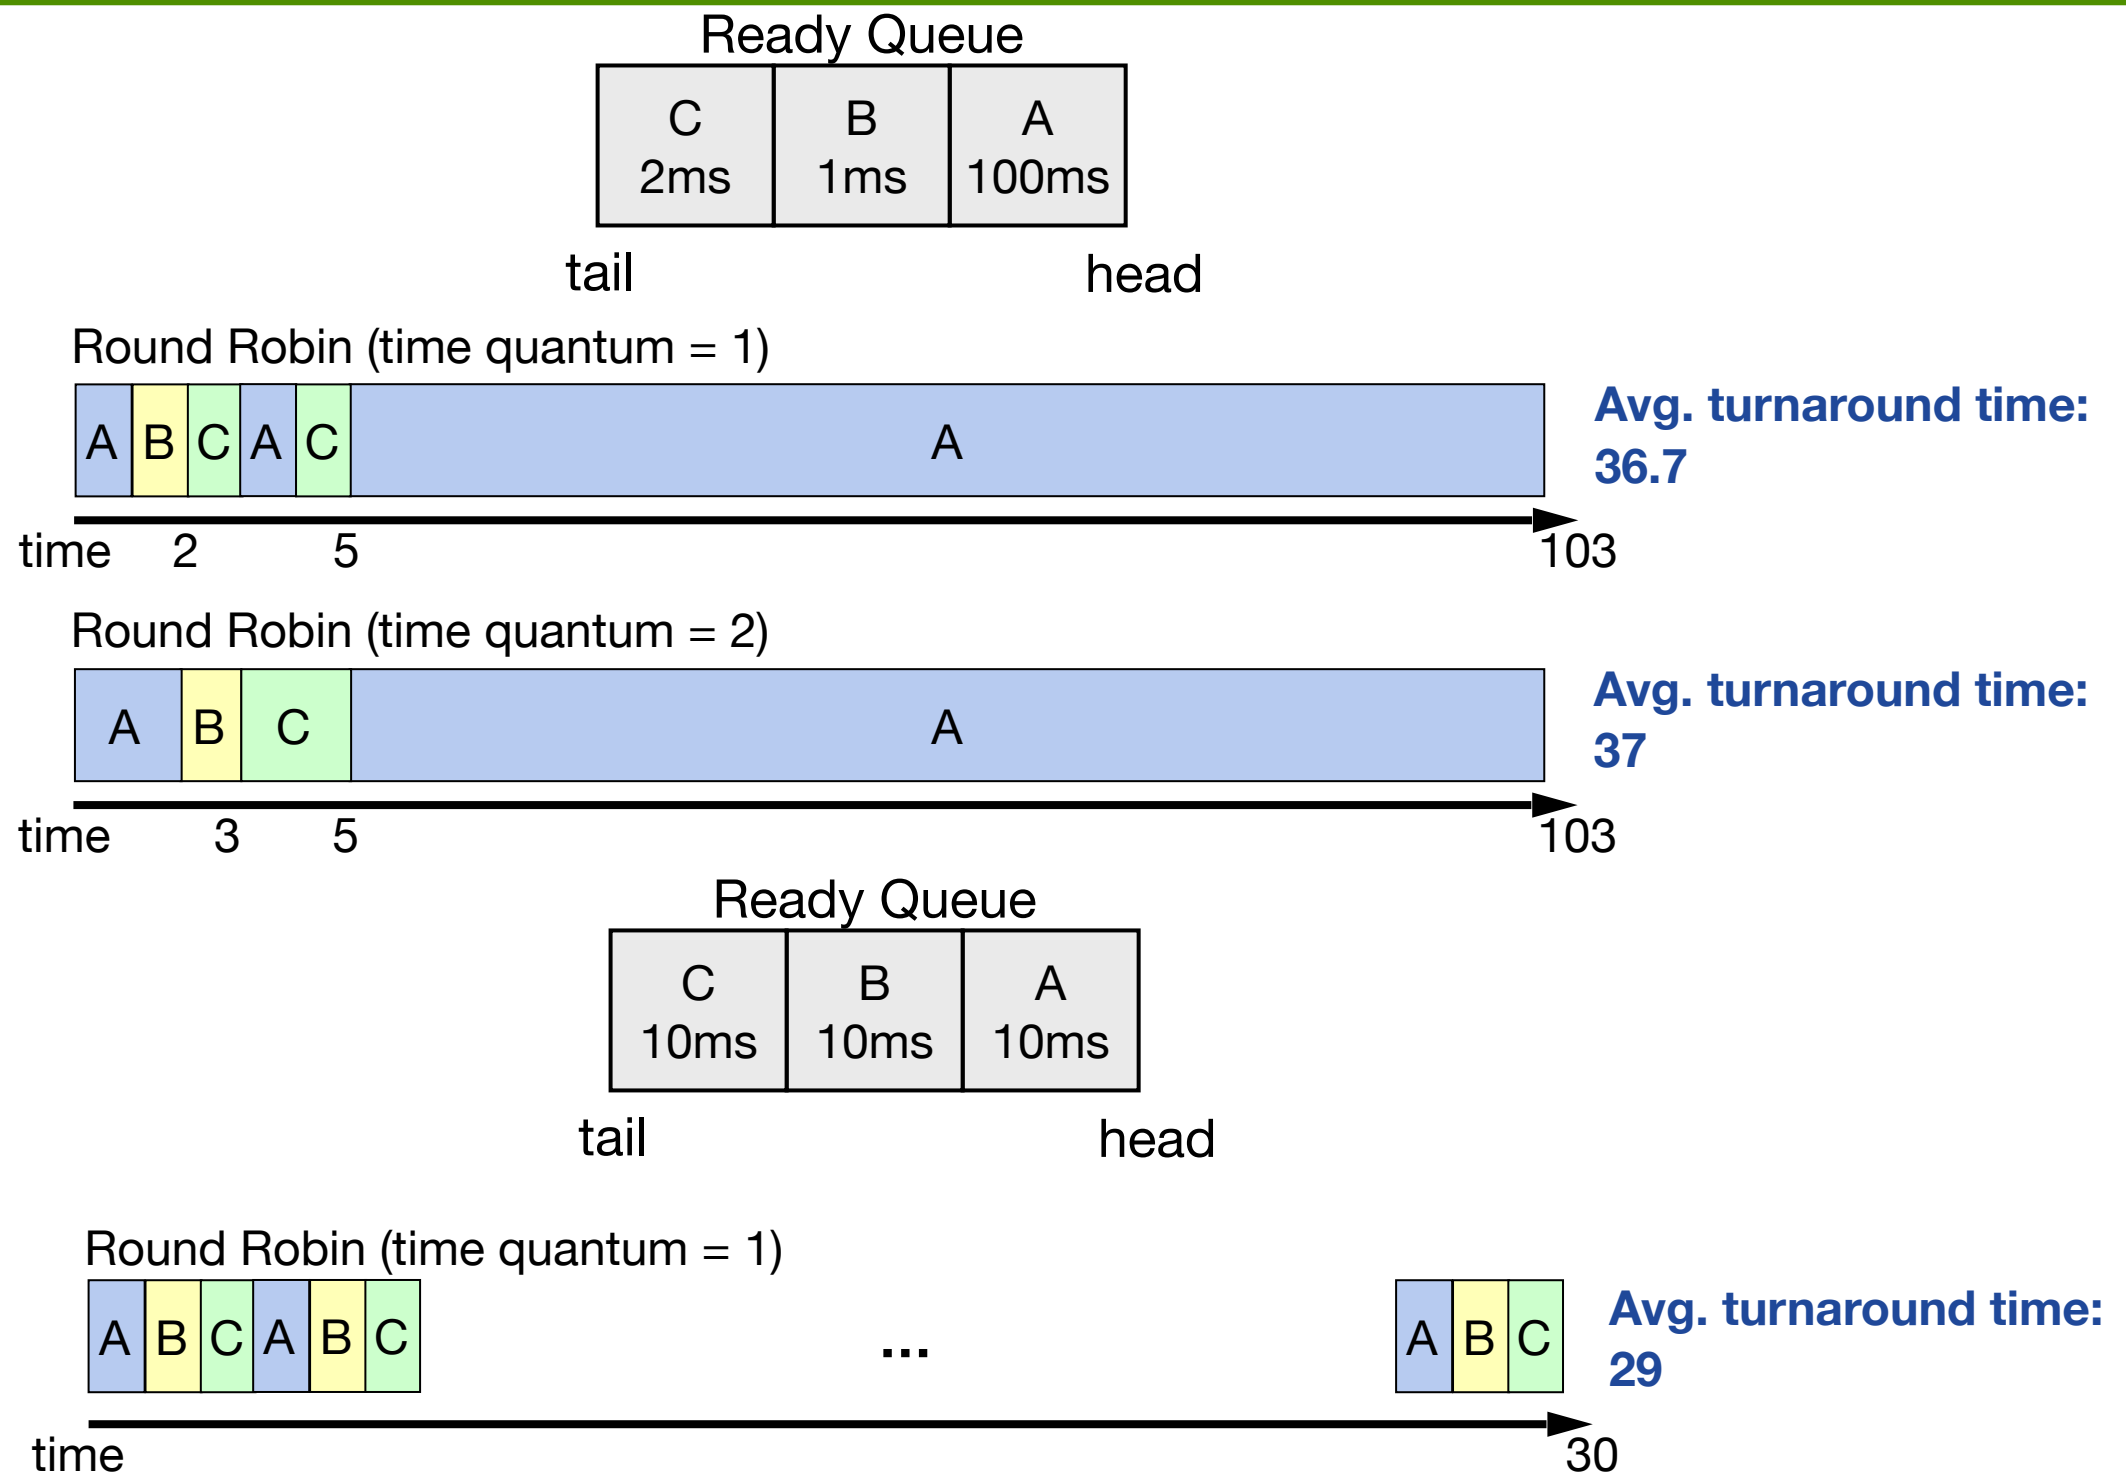
- Choosing the time quantum (Q) is key
 - if too long - waiting time suffers, degenerates to FCFS if processes are never preempted
 - if too short - throughput suffers because too much time is spent context switching (high overhead)
- It is possible to strike a balance between waiting time and throughput by selecting a time slice where context switching is roughly 1% of the time slice
 - today: typical time slice is 10-100ms, context switch time is 0.1-1ms



Pros and cons of Round Robin scheduling

- Variants of the round robin scheduling are used in most time-sharing systems
- Advantage: round robin is **fair**; each job gets an equal shot at the CPU
- Disadvantage: average waiting time can be bad if tasks are equal in size


RR example



assume all processes arrive at t=0

Setting the time quantum

suppose there is no context switching overhead and all process arrive at $t=0$




process	time
P_1	6
P_2	3
P_3	1
P_4	7

increasing the time quantum
can have different effects on
the average turnaround time


SJF scheduling

- Schedule the job that has the least (expected) amount of work (CPU time) to do until its next I/O request or termination
 - much less sensitive to the arrival order!
- If tasks are variable in size, Round Robin approximates SJF
- Advantages
 - provably optimal with respect to **minimizing the average waiting time**
 - works for preemptive and non-preemptive systems
 - preemptive SJF is called Shortest Remaining Time First (SRTF)
- Disadvantages
 - it is not possible to accurately predict the amount of CPU time that a job needs
 - with SRTF, long running CPU bound jobs can **starve** (if new short jobs keep arriving)


Scenario A



FCFS




Round Robin (time quantum = 1)




SJF



Scenario B




FCFS




Avg. turnaround time:
 $(10+20+30)/3=20$

Round Robin (time quantum = 1)



Avg. turnaround time:
29



SJF



Avg. turnaround time:
20

Comparing SRTF and SJF

Process	Burst length	Arrival time
A	9	0
B	10	0
C	3	2





But how do we know the burst length of a process?

- Burst length can be estimated based on previous burst lengths
 - the exponential moving average estimator gives more importance to the recent past
 - $\hat{b}[1] = b[1]$
 - $\hat{b}[t] = \eta b[t] + (1 - \eta)\hat{b}[t - 1]$ for $\eta \in (0,1]$
- Users can provide a burst length of their task
 - they can lie (declare a shorter burst length) to game the system
 - how to encourage truthfulness? terminate execution after the specified burst length has passed

SJF is starvation prone

Is SRJT starvation prone too?

Process	Burst length	Arrival time
A	9	0
B	10	0
C	3	2
D	4	11
E	6	14



it may remain in the ready queue forever

Homework

- Suppose 1 long process (with burst size of 100 units) and n short processes (with burst size of 1 unit each) arrive in the system in a random order at time 0. Under the SJF policy
 - what is the probability that a short process gets stuck behind the large task?
 - how long is the long process expected to wait?