

Operating System Concepts

Lecture 18: Synchronization Primitives

Omid Ardakanian
oardakan@ualberta.ca
University of Alberta

Today's class

- Synchronization primitives
 - Mutex locks
 - Condition variables

Programming abstractions for synchronization

- With low-level hardware support, programming languages provide **atomic operations for synchronization**
 - locks: can be held by at most one process/thread at a time; it is obtained before entering a critical section and released after exiting it
 - condition variables: provide conditional synchronization
 - semaphores: more general version of locks
 - monitors: connect shared data to synchronized primitives

Mutex locks

- A high-level programming abstraction; an object that only one thread can hold at a time
 - can be implemented as a **spinlock** which does busy waiting

```
void acquire() {  
    while(test_and_set(&lock))  
        ; /* spin */  
    /* critical section */  
    lock = false;  
}
```

- can be implemented as a blocking operation (next slide)

Blocking implementation of locks

```
class Lock {
public:
    void Acquire(); ← waits until lock is free and then grabs it
    void Release(); ← releases the lock and wakes up any waiters
private:
    int locked;    // lock state
    Queue Q;      // lock waiters queue
}
```

Mutual exclusion can be supported using locks (symmetric solution)

```
Lock milklock; ← initially free (not held by any process)
...
milklock.Acquire( ) ← acquired before
                    accessing shared data
if(milk == 0)
    buy_milk(); ← critical section
milklock.Release( ) ← released after
                    accessing shared data
```

Example of using locks (loose syntax)

```
void *malloc(size_t size) {
    heaplock.acquire();
    p = allocate memory of the specified size
    heaplock.release();
    return p;
}

void free(void *p) {
    heaplock.acquire();
    deallocate memory & put it back on free list
    heaplock.release();
}
```

recall that threads of a process share the heap section

How to implement locks on uniprocessors?

```
Lock::Acquire() {
    intr_disable();
    if (locked == 0) {    // lock is free
        locked = 1;
    } else {              // lock is held by another thread
        queue_add(Q, gettid());
        thread_block();  // put this thread to sleep
    }
    intr_enable();
}
```

```
Lock::Release() {
    intr_disable();
    if(queue_empty(Q)) {
        locked = 0;                // release the lock
    } else {
        thread_unblock(queue_remove(Q)); // put on ready queue
    }
    intr_enable();
}
```

CLI and STI (privileged) instructions are used to clear and set interrupts respectively

How to implement locks on multiprocessors?

- A thread/process executing a CLI instruction does not disable interrupts on other processors!
- So we have to use other hardware support to implement acquire and release methods
 - test_and_set
 - compare_and_swap

Implementation with compare_and_swap

- Compare the value against some expected value (in register), if they are the same, set the value in memory to a different value (in register)
 - if [addr] == r1 then [addr] = r2;
- Report either a boolean response or the old value
 - there are two variants

```
Lock::Lock {  
    locked = 0;  
}
```

```
Lock::Acquire() {  
    while(compare_and_swap(&locked, 0, 1) != 0)  
        ; // if busy, do nothing  
}
```

```
Lock::Release() {  
    locked = 0;  
}
```

Implementation with test_and_set

- If lock is free (value = 0), test&set reads 0, sets value to 1, and returns 0
 - the Lock is now busy: the test in the while fails (Acquire is complete)
- If lock is busy (value = 1), test&set reads 1, sets value to 1, and returns 1
 - continues to loop until a Release is executed

```
Lock::Lock {  
    locked = 0;  
}
```

```
Lock::Acquire() {  
    while (test_and_set(&locked) == 1) {  
        ; // if busy, do nothing  
    }  
}
```

```
Lock::Release() {  
    locked = 0;  
}
```

Avoiding busy-waiting as much as possible

- We can't eliminate busy waiting entirely but we can minimize its use to build a more efficient lock
 - instead of busy-waiting until lock is free, we busy-wait to atomically check the lock state and give up CPU if we find that the lock is busy
 - updating the lock state is a short critical section than the critical section in which shared data are updated and protected by the lock

```
class Lock {
public:
    void Acquire();
    void Release();
private:
    int locked; // lock state
    int guard; // safe to check the lock state
    Queue Q;
}

Lock::Lock {
    locked = 0; // lock is free initially
    guard = 0;
}
```

Test_and_set — minimal waiting

```
Lock::Acquire() {
    while(test_and_set(guard) == 1)
        ;    // spin until guard can be acquired
    if(locked != 0) {    // lock is busy
        queue_add(Q, gettid());
        guard = 0;    // set guard to 0 before blocking thread
        thread_block();    // block this thread
    } else {    // lock is free
        locked = 1;    // lock is acquired
        guard = 0;
    }
}
```

```
Lock::Release() {
    while(test_and_set(guard) == 1)
        ;
    if(!queue_empty(Q)) {
        // take a thread off the queue and pass the lock directly to it
        thread_unblock(queue_remove(Q));
    } else {
        locked = 0; // let go of lock as there is no waiting thread
    }
    guard = 0;
}
```

Observations

- Why does a thread set guard to 0 before blocking itself?
 - so that another thread can obtain guard to release the lock (**liveness issue**)
- Why does a thread that is releasing the lock pass it to next waiting thread (if any) rather than just releasing it and putting the awakened thread on the ready queue?
 - the awakened thread doesn't hold guard when it wakes up so it can grab the lock before entering the critical section (**mutual exclusion issue**)
- What if there's a context switch right before calling thread_block and the next thread releases the lock?
 - the first thread would be blocked for ever (**wakeup/waiting race**)
 - one solution is to ensure that two operations, i.e., releasing guard and blocking thread, are implemented as one atomic operation

Comparing to the “interrupt disable” solution

```
Lock::Acquire(Thread T) {  
    intr_disable();  
    if (locked == 0) {  
        locked = 1;  
    } else {  
        queue_add(Q, T);  
        thread_block(T);  
    }  
    intr_enable();  
}
```

```
Lock::Release() {  
    intr_disable();  
    if(queue_empty(Q)) {  
        locked = 0;  
    } else {  
        thread_unblock(queue_remove(Q));  
    }  
    intr_enable();  
}
```

Replace:

- `intr_disable()` with `while(test&set(guard));`
- `intr_enable()` with `guard = 0;`

Two-phase locks

- Spin for a small amount of time (spin phase) and if the lock cannot be acquired then put caller to sleep (sleep phase)
- What's the advantage?

Beyond locks

- Locks provide mutual exclusion but sometimes a thread has to wait only if a certain condition is true (synchronizing on a condition)
- Example: A producer thread puts data in a bounded buffer, a consumer thread takes them out.
What are the constraints for the bounded buffer?
 1. only one thread can manipulate buffer queue at a time (*mutual exclusion*)
 2. consumer must wait for producer to fill buffers **if all empty** (*scheduling constraint*)
 3. producer must wait for consumer to empty buffers **if all full** (*scheduling constraint*)

Condition variable

- An abstraction that supports **conditional synchronization**
 - a queue of threads waiting for a specific **event inside the critical section**
 - free memory is getting low, run the garbage collector
 - new data has arrived in a port, process it
 - the condition of the condition variable is defined based on the data protected by a mutex lock

Operations on condition variables

- Support three operations
 - `Wait()` usually takes a lock (to be released) as a parameter
 - atomically release lock and go to sleep (block the thread until signalled)
 - reacquire lock upon waking up
 - `Notify()` — historically called `Signal()`
 - wake up a waiting thread (if any) and put it on the ready queue (may not run immediately)
 - `NotifyAll()` — historically called `Broadcast()`
 - wake up all waiting threads and put them on the ready queue
- A thread **must hold the lock** when doing these condition variable operations because
 - first, these operations may update the state
 - second, to ensure signal and wait operations are not interleaved (by two threads)

Protocol for using condition variables

- Acquire the lock to enter the critical section
- Check the condition inside the critical section
 - if condition is true: block the thread and release the lock
 - if condition is false: only release the lock

Example: the coke machine (loose syntax)

Condition variables are used with a mutex lock and in a loop (to check the condition)

```
Class CokeMachine{
    ...
    storage for cokes (buffer)
    Lock lock;
    int count = 0;
    Condition notFull, notEmpty;
}

CokeMachine::Deposit(){
    lock->acquire( ); // entering the critical section
    while(count == n)
        notFull.wait(&lock); // release lock before blocking; reacquire when waking up
    add coke to the machine;
    count++;
    notEmpty.notify();
    lock->release();
}

CokeMachine::Remove(){
    lock->acquire(); // entering the critical section
    while(count == 0)
        notEmpty.wait(&lock); // release lock before blocking; reacquire when waking up
    remove coke from the machine;
    count--;
    notFull.notify(); // always hold a lock while signalling to avoid a race condition
    lock->release();
}
```