

Devices, Drivers and PCI

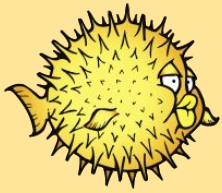
COMP3301 - 5 Week Applied Class



COMP3301 - 2025

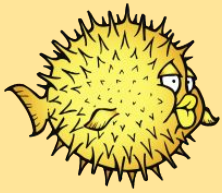


THE UNIVERSITY
OF QUEENSLAND
AUSTRALIA



Devices and Drivers

- A **device** is any **external hardware** that is attached to your computer.
 - Hard Disks, Graphics Cards, Mice, Keyboards, Monitors
 - Network Devices, Encryption Devices,
- Devices participate in the **autoconf(9)** system
 - a framework which enables devices on a system and decides what driver to use for a particular device
- Supported devices are **architecture specific**
- There are two types of devices
 - Block devices (bdev)
 - You write to them in blocks / “random access”
 - (usul. File systems, formats, hard disks swap devices)
 - Character devices (cdev)
 - You write to them in characters
 - Pretty much everything else
- Each cdev has a unique magic number among cvdevs and bdevs among bdevs




Device Special Files

- Device special files (a.k.a. device nodes) in the `/dev/` directory
- Provide a user-space interface to kernel device drivers
- Have an associated major or minor number that connects to the driver


Example:

- `/dev/ttyACM0` (cdev)
 - USB serial device “0” (ACM = Abstract Control Model)
 - Example: Arduino, USB modem
 - `/dev/sda` (bdev)
 - SCSI disk “a” (Small Computer System Interface)
-
- devices usually create these on attaching
 - `mknod(8)` can be used to make device special files manually



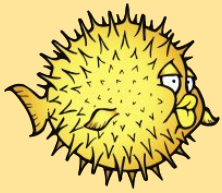
OpenBSD Basic Driver Anatomy

- `const struct cfattach`
- `struct cfdriver`
- `int match(struct device *parent, void *match, void *aux)`
 - Asks the driver “Can you drive this device?”
- `void attach(struct device *parent, struct *device self, void *aux))`
 - What do I need to set up for a device driver when a device is attached
- `int detach(struct device *parent, void *match, int flags)`
 - What do I need to clean up for a device driver when the device is detached
- Software context (`struct softc`)
 - Struct that contains context information for a particular instance of a device.
 - You decide what goes in here – anything you want to be kept by the driver



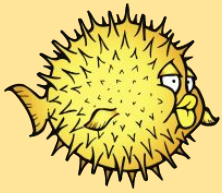
OpenBSD cdev entry points - Basics!

- `int open(dev_t dev, int oflags, int devtype, struct proc *p)`
 - a. How do I open a session to my device?
- `int close(dev_t dev, int iflags, int devtype, struct proc *p)`
 - a. How do I clean up a session to my device?
- `int read(dev_t dev, struct uio *uio, int ioflag)`
 - a. How do I define “reading” from my device
- `int write(dev_t dev, struct uio *uio, int ioflag)`
 - a. How do I define “writing to my device”
- `int ioctl(dev_t dev, u_long *cmd, caddr_t data, int fflags, struct proc *p)`
 - a. “Input/Ouput Control”
- Not all syscalls are used for every device
 - E.g. for some devices read/write might not make sense. You may purely interface with a device via `ioctls`
 - Open/Close might not make sense for a device that doesn’t need per-connection states or isolation




OpenBSD cdev entry points (Continued)

- `int kqueue(dev_t dev, struct knote *kn)`
 - Scalable event notification interface
 - Can be used to allow non-blocking behaviour for long requests
- `tty(struct tty *tp, int rw)`
 - Returns the tty struct associated with the device (for pseudo-terminals, serial ports, console devices, etc.)
- `mmap(dev_t dev, off_t offset, int flags)`
 - Defined how or if a user process can memory-map your device into their address space.
 - Might be used by processes like x11 etc.
- We will go into kqueue in more detail in a future contact but tty and mmap won't be discussed in this course.



OpenBSD bdev entry points

- `int open(dev_t dev, int oflags, int devtype, struct proc *p)`
 - How do I open a session to my device?
- `int close(dev_t dev, int iflags, int devtype, struct proc *p)`
 - How do I clean up a session to my device?
- `void *strategy(struct buf *bp)`
 - How do I access a block of my device for I/O
- `int ioctl(dev_t dev, u_long *cmd, caddr_t data, int fflags, struct proc *p)`
 - “Input/Output Control”
- `int dump(dev_t dev, u_long *cmd, caddr_t data, int fflags, struct proc *p)`
 - Special: writes device contents for kernel crash dumps.



open() and close()

- `int open(dev_t dev, int oflags, int devtype, struct proc *p)`
 - a. How do I open a session to my device?
- `int close(dev_t dev, int iflags, int devtype, struct proc *p)`

`dev_t dev`

The device being accessed (from `/dev/`) with associated Major and Minor Numbers

`int oflags /`
`int iflags`


Flags specifying access mode (read, write, non-blocking, etc.)

`int devtype`

Device type (`cdev` or `bdev`)

`struct proc *p`

Pointer to the calling process



read() and write()

- `int read(dev_t dev, struct uio *uio, int ioflag)`
 - a. How do I define “reading” from my device
- `int write(dev_t dev, struct uio *uio, int ioflag)`
 - a. How do I define “writing to my device”

`dev_t dev`


The device being accessed (from `/dev/`) with associated Major and Minor Numbers

`struct uio *uio`

Pointer to an array of io requests.


`int ioflags`

Flags specifying access mode, e.g., `IO_UNIT`, `IO_NDELAY`



Struct UIO

```
struct uio {  
    struct iovec *uio_iov; /* Pointer to array of I/O vectors */  
    int      uio_iovcnt; /* Number of elements in uio_iov */  
    off_t     uio_offset; /* Offset in the device or file */  
    ssize_t    uio_resid; /* Remaining bytes to transfer */  
    enum uio_rw uio_rw; /* UIO_READ or UIO_WRITE */  
    struct proc *uio_procp; /* Pointer to process performing the I/O */  
};
```



ioctl()

- `int ioctl(dev_t dev, u_long *cmd, caddr_t data, int fflags, struct proc *p)`


`dev_t dev` The device being accessed (from `/dev/`) with associated Major and Minor Numbers

`u_long cmd` Command code — identifies the operation requested by the user program. Each driver can define its own set of commands.

`caddr_t data` Pointer to a data structure in user space containing arguments for the command. Driver reads/writes this to get/set device state.

`int fflags` Flags passed from the open file descriptor (e.g., read/write mode)

`struct proc *p` Pointer to the calling process.



ioctl() - commands

- Each ioctl() has a unique number for it's device
- Up to 13 bits
- Cmd numbers encode
 - Length of params
 - Direction of params
 - Subsystem of device
 - Group of device

31 30 29 28 27 ... 16 15 ... 8 7 ... 0

+-----+-----+-----+-----+

| Direction | Size | Group | Number |

+-----+-----+-----+-----+

| Read/Out | len | '5' | cmd num |

+-----+-----+-----+-----+


```
#if !defined(_SYS_P5D_H)
#define _SYS_P5D_H

#include <sys/ioctl.h>
#include <sys/ioccom.h>
#include <sys/types.h>

struct p5d_status_params {
    uint    psp_is_num_waiting;
};

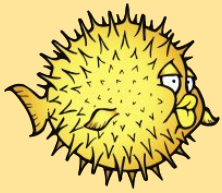
#define      P5D_IOC_STATUS      _IOR('5', 1, struct p5d_status_params)

#endif /* _SYS_P5D_H */
```



Peripheral Controller Interface (PCI)

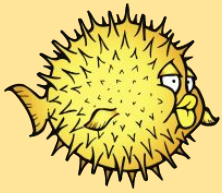
- A bus device
 - Transfers data in parallel
 - High speed access to all devices on the bus (For ye olden days)
 - 133MB/s to 533MB
- PCI Express (PCIe)
 - Even faster!
 - up to 242 GB/s
- In Comparison
 - AXI -> 1GB/s
 - USB FS -> 48MB/s
 - USB HS -> 480MB/s
 - I2c -> 400KB/s
 - SPI -> 4MB/s
- Registers in the form of “Bars”
- Can have capabilities like MSIx or DMA (more on this next time)
- `pcidump -v`



Bars and Bus Space


- The main “registers” for pci devices are found in their “bars”
- These could be settings, or small fields for entry
- They could also be triggers (e.g. DBELL regs)

Offset	Size	Name	Description
0x00	64 bits	<i>A</i>	The first number to be added (R/W)
0x08	64 bits	<i>B</i>	The second number to be added (R/W)
0x10	64 bits	<i>SUM</i>	The sum of A and B (read-only) : $SUM = A + B$




Memory Mapping and Barriers

- You can directly map memory for bars to an address in your OS
- `pci_mapreg_map`
- For linear mapping you can organise this with a struct so you only need to map once
- Most modern systems have “caching”
 - IO operations may not be resolved immediately to improve performance
 - When external sources modify these, or try to read them they may not be up to date
- “Barriers” can be used to force IO operations to occur immediately
 - A write barrier will force the CPU to write to the register instead of leaving it in it’s cache
 - A read barrier will force a CPU to discard it’s cached value and read the memory again



Practicals for this Week and Next Week

- **Prac 4 - Basic PCI Device Driver (Content from Today)**
 - Writing a basic PCI driver for the emulated p4d to add two numbers
 - Redefining add2 syscall in the pci driver to interface with userland
 - End result functionally the same as prac 2
- **Prac 5 - Writing a Pseudo Device (Content from Today)**
 - Define a “pseduo device” in the kernel as a cdev
 - gives you an understanding of how devices are defined in the kernel
 - You don’t usually have to go through all these steps if a type of device is already defined for it (e.g. pci, usb)
 - End result functionally the same as prac 3
- **Prac 6 - More Advanced PCI Device Driver (Next week’s content)**
 - Writing a PCI device driver to interface with a more advanced device
 - Using DMA to write to ring buffers
 - Sending requests to a device
 - Blocking – waiting on the device for the result
 - Non-blocking - making the request then polling the device for completions
 - Catching MSIx generated interrupts for completions and handling them
 - Using kqueuefilter() to poll for completions (handling non-blocking requests)



Coming Weeks - Device Events and DMA

- PCIe - Events and Interrupts
- DMA - Direct Memory Access
- MSIx -Message Signaled Interrupts
- kqueue() - scalable event notification interface
- Assignment 2 - Write a device driver for an emulated PCIe device