



# Interrupts

ECE 391



# Topics and Concepts

- Control and Status Registers
- RISC-V execution modes (M, S, and U)
- Interrupt controller (PLIC on RISC-V)
- Interrupt priorities and masking
- CPU interrupt delivery and trap handling



# Know how to ...

- Read, write, and modify CSRs
- Configure the PLIC to deliver interrupts from source
- Configure CPU to receive external interrupts
- Write a RISC-V interrupt handler

# Know these terms ...

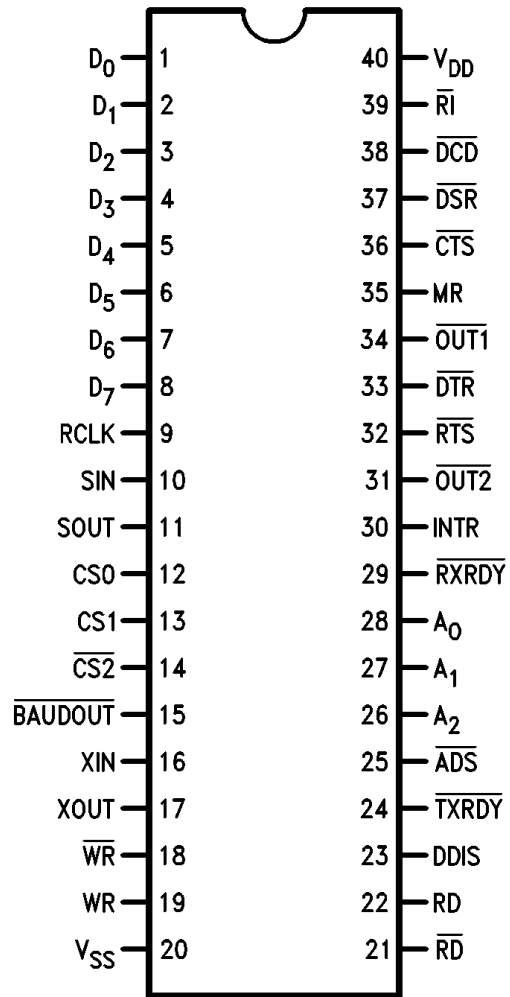
- CSRs
  - Exceptions
  - Interrupts
  - Traps
  - PLIC
  - Hart
  - Context (PLIC interrupt sink)
  - Atomic
- CSRs
    - mstatus
    - mip and mie
    - mtvec
    - mcause
    - mepc



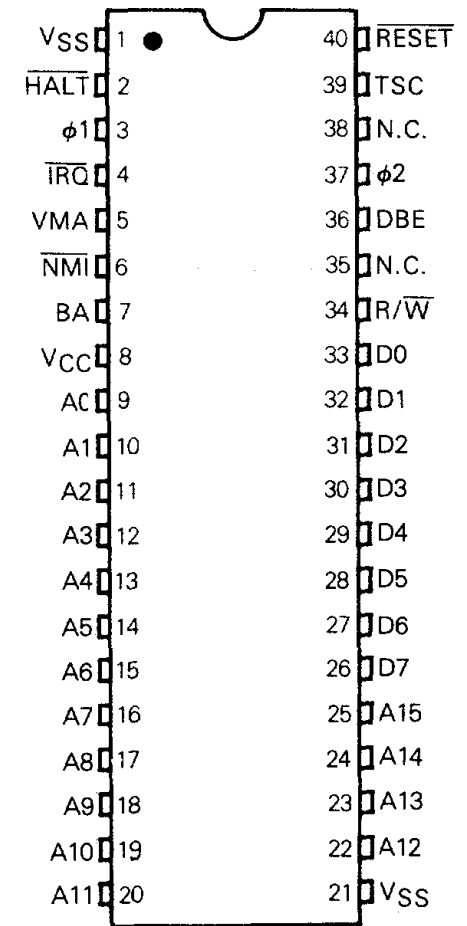
# Interrupts

- The CPU has one interrupt request pin: *what if we have multiple devices?*

# Simplest Interrupt Wiring



NS16550A

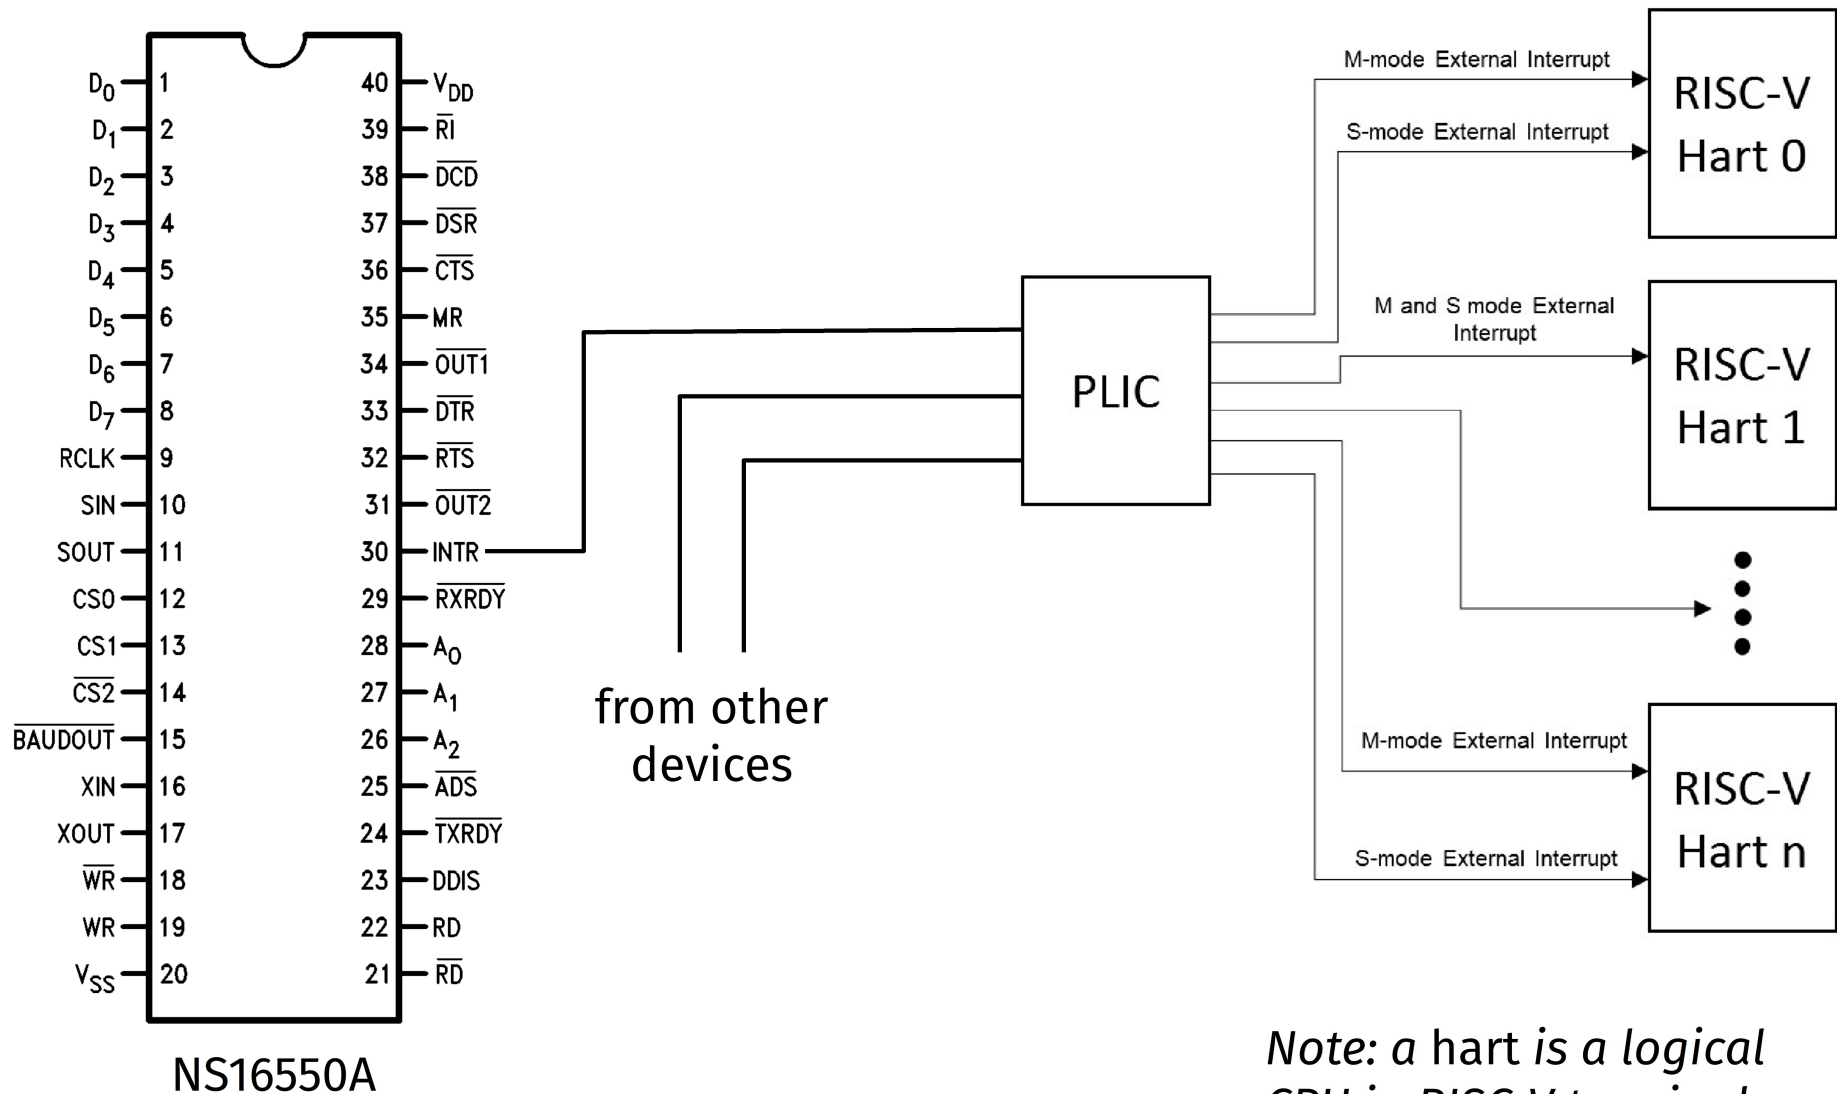


MC6800

# Interrupt Controller

- The one-gate-to-or-them-all interrupt wiring scheme works for simple systems
  - Check every device to know who raised the interrupt
  - No way to disable some interrupt sources
- Modern systems use an **interrupt controller** to give the CPU more control over external interrupts
- RISC-V systems use an interrupt controller called a *Platform Level Interrupt Controller (PLIC)*

# PLIC Wiring





# Hart to Heart

- RISC-V documentation uses the term *hart* to refer to a computing resource that behaves like a discrete CPU from the programmer's point of view
  - I will refer to this as a *logical CPU* and often just *CPU* when the context is software (rather than architecture)
- Logical CPUs may share units (such as ALUs) but appear as separate processing units
- From the programmer's point of view, these are just CPUs

# RISC-V Privilege Modes

- At any given time, a RISC-V CPU is operating in one of three modes:
  - **M**achine mode: most privileged, CPU starts in this mode
  - **S**upervisor mode: privileged mode intended for the OS
  - **U**ser mode: least privileged mode for user programs

*RISC-V privilege levels.*

Level	Encoding	Name	Abbreviation
0	00	User/Application	U
1	01	Supervisor	S
2	10	<i>Reserved</i>	
3	11	Machine	M

# RISC-V Privilege Modes

- At any given time, a RISC-V CPU is operating in one of three modes:
  - **M**achine mode: most privileged, CPU starts in this mode
  - **S**upervisor mode: privileged mode intended for the OS
  - **U**ser mode: least privileged mode for user programs
- An OS can run in M-mode or S-mode
  - In the course we will use M mode for the OS and U mode for user-space programs (later in the course)
  - So far we have been running all code in M mode

# RISC-V Control and Status Registers

- In addition to general purpose registers and PC, a RISC-V CPU has *Control and Status Registers (CSRs)*
- Special registers that indicate CPU status and control its behavior
- Accessed using six instructions
  - CSRRW, CSRRS, CSRRC (register source operand)
  - CSRRWI, CSRRSI, CSRRCI (immediate source operand)
- *Access may have side-effects*

# CSR Atomic Read and Write

$rd \leftarrow csr \leftarrow rs$

`csrrw rd, csr, rs`

$rd \leftarrow csr \leftarrow imm$

`csrrwi rd, csr, imm`

- If *rd* is not zero: Read named CSR, write the value read into general purpose register *rd*
  - If *rd* is zero, CSR is not read
- Write contents of *rs* (`csrrw`) or sign-extended immediate argument (`csrrwi`) to named CSR
- Read and write operations happen *atomically*
  - In a single indivisible action without an observable intermediate state

# CSR Atomic Read and Set Bits

$rd \leftarrow csr \leftarrow csr \mid rs$

`csrrs rd, csr, rs`

$rd \leftarrow csr \leftarrow csr \mid imm$

`csrrsi rd, csr, imm`

- Read named CSR, write the value read into general purpose register *rd*
  - If *rd* is zero, CSR is still read (but result is ignored)
- If *rs* is not zero: Bitwise-OR the CSR with the contents of *rs* (`csrrs`)
- If *imm*  $\neq$  0: Bitwise-OR the CSR with the zero-extended immediate value (`csrrsi`)
- If *rs* is zero or *imm* = 0: no CSR write takes place

# CSR Atomic Read and Clear Bits

$rd \leftarrow csr \leftarrow csr \& \sim rs$

`csrrc rd, csr, rs`

$rd \leftarrow csr \leftarrow csr \& \sim imm$

`csrrci rd, csr, imm`

- Read named CSR, write the value read into general purpose register *rd*
  - If *rd* is zero, CSR is still read (but result is ignored)
- If *rs* is not zero: Bitwise-AND the CSR with the bitwise-negated contents of *rs* (`csrrs`)
- If *imm*  $\neq 0$ : Bitwise-AND the CSR with the bitwise-negated zero-extended immediate value (`csrrsi`)
- If *rs* is zero or *imm* = 0: no CSR write takes place



# CSR Immediate Instructions

- The immediate value in the CSRRxI instructions is only 5 bits (zero-extended to 32 or 64 bits)
  - Can only modify bits 0 to 4 of the CSR



# CSR Pseudo-Instructions

- `csrw csr, rs`  $\equiv$  `csrrw zero, csr, rs` (CSR write)
- `csrwi csr, imm`  $\equiv$  `csrrwi zero, csr, imm` (CSR write)
- `csrr csr, rs`  $\equiv$  `csrrs rd, csr, zero` (CSR read)
- `csrs csr, rs`  $\equiv$  `csrrs zero, csr, rs` (CSR set bit)
- `csrsi csr, imm`  $\equiv$  `csrrs zero, csr, imm` (CSR set bit)
- `csrc csr, rs`  $\equiv$  `csrrc zero, csr, rs` (CSR clear bit)
- `csrci csr, imm`  $\equiv$  `csrrc zero, csr, imm` (CSR clear bit)

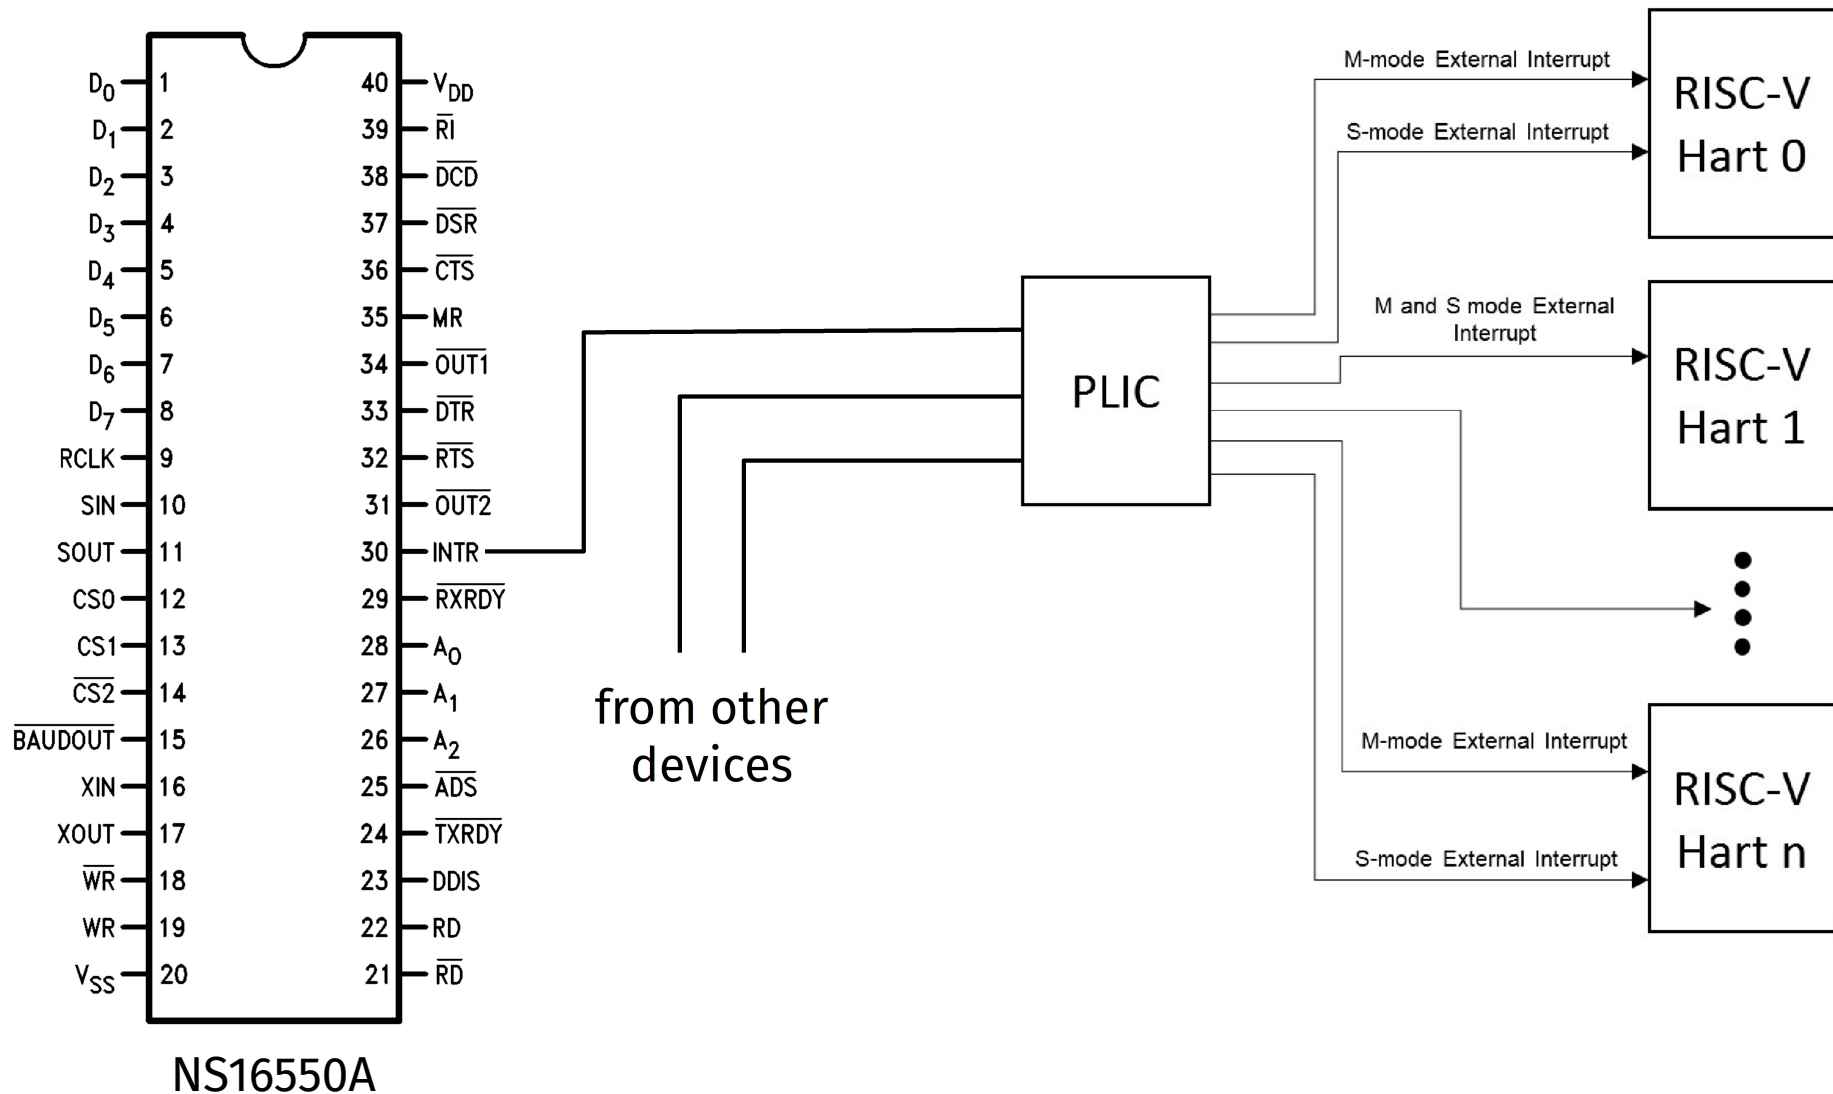
# CSR Instructions Summary

<i>Register operand</i>				
Instruction	<i>rd</i> is x0	<i>rs</i> is x0	Reads CSR	Writes CSR
CSRRW	Yes	-	No	Yes
CSRRW	No	-	Yes	Yes
CSRRS/CSRRC	-	Yes	Yes	No
CSRRS/CSRRC	-	No	Yes	Yes
<i>Immediate operand</i>				
Instruction	<i>rd</i> is x0	<i>imm</i> = 0	Reads CSR	Writes CSR
CSRRWI	Yes	-	No	Yes
CSRRWI	No	-	Yes	Yes
CSRRSI/CSRRCI	-	Yes	Yes	No
CSRRSI/CSRRCI	-	No	Yes	Yes

# RISC-V Terminology

- An **exception** is “an unusual condition occurring at run time associated with an instruction in the current RISC-V hart.”
- An **interrupt** is “an external asynchronous event that may cause a RISC-V hart to experience an unexpected transfer of control.”
- A **trap** is “the transfer of control to a trap handler caused by either an exception or an interrupt.”

# Back to Interrupts



# What Can PLIC Do for You?

- Assign a priority to each interrupt source
  - Priority 0 means interrupt is disabled (*masked*)
- Route interrupts to different CPUs
  - Can have dedicated CPU for each interrupt or spread interrupt delivery between multiple CPUs
  - In this course, we are working with a single-CPU system
- Set a priority threshold for each CPU
  - Only interrupts with priority above threshold delivered
- Indicates which interrupt source raised the interrupt

# Interrupt Sources on virt

Device	PLIC source no.
VirtIO device 0	1
VirtIO device 1	2
VirtIO device $i$	$1+i$
Virtio device 7	8
RTC	9
UART 0	10
UART 1	11
UART $j$	$10+j$
UART 7	17

# CPU Modes and Interrupts

- PLIC can deliver interrupts to M mode or S mode
- By default, all interrupts handled in M mode
- M mode can delegate interrupts to S mode
- PLICs route interrupts to *PLIC contexts*
  - Context is overloaded: we will say “PLIC context” when we mean context as defined by PLIC spec
- PLIC contexts map to M or S mode on CPUs in a system-dependent manner (not specified by spec)

# PLIC Context to CPU and Mode

- On the virt system we are using:
  - PLIC context 0 maps to M mode on hart 0
  - PLIC context 1 maps to S mode on hart 0
  - PLIC context 2 maps to M mode on hart 1
  - PLIC context 3 maps to S mode on hart 1
  - PLIC context  $2k$  maps to M mode on hart  $k$
  - PLIC context  $2k+1$  maps to S mode on hart  $k$
- In this course: *we will emulate one CPU and will not use S mode: only PLIC context 0 matters*



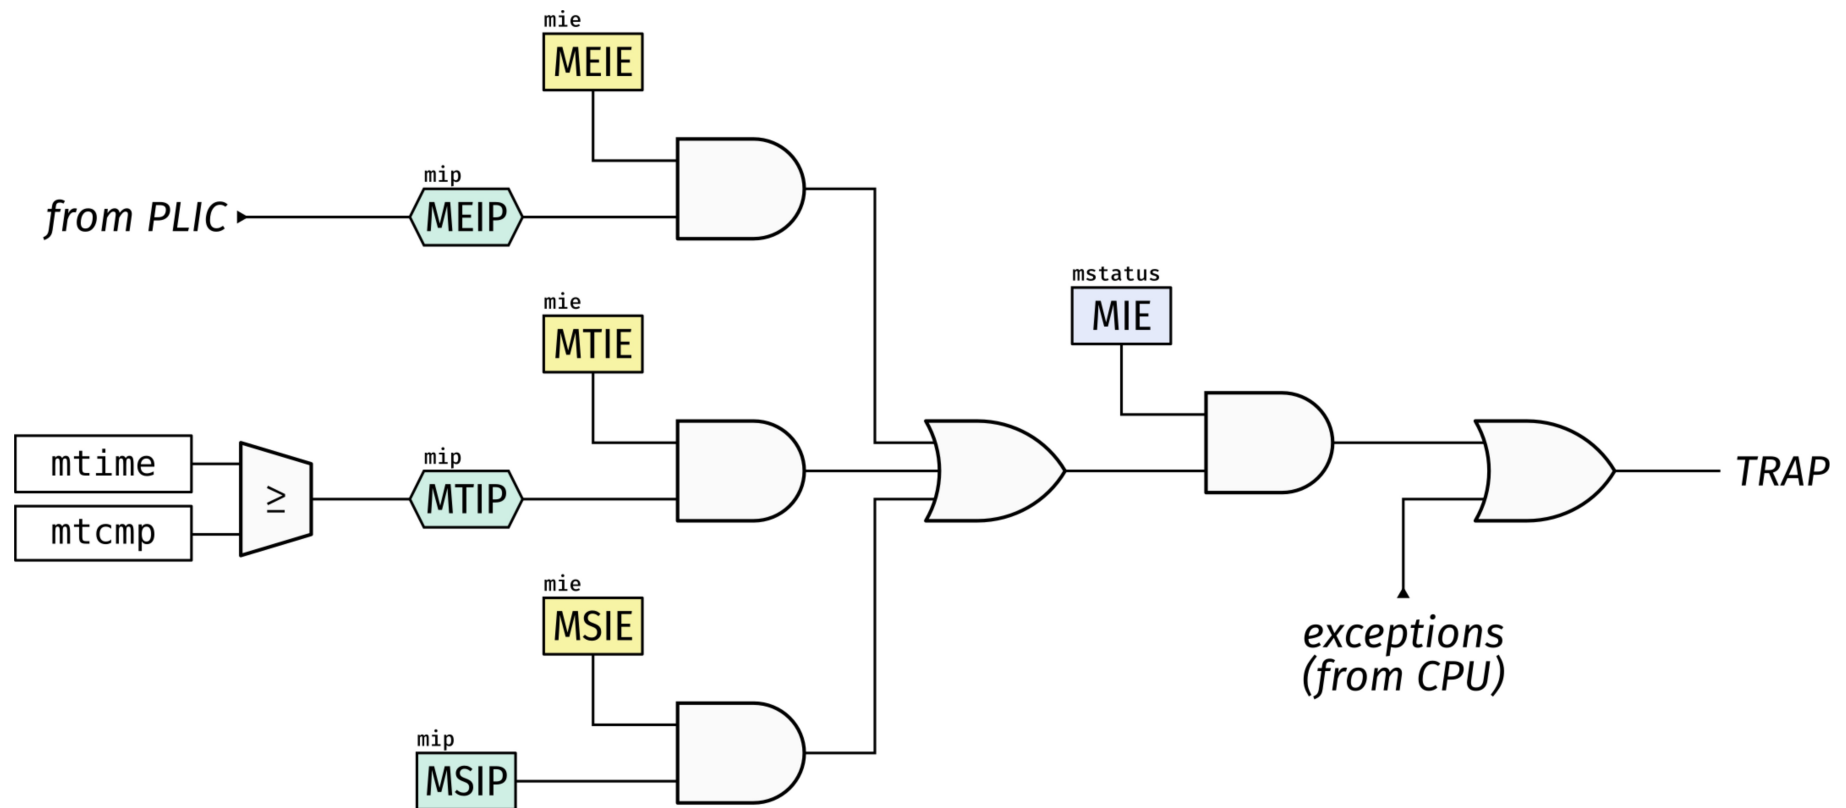
# Our PLIC API

- You will write some functions for controlling the PLIC
  - `void plic_set_src_prio (unsigned int srcid, unsigned int prio);`
  - `void plic_enable_ctx_src (unsigned int ctxid, unsigned int srcid);`
  - `void plic_disable_ctx_src (unsigned int ctxid, unsigned int srcid);`
  - `void plic_set_ctx_threshold (unsigned int ctxid, unsigned int prio);`

# The Picture So Far

- External devices are interrupt *sources* for the PLIC
  - One device could originate multiple interrupts
- When an interrupt is asserted, the PLIC decides whether to send it to a PLIC context based on:
  - Whether the source is enabled (priority non-zero)
  - The set of sources enabled for each PLIC context
  - Whether (source priority) > (PLIC context threshold)
- *We will only use context 0 in the course*

# Interrupt Signals Inside the CPU



# mie and mip

mip - Pending interrupts (status register)


15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	LCOFIP	0	MEIP	0	SEIP	0	MTIP	0	STIP	0	MSIP	0	SSIP	0	0
2	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1

mie - Enable interrupts (control register)

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	LCOFIE	0	MEIE	0	SEIE	0	MTIE	0	STIE	0	MSIE	0	SSIE	0	0
2	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1

- Bits mip.MEIP and mie.MEIE are the interrupt-pending and interrupt-enable bits for machine-level external interrupts. MEIP is read-only in mip, and is set and cleared by a platform-specific interrupt controller.
- Bits mip.MTIP and mie.MTIE are the interrupt-pending and interrupt-enable bits for machine timer interrupts. MTIP is read-only in mip, and is cleared by writing to the memory-mapped machine-mode timer compare register.

# Interrupt Signals Inside the CPU



# mstatus

mstatus - Machine status register

22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
	TSR	TW	TVM	MXR	SUM	MPRV	XS[1:0]	FS[1:0]	MPP[1:0]	VS[1:0]	SPP	MPIE	UBE	SPIE	WPRI	MIE	WPRI	SIE	WPRI			
	1	1	1	1	1	1	2	2	2	2	1	1	1	1	1	1	1	1	1	1	1	

- MIE: global interrupt enable for M mode
- MIE can be set and cleared using CSRRSI and CSRRCI instructions


# mtvec



- MODE:

Value	Name	Description
0	Direct	All traps set <b>pc</b> to BASE.
1	Vectored	Asynchronous interrupts set <b>pc</b> to BASE+4×cause.
≥2	---	<i>Reserved</i>

# mcause



*Machine cause register (mcause) values after trap.*

Interrupt	Exception Code	Description
1	0	<i>Reserved</i>
1	1	Supervisor software interrupt
1	2	<i>Reserved</i>
1	3	Machine software interrupt
1	4	<i>Reserved</i>
1	5	Supervisor timer interrupt
1	6	<i>Reserved</i>
1	7	Machine timer interrupt
1	8	<i>Reserved</i>
1	9	Supervisor external interrupt
1	10	<i>Reserved</i>
1	11	Machine external interrupt
1	12	<i>Reserved</i>
1	13	Counter-overflow interrupt
1	14-15	<i>Reserved</i>
1	≥16	<i>Designated for platform use</i>



# Enabling Interrupts

1. Set trap handler address (`mtvec`)
2. Enable interrupts on device (e.g. set `DRIE=1` on NS16550A)
3. Set source priority on PLIC to non-zero value
4. Enable source for context 0 (M mode on hart 0)
5. Set context 0 interrupt priority threshold
  - Only sources whose priority is greater than context's priority threshold will be sent to the context
6. Enable M mode external interrupts on CPU (`mie.MEIE = 1`)
7. Enable M mode interrupts (`mstatus.MIE = 1`)

# Accessing CSRs from C

- Create functions to access CSRs using C inline assembly:

```
int64_t csrr_mcause(void) {  
    int64_t val;  
  
    asm inline ("csrr %0, mcause" : "=r" (val));  
    return val;  
  
}  
  
void csrw_mtvec(void (*handler)(void)) { // MODE 0  
    asm inline ("csrwr mtvec, %0" :: "r" (handler));  
  
}
```

<https://gcc.gnu.org/onlinedocs/gcc/extensions-to-the-c-language-family/how-to-use-inline-assembly-language-in-c-code.html>

# Enabling Interrupts

1. Set trap handler address (mtvec)
  - `csrw_mtvec(trap_handler);`
2. Enable interrupts on device (e.g. set DRIE=1 on NS16550A)
  - *E.g.* `UART0.ier = DRIE;`
3. Set source priority on PLIC to non-zero value
  - `plic_set_src_prio(srcid, prio);`
4. Enable source for context 0 (M mode on hart 0)
  - `plic_enable_ctx_src(ctxid, srcid);`

# Enabling Interrupts

5. Set context 0 interrupt priority threshold
  - `plic_set_ctx_threshold(ctxid, level);`
6. Enable M mode external interrupts on CPU (`mie.MEIE = 1`)
  - `csrs_mie(MEIE);`
7. Enable M mode interrupts (`mstatus.MIE = 1`)
  - `csrs_mstatus(MIE);`

# Interrupt Handler

- CPU jumps to trap handler when PLIC interrupts
  - Can happen at any time when `mie.MEIE = 1` and `mstatus.MIE = 1`
- All general purpose registers contain whatever values they had when CPU was interrupted
- After servicing the interrupt, CPU should resume execution at the point where the interrupted
- Need to restore register values to their original state
  - Interrupted program assumes register values will not mysteriously change between instructions!
- This is tricky!

# Execution State

- We need to save the current program execution state on entry so we can resume execution later:
  - PC and general purpose registers
- Need stack for handler (unless handler very simple):
  - Use separate stack for trap handlers
  - Use current kernel stack
- Where to save PC and registers?
  - To special area
  - To handler stack (which may be the main kernel stack)

# Trap Entry

- Interrupts and exceptions both handled as *traps*
- Traps are handled in M mode
  - M mode can delegate to S mode, but we won't use this
- On trap entry, the CPU does the following:
  - Places current PC into mepc CSR
  - Sets mstatus.MPP to previous privilege mode
  - Sets mstatus.MPIE to mstatus.MIE
  - Sets mstatus.MIE to 0 (interrupts disabled on entry)
  - Sets PC to mtvec.BASE (if mode 0) or mtvec.BASE + 4×*cause*

# Trap Entry

Value	Name	Description
0	Direct	All traps set <b>pc</b> to BASE.
1	Vectored	Asynchronous interrupts set <b>pc</b> to BASE+4×cause.
≥2	---	<i>Reserved</i>

- On trap entry, the CPU does the following:
  - Places current PC into mepc CSR
  - Sets mstatus.MPP to previous privilege mode
  - Sets mstatus.MPIE to mstatus.MIE
  - Sets mstatus.MIE to 0 (interrupts disabled on entry)
  - Sets PC to mtvec.BASE (if mode 0) or mtvec.BASE + 4×cause



# Trap Exit

- The mret instruction is used to exit a trap:
  - Sets mstatus.mie to mstatus.MPP; sets MPIE to 1
  - Changes privilege mode to mstatus.MPP; sets MPP to 0
  - Set PC to mepc

*RISC-V privilege levels.*

Level	Encoding	Name	Abbreviation
0	00	User/Application	U
1	01	Supervisor	S
2	10	<i>Reserved</i>	
3	11	Machine	M

# Trap Handler Overview

1. Save all registers to the stack
  - We will re-use the kernel stack
2. Save mepc and mstatus
  - Optional, will be useful later
3. Call exception handler or interrupt handler
  - These can be normal C functions
4. Restore mepc and mstatus
5. Restore registers from stack (incl. stack pointer)
6. Execute mret instruction (this re-enables interrupts)

# trap.s

```
# trap.s
#

.text
.global _trap_entry
.type _trap_entry, @function

.align 4 # Trap entry must be 4-byte aligned for mtvec CSR

_trap_entry:

# Save current state to the stack.
#
# struct trap_frame {
#     uint64_t x[32]; // x[0] unused
#     uint64_t mstatus;
#     uint64_t mepc;
# };
#
```

# trap.s

`_trap_entry:`

```
# Save current state to the stack.
#
#   struct trap_frame {
#       uint64_t x[32]; // x[0] unused
#       uint64_t mstatus;
#       uint64_t mepc;
#   };
#
```

```
nop                # x0 is zero
sd      x1, 1*8(sp) # x1 is ra
nop                # x2 is sp
sd      x3, 3*8(sp) # x3 is gp
sd      x4, 4*8(sp) # x4 is tp
sd      x5, 5*8(sp) # x4 is t0
```

```
# and so on
```

# trap.s

```
sd      x30, 30*8(sp) # x30 is t5
sd      x31, 31*8(sp) # x31 is t6
```

```
# Save mstatus and mepc
```

```
csrr    t0, mstatus
sd      t0, 32*8(sp)
csrr    t0, mepc
sd      t0, 33*8(sp)
```

```
# ALL HANDLERS RETURN HERE
```

```
# ... restore registers
```

```
1:      csrr    a0, mcause          # put mcause in a0
        mv      a1, sp             # put pointer to trap_frame in a1
        bgez    a0, fault_handler
        j       intr_handler
```