



Data Structures and Algorithms

Bloom Filters

CS 225
G Carl Evans

April 25, 2025



Department of Computer Science

Memory-Constrained Data Structures

What method would you use to build a search index on a collection of objects *in a memory-constrained environment*?

Constrained by Big Data (Large N)




Image: <https://doi.org/10.1038/nature03597>

Sky Survey Projects	Data Volume
DPOSS (The Palomar Digital Sky Survey)	3 TB
2MASS (The Two Micron All-Sky Survey)	10 TB
GBT (Green Bank Telescope)	20 PB
GALEX (The Galaxy Evolution Explorer)	30 TB
SDSS (The Sloan Digital Sky Survey)	40 TB
SkyMapper Southern Sky Survey	500 TB
PanSTARRS (The Panoramic Survey Telescope and Rapid Response System)	~ 40 PB expected
LSST (The Large Synoptic Survey Telescope)	~ 200 PB expected
SKA (The Square Kilometer Array)	~ 4.6 EB expected

Table: <http://doi.org/10.5334/dsj-2015-011>

Estimated total volume of one array: 4.6 EB

Bloom Filter: Insertion

S = { 16, 8, 4, 13, 29, 11, 22 }


$h(k) = k \% 7$

0	0
1	0
2	0
3	0
4	0
5	0
6	0

Bloom Filter: Insertion

An item is inserted into a bloom filter by hashing and then setting the hash-valued bit to 1

If the bit was already one, it stays 1



Bloom Filter: Search

S = { 16, 8, 4, 13, 29, 11, 22 }

$h(k) = k \% 7$

0	0
1	1
2	1
3	0
4	1
5	0
6	1

`_find(16)`

`_find(20)`

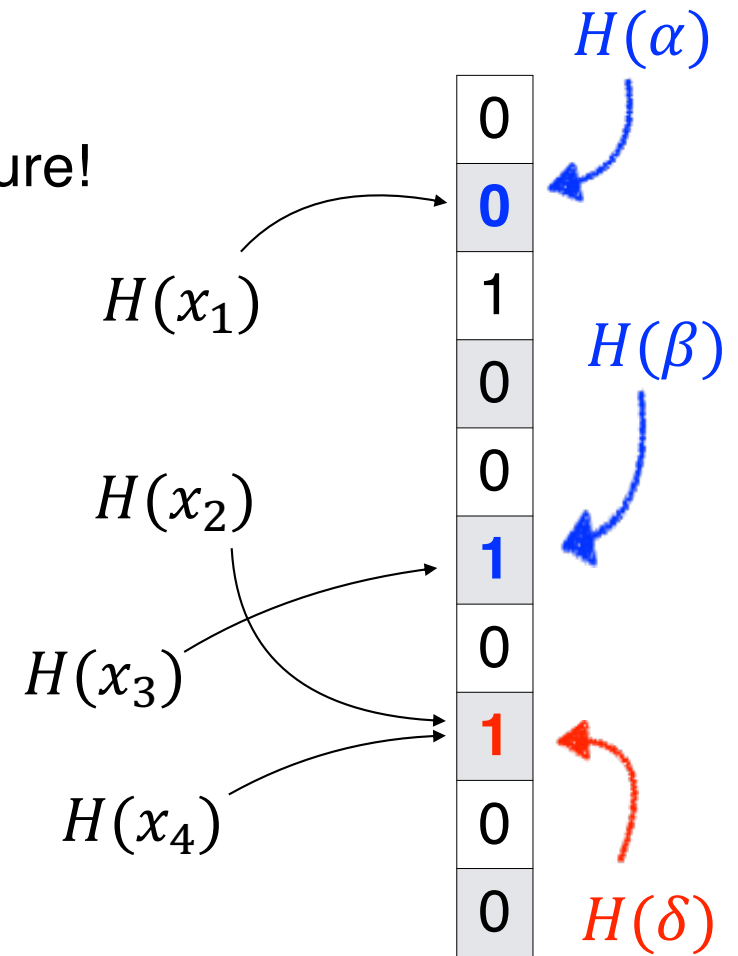
`_find(3)`

Bloom Filter: Search

The bloom filter is a *probabilistic* data structure!



If the value in the BF is 0:

If the value in the BF is 1:





Probabilistic Accuracy: Malicious Websites

Imagine we have a detection oracle that identifies if a site is malicious



"Not malicious"



"Malicious"



Probabilistic Accuracy: Malicious Websites

Imagine we have a detection oracle that identifies if a site is malicious

True Positive:

False Positive:

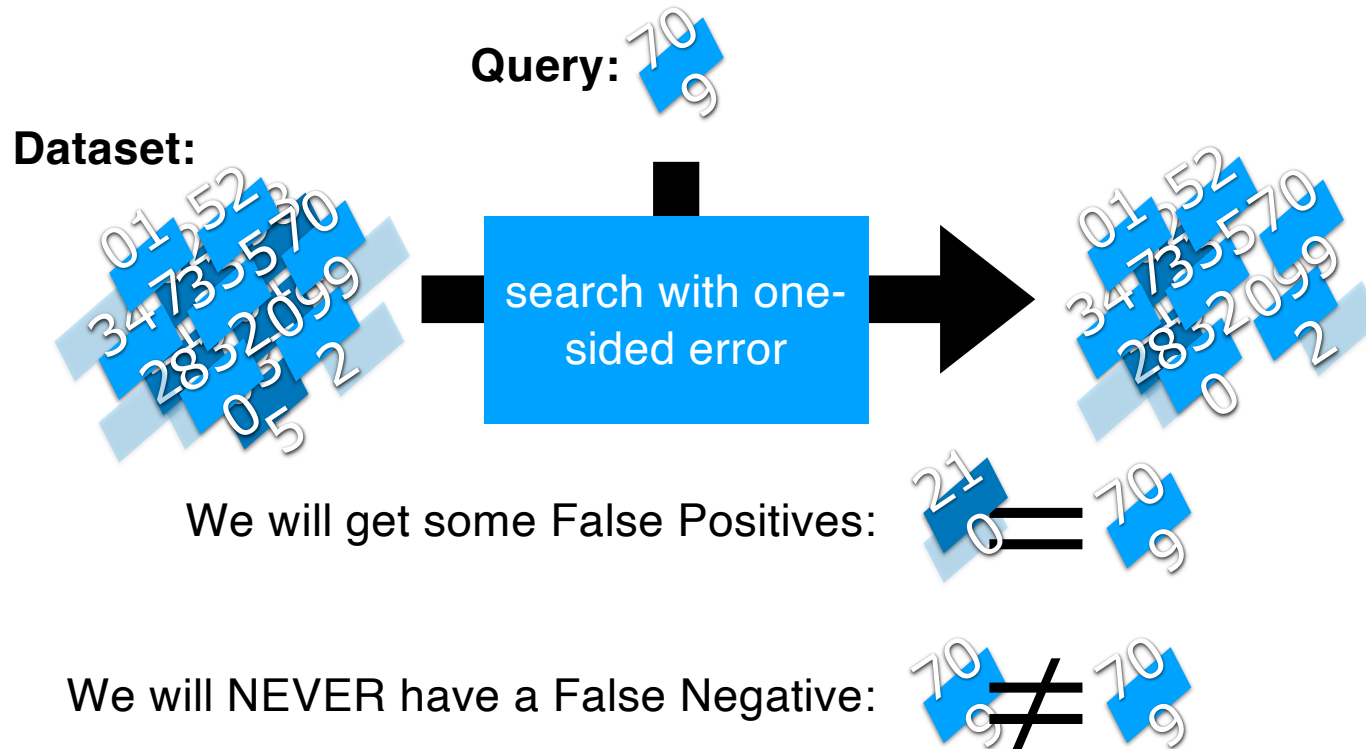
False Negative:

True Negative:


Imagine we have a **bloom filter** that **stores malicious sites...**

	Bit Value = 1	Bit Value = 0
Item Inserted	<div><div>$H(z)$</div><div><div>0</div><div>1</div><div>0</div><div>0</div><div>1</div></div><div>'Yes'</div><div>True Positive</div></div>	<div><div>$H(z)$</div><div><div>0</div><div>0</div><div>0</div><div>0</div><div>1</div></div><div>'No'</div><div>False Negative</div></div>
Item NOT inserted	<div><div>$H(z)$</div><div><div>0</div><div>1</div><div>0</div><div>0</div><div>1</div></div><div>'Yes'</div><div>False Positive</div></div>	<div><div>$H(z)$</div><div><div>0</div><div>0</div><div>0</div><div>0</div><div>1</div></div><div>'No'</div><div>True Negative</div></div>

Probabilistic Accuracy: One-sided error



Probabilistic Accuracy: One-sided error



Bloom Filter: Repeated Trials

Use many hashes/filters; add each item to each filter

h_1

0
1
0
1
0
0
0
1
0
1
1
0
1
0
1
1
0
1
0
1



Bloom Filter: Repeated Trials

Use many hashes/filters; add each item to each filter

h_1

0
1
0
1
0
0
0
1
0
1
1
0
1
0
1
0
1
1
0
1
0
1

h_2

0
0
0
1
0
0
0
1
0
1
0
0
1
1
0
0
1
1
0
0
1



Bloom Filter: Repeated Trials


Use many hashes/filters; add each item to each filter

h_1	0	h_2	0	h_3	0
	1		0		1
	0		0		1
	1		1		1
	0		0		0
	0		0		0
	0		0		1
	1		1		1
	0		0		0
	1		1		1
	1		0		1
	0		0		0
	1		1		1
	0		1		0
	1		1		1
	1		0		1
	0		0		0
	1		1		1
	0		0		0
	1		0		1



Bloom Filter: Repeated Trials

Use many hashes/filters; add each item to each filter




Bloom Filter: Repeated Trials


0	0	0	...	0
1	0	1		1
0	0	1		1
1	1	1		1
0	0	0		1
0	0	0		1
0	0	1		0
1	1	1		0
0	0	0		1
1	1	1		0
1	0	1		1
0	0	0		1
1	1	1		0
0	1	0		1
1	1	1		1
0	0	0		1
1	0	1		1

$$h_{\{1,2,3,\dots,k\}}(y)$$

Bloom Filter: Repeated Trials



Bloom Filter: Repeated Trials





Bloom Filter: Repeated Trials


Using repeated trials, even a very bad filter can still have a very low FPR!

If we have k bloom filter, each with a FPR p , what is the likelihood that ***all*** filters return the value '1' for an item we didn't insert?



Bloom Filter: Repeated Trials

Rather than use a new filter for each hash, one filter can use k hashes



Bloom Filter: Repeated Trials

Rather than use a new filter for each hash, one filter can use k hashes

0	0	$h_1(x) = x \% 10$	$h_2(x) = 2x \% 10$	$h_3(x) = (5+3x) \% 10$
1	0			
2	1	<code>_find(1)</code>		
3	1			
4	1			
5	0			
6	1	<code>_find(16)</code>		
7	1			
8	1			
9	1			

Bloom Filter

A probabilistic data structure storing a set of values $H = \{h_1, h_2, \dots, h_k\}$

Built from a bit vector of length m and k hash functions

Insert / Find runs in: _____


Delete is not possible (yet)!

0
0
1
0
0
1
0
1
0
0

Bloom Filter: Error Rate

Given bit vector of size m and k SUHA hash function

What is our expected FPR after n objects are inserted?




Bloom Filter: Error Rate

Given bit vector of size m and 1 SUHA hash function

What's the probability a specific bucket is 1 after one object is inserted?


Same probability given k SUHA hash function?



Bloom Filter: Error Rate

Given bit vector of size m and k SUHA hash function

Probability a specific bucket is 0 after one object is inserted?




After n objects are inserted?

Bloom Filter: Error Rate

Given bit vector of size m and k SUHA hash function

What's the probability a specific bucket is **1** after n objects are inserted?



Bloom Filter: Error Rate


Given bit vector of size m and k SUHA hash function

What is our expected FPR after n objects are inserted?

The probability my bit is 1 after n objects inserted

$$\left(1 - \left(1 - \frac{1}{m}\right)^{nk}\right)^k$$

The number of [assumed independent] trials



Bloom Filter: Error Rate

Vector of size m , k SUHA hash function, and n objects

To minimize the FPR, do we prefer...

(A) large k

(B) small k

$$\left(1 - \left(1 - \frac{1}{m}\right)^{nk}\right)^k$$

