

CS915/435 Advanced Computer Security

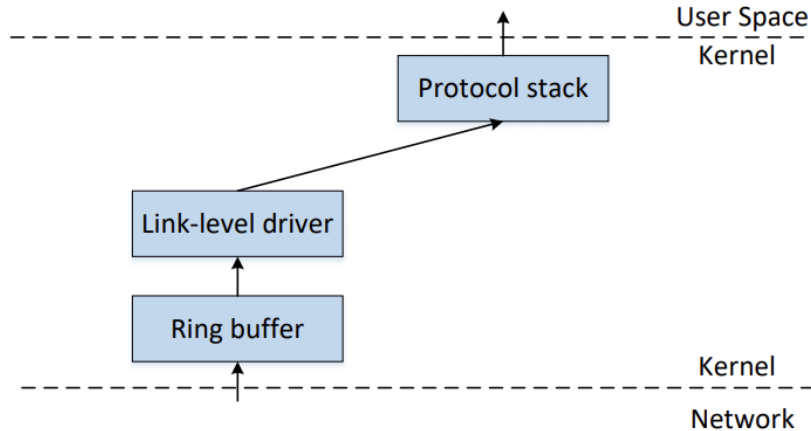
- Network Security (I)

Packet Sniffing and Spoofing

Outline

- Packet Sniffing and Spoofing
 - How packets are received
 - Packet sniffing
 - Packet spoofing
 - Packet sniffing then spoofing
- TCP attacks
- Firewall
- DNS attacks

How Packets Are Received



- Machines are connected to networks through Network Interface Cards
- Each NIC has a MAC address
- Every NIC on the network will hear all the frames on the wire
- If a match is found (destination address in the header), the frame is copied into a buffer and dispatched to user-space programs.

Promiscuous Mode

- The frames that are not destined to a given NIC are discarded
- When operating in promiscuous mode, NIC passes every frame received from the network to the kernel
- If a sniffer program is registered with the kernel, it will be able to see all the packets.
- Normally, enabling Promiscuous Mode requires elevated privilege (root)

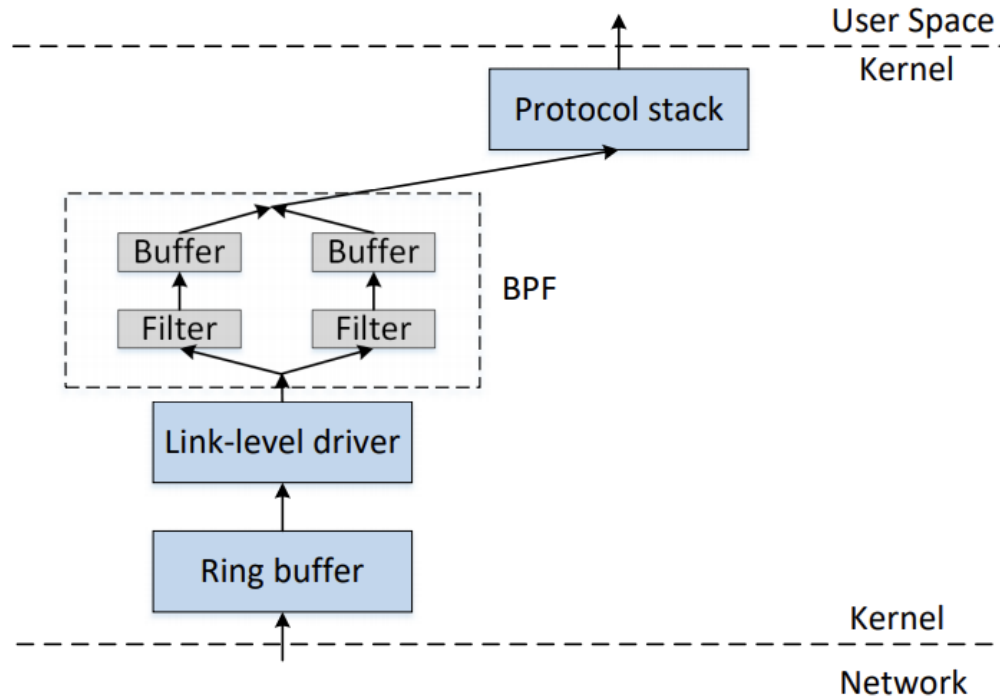
Wireless network

- Promiscuous mode also available in wireless network
- In Wi-Fi, it is called **Monitor Mode**
- However, due to interference, Wi-Fi card can't copy everything (too much data in the air).
- Wi-Fi cards work with different channels, or slices of the spectrum
- If a Wi-Fi card is put in the monitor mode, they will capture 802.11 frames on the channel that they are listening to.
- Hence, you may miss getting information in Monitor Mode if you're on a different channel.

How do NIC cards copy sniffed data to buffers?

- Normally sniffers are only interested in certain types of packets, e.g., TCP, DNS queries
- The system can give all the captured packets to the sniffer program, who will discard unwanted packets.
- But this is rather inefficient.
- It is better to filter unwanted packets as early as possible.

BSD Packet Filter (BPF)



- BPF allows a user-program to attach a filter to the socket, which tells the kernel to discard unwanted packets.
- For example, a filter allows only packets on port 22.
- It is possible to have a combination of filters.

Packet Sniffing

- Packet sniffing describes the process of capturing live data as they flow across a network.
- Packet sniffing tools are also called packet sniffers.
- Let's first see how computers receive packets.

Receiving Packets Using Socket (UDP server)

Create the socket

Provide information
about server

Receive packets

```
// Step ①
int sock = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP);

// Step ②
memset((char *) &server, 0, sizeof(server));
server.sin_family = AF_INET;
server.sin_addr.s_addr = htonl(INADDR_ANY);
server.sin_port = htons(9090);

if (bind(sock, (struct sockaddr *) &server, sizeof(server)) < 0)
    error("ERROR on binding");

// Step ③
while (1) {
    bzero(buf, 1500);
    recvfrom(sock, buf, 1500-1, 0,
            (struct sockaddr *) &client, &clientlen);
    printf("%s\n", buf);
}
```

Receiving Packets Using Raw Socket (sniffer)

Creating a raw socket

Capture all types of packets

```
// Create the raw socket
int sock = socket(AF_PACKET, SOCK_RAW, htons(ETH_P_ALL)); ①

// Turn on the promiscuous mode.
mr.mr_type = PACKET_MR_PROMISC; ②
setsockopt(sock, SOL_PACKET, PACKET_ADD_MEMBERSHIP, &mr, ③
           sizeof(mr));

// Getting captured packets
while (1) {
    int data_size=recvfrom(sock, buffer, PACKET_LEN, 0, ④
                          &saddr, (socklen_t*)sizeof(saddr));
    if(data_size) printf("Got one packet\n");
}
```

Enable the
promiscuous
mode

Wait for packets

Limitation of the Approach

- This program is not portable across different operating systems.
- Setting filters is not easy (not included in the code examples)
- The program does not explore any optimisation to improve performance.
- The PCAP (packet capture) library was thus created.
 - It still uses raw sockets internally, but its API is standard across all platforms. OS specifics are hidden by PCAP's implementation.
 - Allows programmers to specify filtering rules using human readable Boolean expressions.

Packet sniffing using PCAP API

```
char filter_exp[] = "ip proto icmp";
```

Filter

```
// Step 1: Open live pcap session on NIC with name eth3
handle = pcap_open_live("eth3", BUFSIZ, 1, 1000, errbuf); ①

// Step 2: Compile filter_exp into BPF psuedo-code
pcap_compile(handle, &fp, filter_exp, 0, net); ②
pcap_setfilter(handle, &fp); ③

// Step 3: Capture packets
pcap_loop(handle, -1, got_packet, NULL); ④
```

Initialise a raw socket, set the network device into promiscuous mode.

Invoke this function for every captured packet

```
void got_packet(u_char *args, const struct pcap_pkthdr *header,
               const u_char *packet)
{
    printf("Got a packet\n");
}
```

A simple sniffer (sniff.c)

https://github.com/kevin-w-du/BookCode/blob/master/Sniffing_Spoofing/C_sniff/sniff.c

```
$ gcc -o sniff sniff.c -lpcap
```

```
$ sudo ./sniff
```

Note: root privilege required

```
1  #include <pcap.h>
2  #include <stdio.h>
3
4  void got_packet(u_char *args, const struct pcap_pkthdr *header,
5                  const u_char *packet)
6  {
7      printf("Got a packet\n");
8  }
9
10 int main()
11 {
12     pcap_t *handle;
13     char errbuf[PCAP_ERRBUF_SIZE];
14     struct bpf_program fp;
15     char filter_exp[] = "ip proto icmp";
16     bpf_u_int32 net;
17
18     // Step 1: Open live pcap session on NIC with name enp0s3
19     handle = pcap_open_live("enp0s3", BUFSIZ, 1, 1000, errbuf);
20
21     // Step 2: Compile filter_exp into BPF psuedo-code
22     pcap_compile(handle, &fp, filter_exp, 0, net);
23     pcap_setfilter(handle, &fp);
24
25     // Step 3: Capture packets
26     pcap_loop(handle, -1, got_packet, NULL);
27
28     pcap_close(handle);    //Close the handle
29     return 0;
30 }
31
```

Packet Spoofing

- When some critical information in the packet is forged, we refer to it as packet spoofing.
- Many network attacks rely on packet spoofing.
- Let's see how to send packets without spoofing.

Sending Packets Without Spoofing

```
void main()
{
    struct sockaddr_in dest_info;
    char *data = "UDP message\n";

    // Step 1: Create a network socket
    int sock = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP);

    // Step 2: Provide information about destination.
    memset((char *) &dest_info, 0, sizeof(dest_info));
    dest_info.sin_family = AF_INET;
    dest_info.sin_addr.s_addr = inet_addr("10.0.2.5");
    dest_info.sin_port = htons(9090);

    // Step 3: Send out the packet.
    sendto(sock, data, strlen(data), 0,
           (struct sockaddr *)&dest_info, sizeof(dest_info));
    close(sock);
}
```

- The first step creates a socket
- The second step provides the destination information
- In the final step, sendto() is used to send out the UDP packet with the provided payload

Spoofing Packets Using Raw Sockets

- However, using the typical socket doesn't give us much control over the header fields.
- The header fields are set by the OS, e.g., source IP, packet length etc
- We can use a special type of socket called **raw socket**
- With raw socket, we construct the entire packet in a buffer including the IP header before sending it out via the socket

An example: spoofing ICMP packet

Fill in the ICMP Header


```
char buffer[1500];

memset(buffer, 0, 1500);


/*****
    Step 1: Fill in the ICMP header.
    *****/
struct icmpheader *icmp = (struct icmpheader *)
    (buffer + sizeof(struct ipheader));
icmp->icmp_type = 8; //ICMP Type: 8 is request, 0 is reply.

// Calculate the checksum for integrity
icmp->icmp_chksum = 0;
icmp->icmp_chksum = in_cksum((unsigned short *)icmp,
    sizeof(struct icmpheader));
```

Find the starting point
of the ICMP header,
and typecast it to the
ICMP structure



Fill in the ICMP header
fields



Spoofing Packets: Constructing the Packet

Fill in the IP Header

```
/* *****  
Step 2: Fill in the IP header.  
***** */  
struct ipheader *ip = (struct ipheader *) buffer;  
ip->iph_ver = 4;  
ip->iph_ihl = 5;  
ip->iph_ttl = 20;  
ip->iph_sourceip.s_addr = inet_addr("1.2.3.4");  
ip->iph_destip.s_addr = inet_addr("10.0.2.5");  
ip->iph_protocol = IPPROTO_ICMP;  
ip->iph_len = htons(sizeof(struct ipheader) +  
                    sizeof(struct icmphheader));
```

Typecast the buffer to
the IP structure

Fill in the IP header
fields

More on this Host to Network function when we discuss Endianness

Finally, send out the packet

```
send_raw_ip_packet (ip);
```

Spoofing UDP Packets

Constructing UDP packets is similar, except that we need to include the payload data.

Sniffing and Then Spoofing

- In many situations, we need to capture packets first, and then spoof a response based on the captured packets.
- Procedure (using UDP as example)
 - Use PCAP API to capture the packets of interests
 - Make a copy from the captured packet
 - Replace the UDP data field with a new message and swap the source and destination fields
 - Send out the spoofed reply
- Instead of C, we can use Scapy in Python to do the same

Sniffing/Spoofing UDP Packet

```
void spoof_reply(struct ipheader* ip)
{
    const char buffer[1500];
    int ip_header_len = ip->iph_ihl * 4;
    struct udphheader* udp = (struct udphheader *) ((u_char *)ip +
                                                    ip_header_len);

    if (ntohs(udp->udp_dport) != 9999) {
        // Only spoof UDP packet with destination port 9999
        return;
    }

    // Step 1: Make a copy from the original packet
    memset((char*)buffer, 0, 1500);
    memcpy((char*)buffer, ip, ntohs(ip->iph_len));
    struct ipheader * newip = (struct ipheader *) buffer;
    struct udphheader * newudp = (struct udphheader *) (buffer +
                                                         ip_header_len);
    char *data = (char *)newudp + sizeof(struct udphheader);

    // Step 2: Construct the UDP payload, keep track of payload size
    const char *msg = "This is a spoofed reply!\n";
    int data_len = strlen(msg);
    strncpy (data, msg, data_len);
```

Sniffing/Spoofing UDP Packet (Continued)

```
// Step 3: Construct the UDP Header
newudp->udp_sport = udp->udp_dport;
newudp->udp_dport = udp->udp_sport;
newudp->udp_ulen = htons(sizeof(struct udphheader) + data_len);
newudp->udp_sum = 0;

// Step 4: Construct the IP header (no change for other fields)
newip->iph_sourceip = ip->iph_destip;
newip->iph_destip = ip->iph_sourceip;
newip->iph_ttl = 50; // Rest the TTL field
newip->iph_len = htons(sizeof(struct ipheader) +
                        sizeof(struct udphheader) + data_len);

// Step 5: Send out the spoofed IP packet
send_raw_ip_packet(newip);
}
```

Packet Sniffing Using Scapy

```
#!/usr/bin/python3
from scapy.all import *

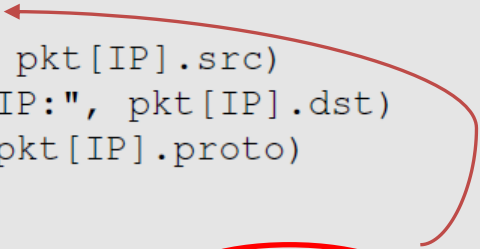
print("SNIFFING PACKETS.....")

def print_pkt(pkt):
    print("Source IP:", pkt[IP].src)
    print("Destination IP:", pkt[IP].dst)
    print("Protocol:", pkt[IP].proto)
    print("\n")

pkt = sniff(filter='icmp', prn=print_pkt)
```

①

②



A red arrow originates from the `print_pkt` argument in the `sniff` function call (line 14) and points to the `def print_pkt(pkt):` function definition (line 7). The `sniff` function call and the `'icmp'` filter string are also circled in red.

Spoofing ICMP & UDP Using Scapy

```
#!/usr/bin/python3
from scapy.all import *

print("SENDING SPOOFED ICMP PACKET.....")
ip = IP(src="1.2.3.4", dst="93.184.216.34") ①
icmp = ICMP() ②
pkt = ip/icmp ③
pkt.show()
send(pkt, verbose=0) ④
```

```
#!/usr/bin/python3
from scapy.all import *

print("SENDING SPOOFED UDP PACKET.....")
ip = IP(src="1.2.3.4", dst="10.0.2.69") # IP Layer
udp = UDP(sport=8888, dport=9090) # UDP Layer
data = "Hello UDP!\n" # Payload
pkt = ip/udp/data # Construct the complete packet
pkt.show()
send(pkt, verbose=0)
```


Sniffing and Then Spoofing Using Scapy

```
#!/usr/bin/python3
from scapy.all import *

def spoof_pkt(pkt):
    if ICMP in pkt and pkt[ICMP].type == 8:
        print("Original Packet.....")
        print("Source IP : ", pkt[IP].src)
        print("Destination IP :", pkt[IP].dst)

        ip = IP(src=pkt[IP].dst, dst=pkt[IP].src, ihl=pkt[IP].ihl)
        icmp = ICMP(type=0, id=pkt[ICMP].id, seq=pkt[ICMP].seq)
        data = pkt[Raw].load
        newpkt = ip/icmp/data

        print("Spoofed Packet.....")
        print("Source IP : ", newpkt[IP].src)
        print("Destination IP :", newpkt[IP].dst)
        send(newpkt, verbose=0)


pkt = sniff(filter='icmp and src host 10.0.2.69', prn=spoof_pkt)
```

Packet Spoofing: Scapy v.s C

- Python + Scapy
 - Pros: constructing packets is very simple
 - Cons: much slower than C code
- C Program (using raw socket)
 - Pros: much faster
 - Cons: constructing packets is complicated
- Hybrid Approach
 - Using Scapy to construct packets
 - Using C to slightly modify packets and then send packets

Endianness

- **Endianness**: a term that refers to the order in which a given multi-byte data item is stored in memory.
 - **Little Endian**: put the small end in memory first
 - **Big Endian**: put the big end in memory first
- Atmel AVR32, IBM z/Architecture mainframes use Big-Endian; x86 uses Little-Endian.



Endianness In Network Communication

- Computers with different byte orders will “misunderstand” each other.
 - Solution: agree upon a common order for communication
 - This is called “**network order**”, which is the same as **big endian** order
- All computers need to convert data between “host order” and “network order” .

Macro	Description
<code>htons()</code>	Convert unsigned short integer from host order to network order.
<code>htonl()</code>	Convert unsigned integer from host order to network order.
<code>ntohs()</code>	Convert unsigned short integer from network order to host order.
<code>ntohl()</code>	Convert unsigned integer from network order to host order.

Summary

- Packet sniffing
 - Using raw socket
 - Using PCAP APIs
- Packet spoofing using raw socket
- Sniffing and the spoofing
 - Using C, or Scapy, or a hybrid
- Endianness