

Topic 5: SDN Controller


Presented by
Dong YUAN

School of Electrical and Computer
Engineering

dong.yuan@sydney.edu.au



THE UNIVERSITY OF
SYDNEY




Contents

- Control plane
- Representative controllers

A few questions

- What should happen when h1 wants to visit h2 but no corresponding rules in sw1?

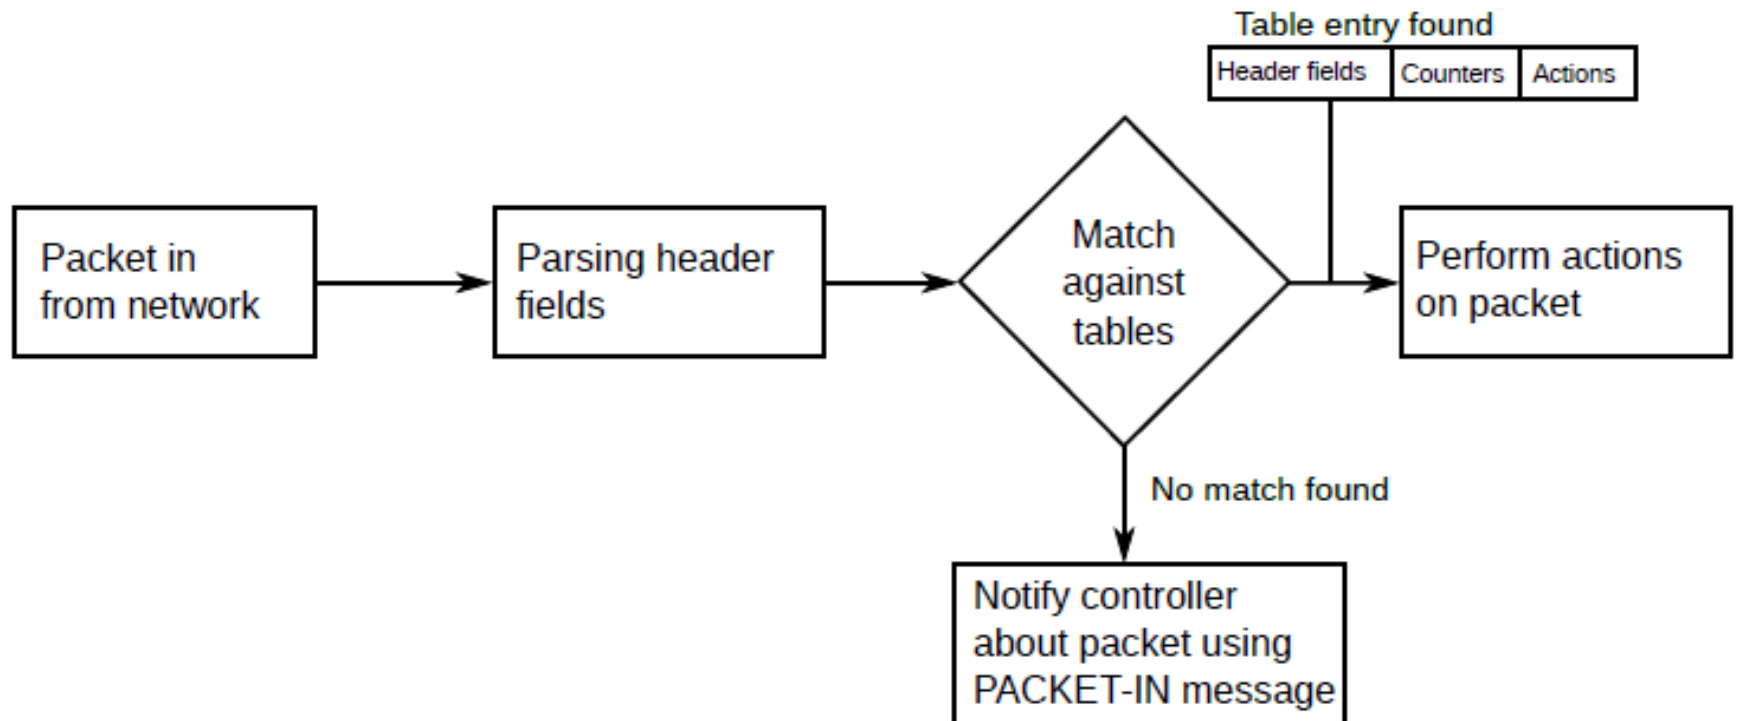


- How does a new switch connect to the Controller in an SDN?
 - Physically and Logically
- How do data communicate among controllers?

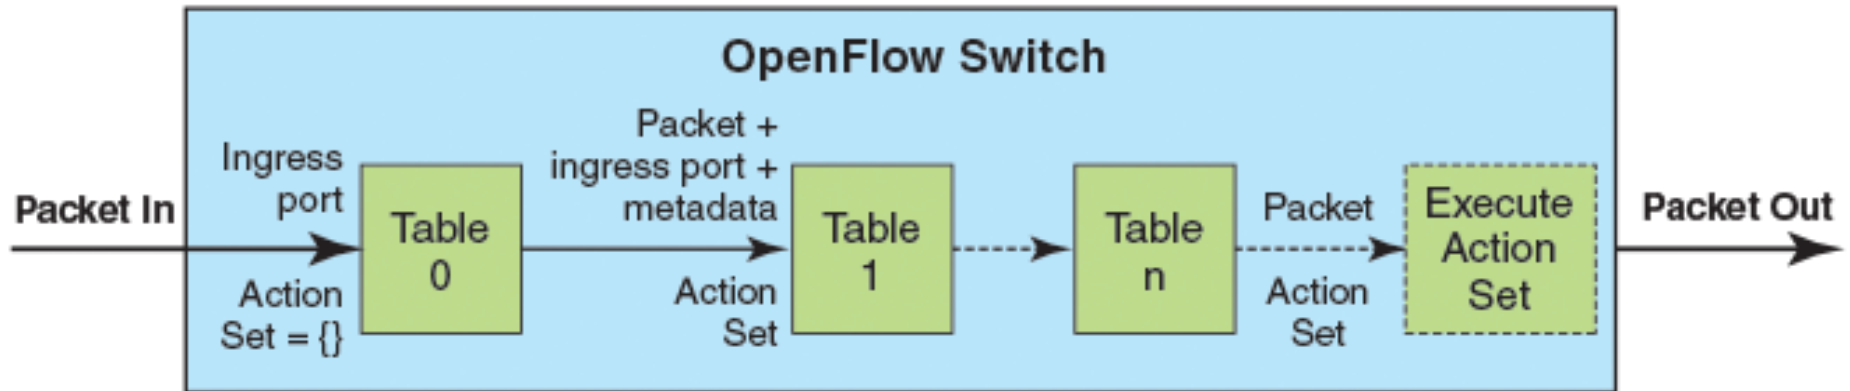
Controller

- In openflow, controller communicates with switch over a secure channel
 - Openflow protocol defines message format
 - Purpose of control channel: update flow table
 - Logic is executed at controller
 - Communication with external controllers
- Flow Table: Packet switching
 - All packets compare to flow table for match
 - Packet header fields matched against one of the N tables
 - Actions depend on match being found
 - Forward, Drop, modify, enqueue
 - If no match, traffic is sent to controller


Basic OpenFlow Matching



Advanced Matching



{a} Packets are matched against multiple tables in the pipeline



① Find highest - priority matching flow entry



② Apply instructions:

- i. Modify packet & update match fields (apply actions instruction)
- ii. Update action set (clear actions and/or write actions instructions)
- iii. Update metadata

③ Send match data and action set to next table

{b} Per-table packet processing

The control plane



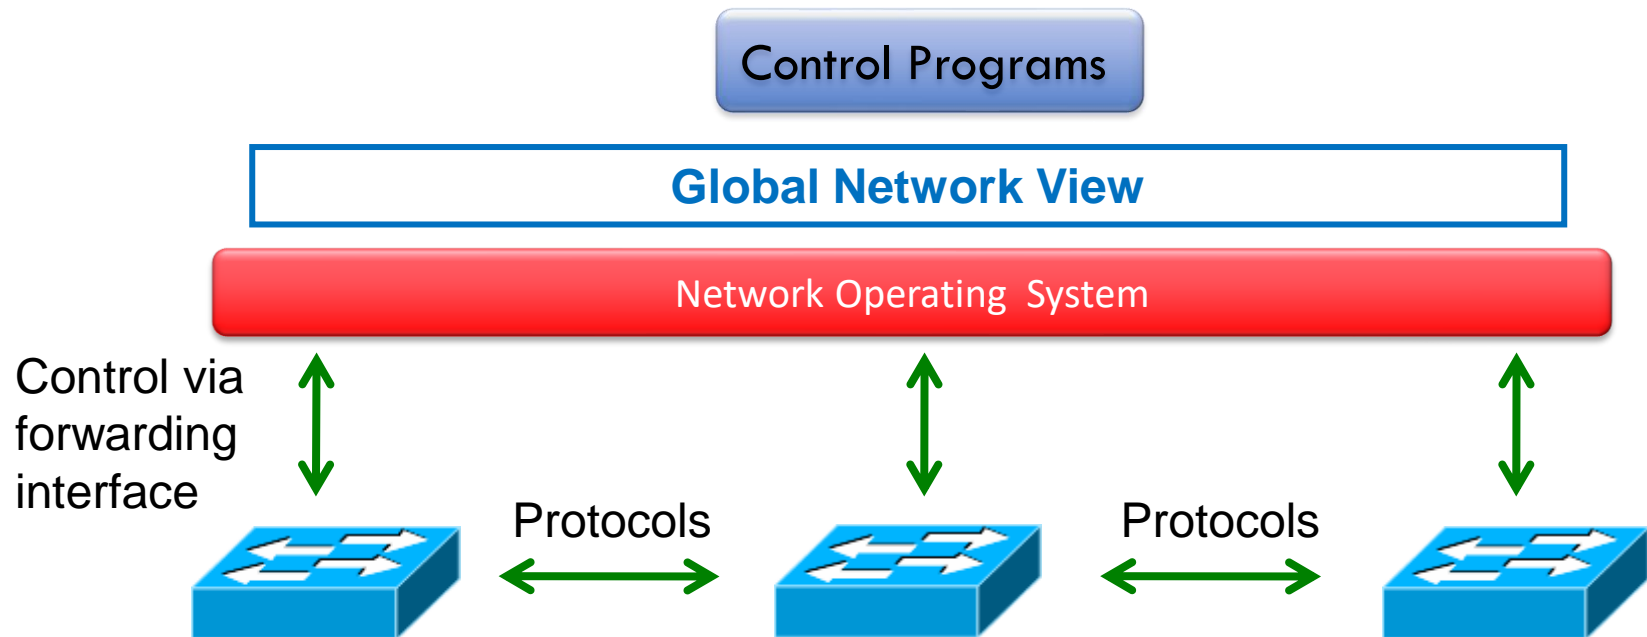
- *Separation of data plane and control plane*

Deriving the Control Plane

- What does it take to control a network?
 - Learn the network state (e.g., topologies, etc.)
 - Decide how to configure it (routing, isolation, traffic management, etc.)
 - Push configuration to the network element
- What processes are re-usable?
 - Build the network topology
 - Push configuration to network elements

Idea: An OS for Networks

Software-Defined Networking (SDN)



Network OS

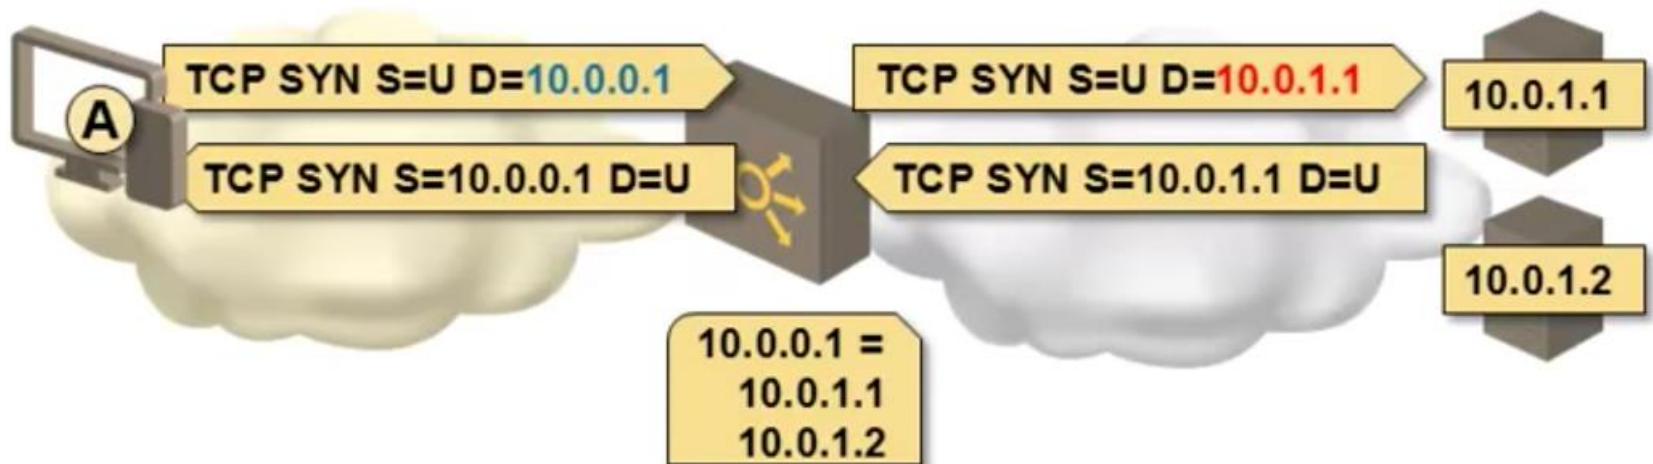
- Maintain an up-to-date view of the network state (e.g., topology, etc.)
 - Next page
- Configure network elements
 - A.k.a “south-bound” interface or API
- Provide a graph abstraction to the applications on the top
 - A.k.a “north-bound” interface or API

Network View (graph abstraction)

- In the network view:
 - Switch-level topology
 - Location of host, middlebox and other network elements
 - Location of users
 - Namespace: bindings between names and addresses
- Not in the network view:
 - Some states of network, e.g., traffic.
- Network Information Base (NIB)
 - Graph and abstraction


Example: Scale-out load balancers

- Map incoming sessions to a single virtual IP address into sessions distributed across a server farm
- Challenge
 - NAT translation record for every active session
 - Large amount of states




Example: Scale-out load balancers

- Stateful services are hard to scale out
 - Forward and reverse traffic flows might not match
 - Traffic distribution might change with adding or removing of devices



OpenFlow-assisted scale-out load balancer

- OpenFlow switches use coarse-grained flows, based on source/dest IP address ranges
- Load Balancing appliances perform fine-grained load balancing
- Load Balancing controller modifies OpenFlow flows and traffic distribution based on node availability



Challenges in building an NOS

- Scalability
- Reliability
- Good performance (fast, etc.)
- Generality and Simplicity of the “north-bound” API.
 - Easily to be used by applications

Controller's Diversification

- Different SDN controllers
 - NOX/POX
 - ONIX
 - Ryu
 - Beacon
 - Floodlight
 - Pyretic, Frenetic, Procera, etc.
- The Most Popular Two Controllers: ONOS and OpenDayLight

How to choose a controller

- Programming language
 - Can affect performance
- Focus
 - Southbound API
 - Northbound API
 - Support for Openstack/Kubernetes
 - Research or production
- Learning curve
- User base and community support

NOX

- First-generation OpenFlow controller
 - Open source, stable, widely used
 - Fast, well maintained
 - C++, Openflow 1.0
 - Some fork versions support higher OpenFlow version
- Programming model
 - Controller registers for events
 - Programmer writes event handlers
 - This is a very common SDN programming model

When to use NOX

- You know C++
- You are willing to use low-level facilities and semantics of OpenFlow
- You need good performance

POX

- NOX in Python
 - Also only support OpenFlow 1.0
- Advantages:
 - Widely used, well maintained and supported.
 - Relatively easy to read and write code
- Disadvantages:
 - Performance
- When to use POX
 - You know python and do not care about performance
 - Rapid prototyping and experimentation
 - For research, demonstration and learning concept

ONIX

- A distributed controller
- Closed source
- Production quality (used by google)
- Distributed Design
 - Scalability
 - Robustness

Ryu (pronounced "ree-yooh")

- Open source Python controller
 - Support OpenFlow 1.0-1.4
 - Works with openstack/kubernetes
- Aims to be an “Operating System” for SDN
 - Provide well-defined North-Bound APIs.
- Advantages
 - Openstack integration
 - Support newest OpenFlow
 - Easy to use, and often used in research for prototyping
- Disadvantages
 - Performance

Beacon

- From Stanford
- First (?) Java based controller
 - Fast
 - Cross platforms

Floodlight

- Open source Java controller
 - Support OpenFlow 1.0
 - For from Beacon
 - Maintained by Big Switch Networks
- Advantage
 - Good documentation
 - Integration with REST API, OpenStack
 - Production-level performance
- Disadvantage
 - Deep learning curve

Today's Two Popular controllers

- ONOS and OpenDayLight (ODL)
- ONOS (2014)
 - From Open Networking Foundation
 - Previously ON.LAB funded by Stanford and Berkeley
- ODL (2013)
 - From Linux Foundation
- Both ONOS and ODL are written in Java and designed for modular use with a customizable infrastructure
- Both support OpenStack
- Every ONOS partner is also an ODL member

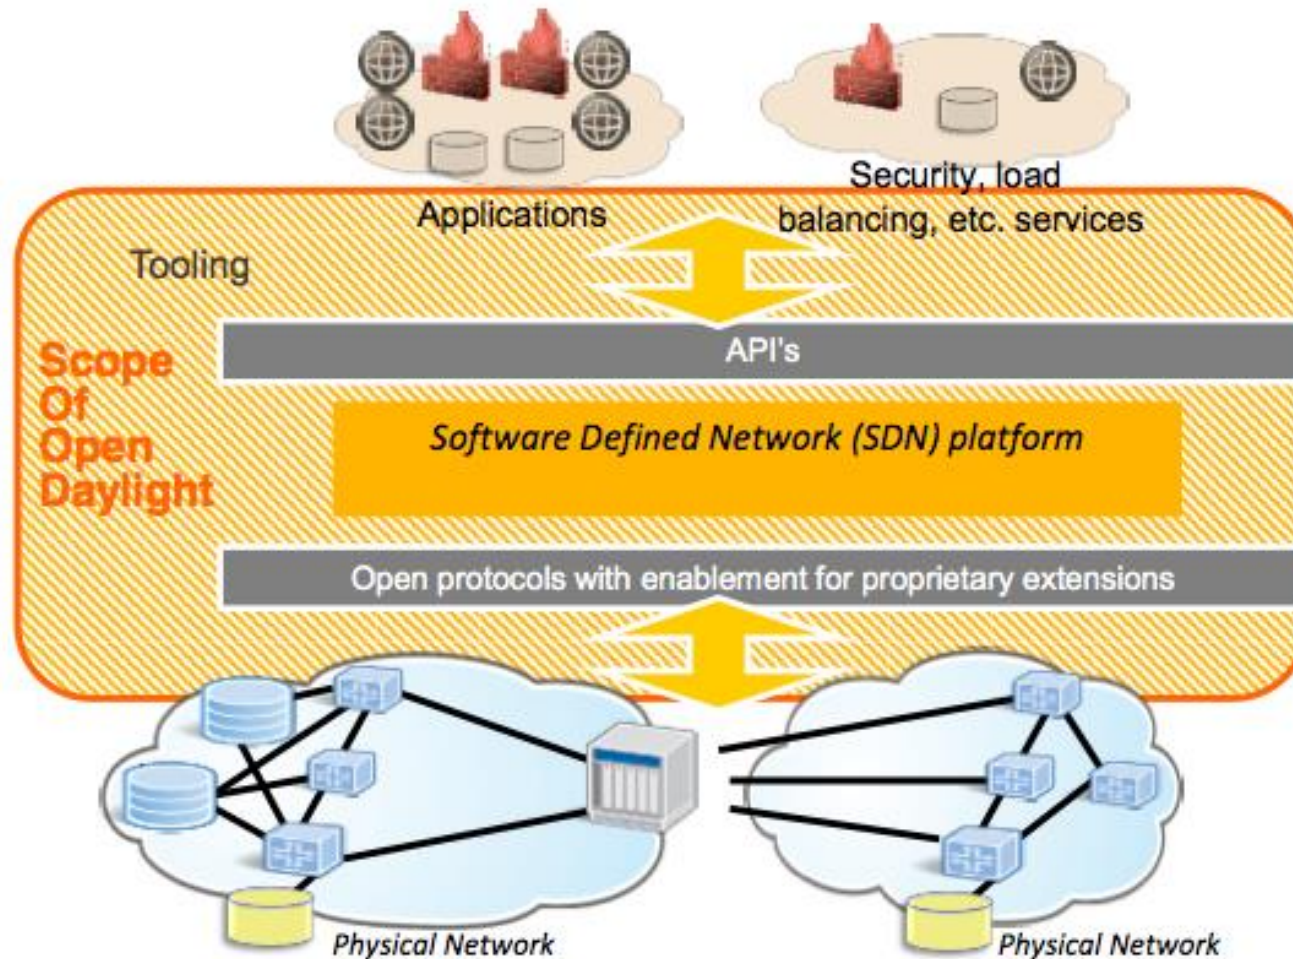
Differences

- ONOS vs. ODL
 - Carrier-grade networks vs. Cloud provider
 - Pure SDN vs. Legacy
 - Academic initiated vs. Corporate initiated

OpenDayLight Controller

OpenDaylight Scope

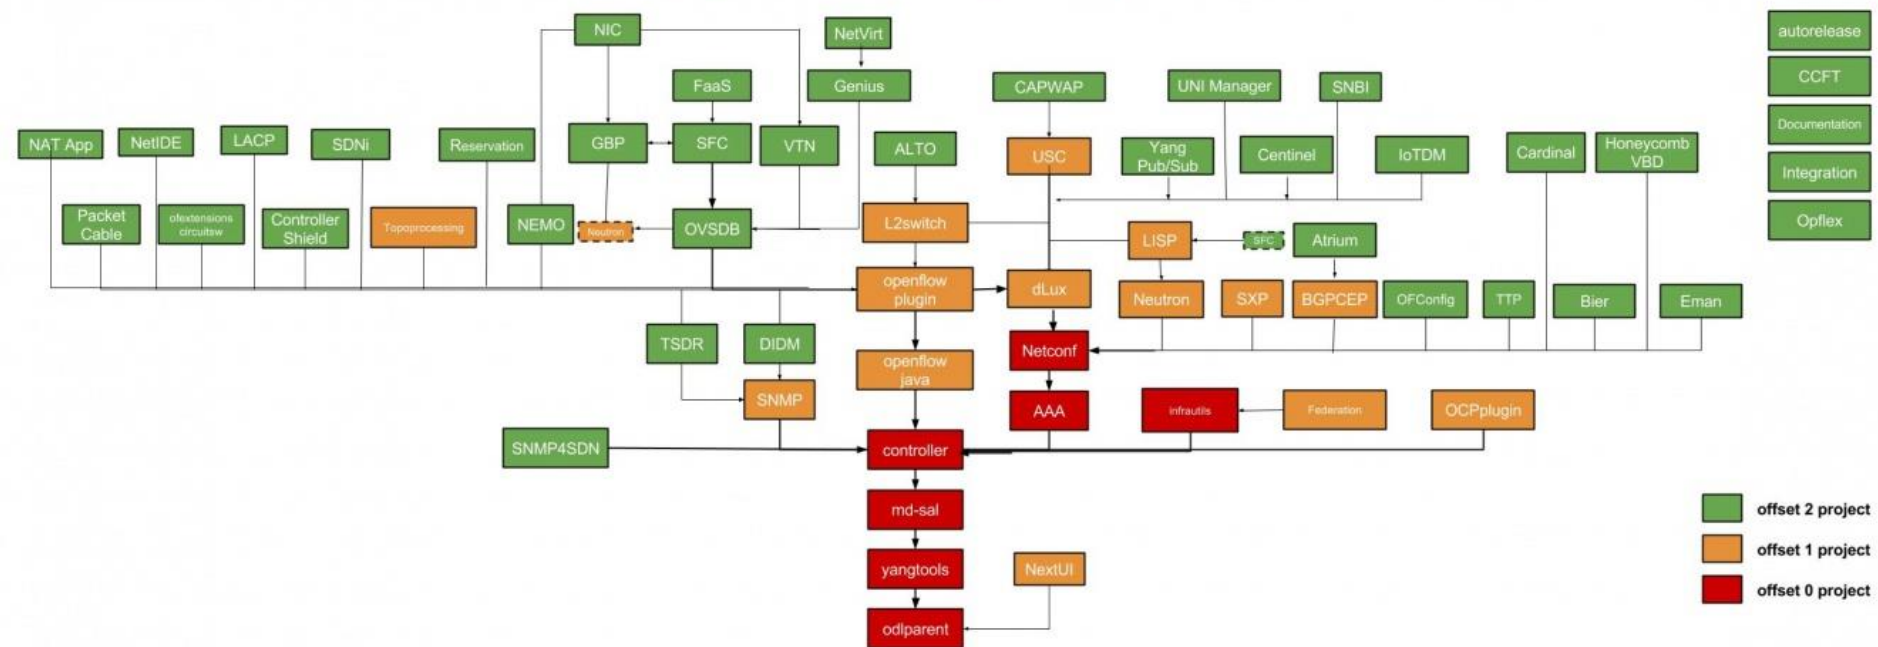
– OpenDaylight Scope



OpenDaylight Projects

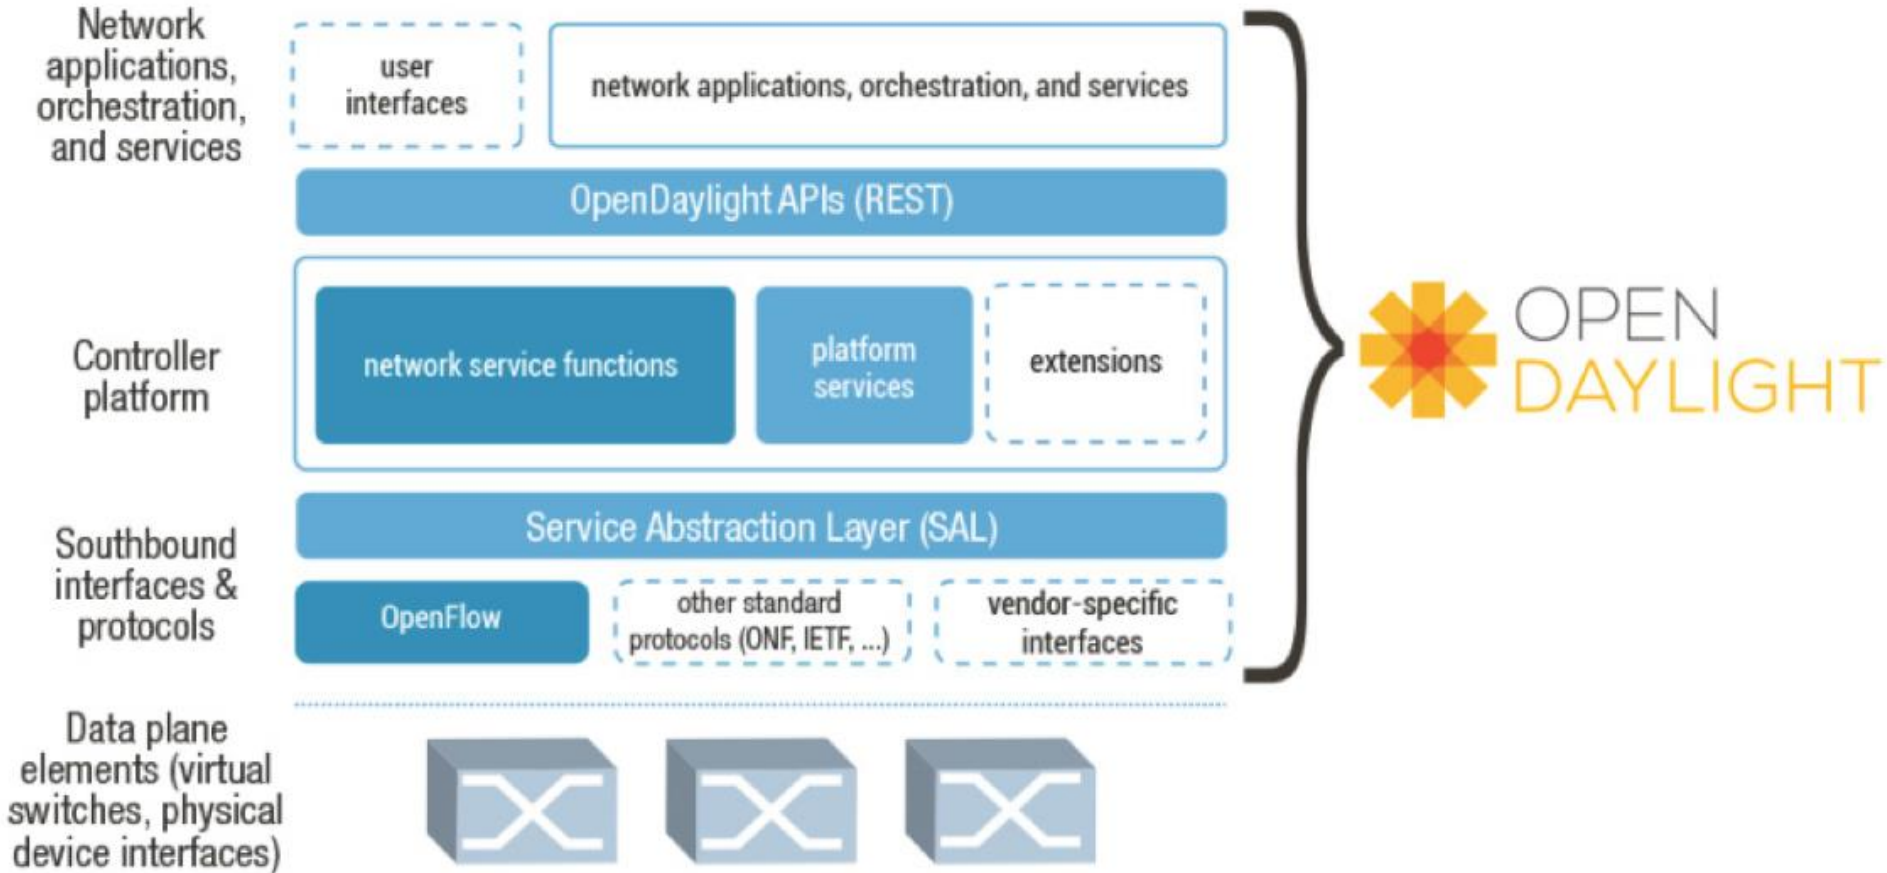
- OpenDaylight Projects
 - Classification: kernel projects, protocol projects, app projects, service projects, support projects, etc.
 - More projects indicate the OpenDaylight community is active

OpenDaylight Carbon Project Dependencies



OpenDaylight Framework

– Project Framework




Key Technologies in ODL

- **OSGI** - This framework is the back-end of OpenDaylight as it allows dynamically loading of bundles and packages JAR files, and binding bundles together for exchanging information.
- **Karaf** - Application container built on top of OSGI, which simplifies operational aspects of packaging and installing applications. Runtime environment for well designed code modules, aka, “modulith runtime”.
- **YANG** - a data modeling language used to model configuration and state data manipulated by the applications, remote procedure calls, and notifications.

Java OSGi Framework

- OSGi (Open Service Gateway Initiative) is a Java framework for developing and deploying modular software programs and libraries.
- Applications or components, coming in the form of bundles for deployment, can be remotely installed, started, stopped, updated, and uninstalled without requiring a reboot.



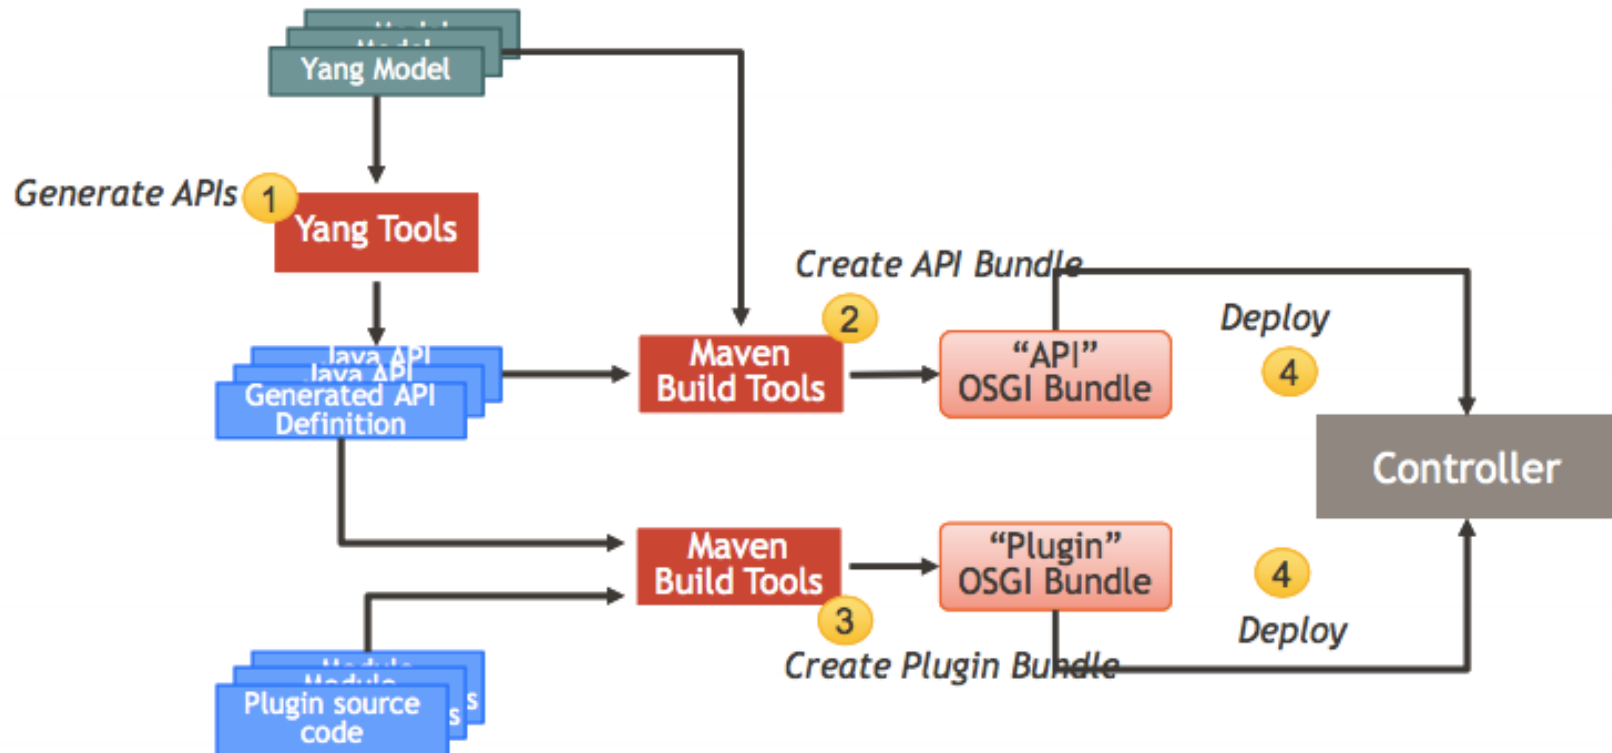
OpenDaylight Controller

– Features

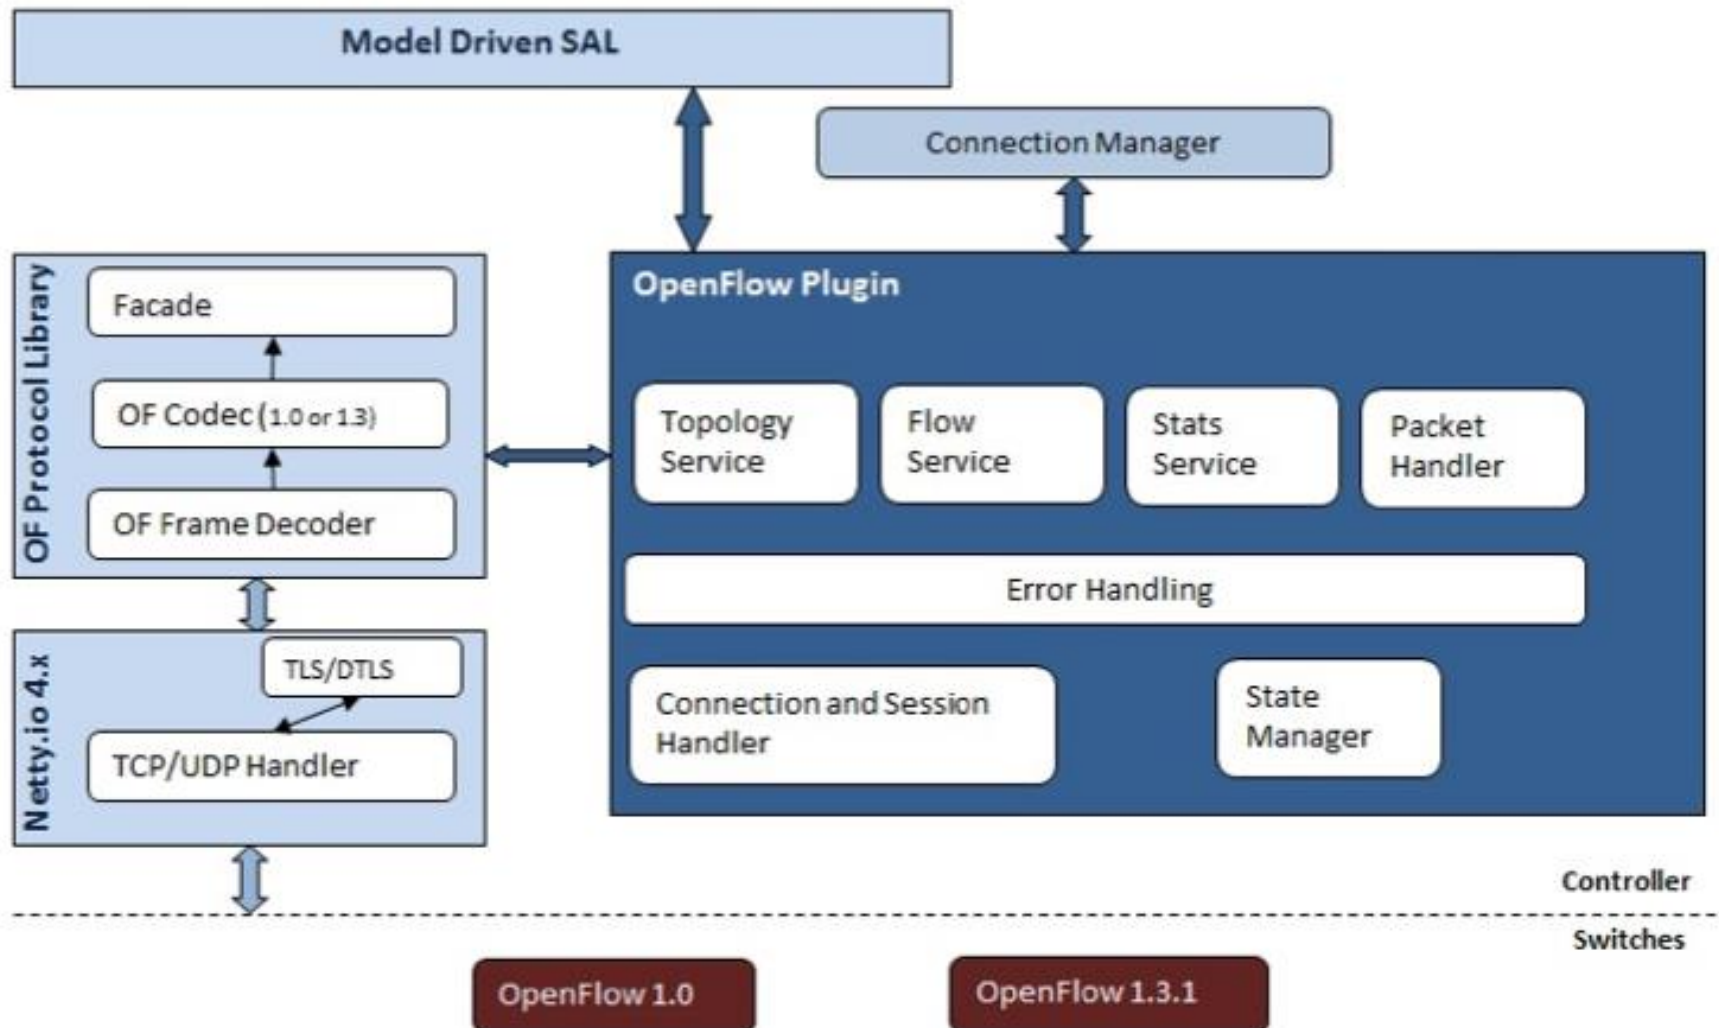
- Built using Java OSGi Framework
 - Provides Modularity & Extensibility
 - Bundle Life-cycle management
 - In-Service-Software Upgrade (ISSU) & multi-version support
- Service Abstraction Layer (SAL)
 - Provides Multi-Protocol Southbound support
 - Abstracts/hides southbound protocol specifics from the applications
- High availability & horizontal scaling using clustering
- Version evolve: Hydrogen, Helium, Lithium, Beryllium, Boron, Carbon, Nitrogen, Oxygen, Fluorine, ...

OpenDaylight Controller

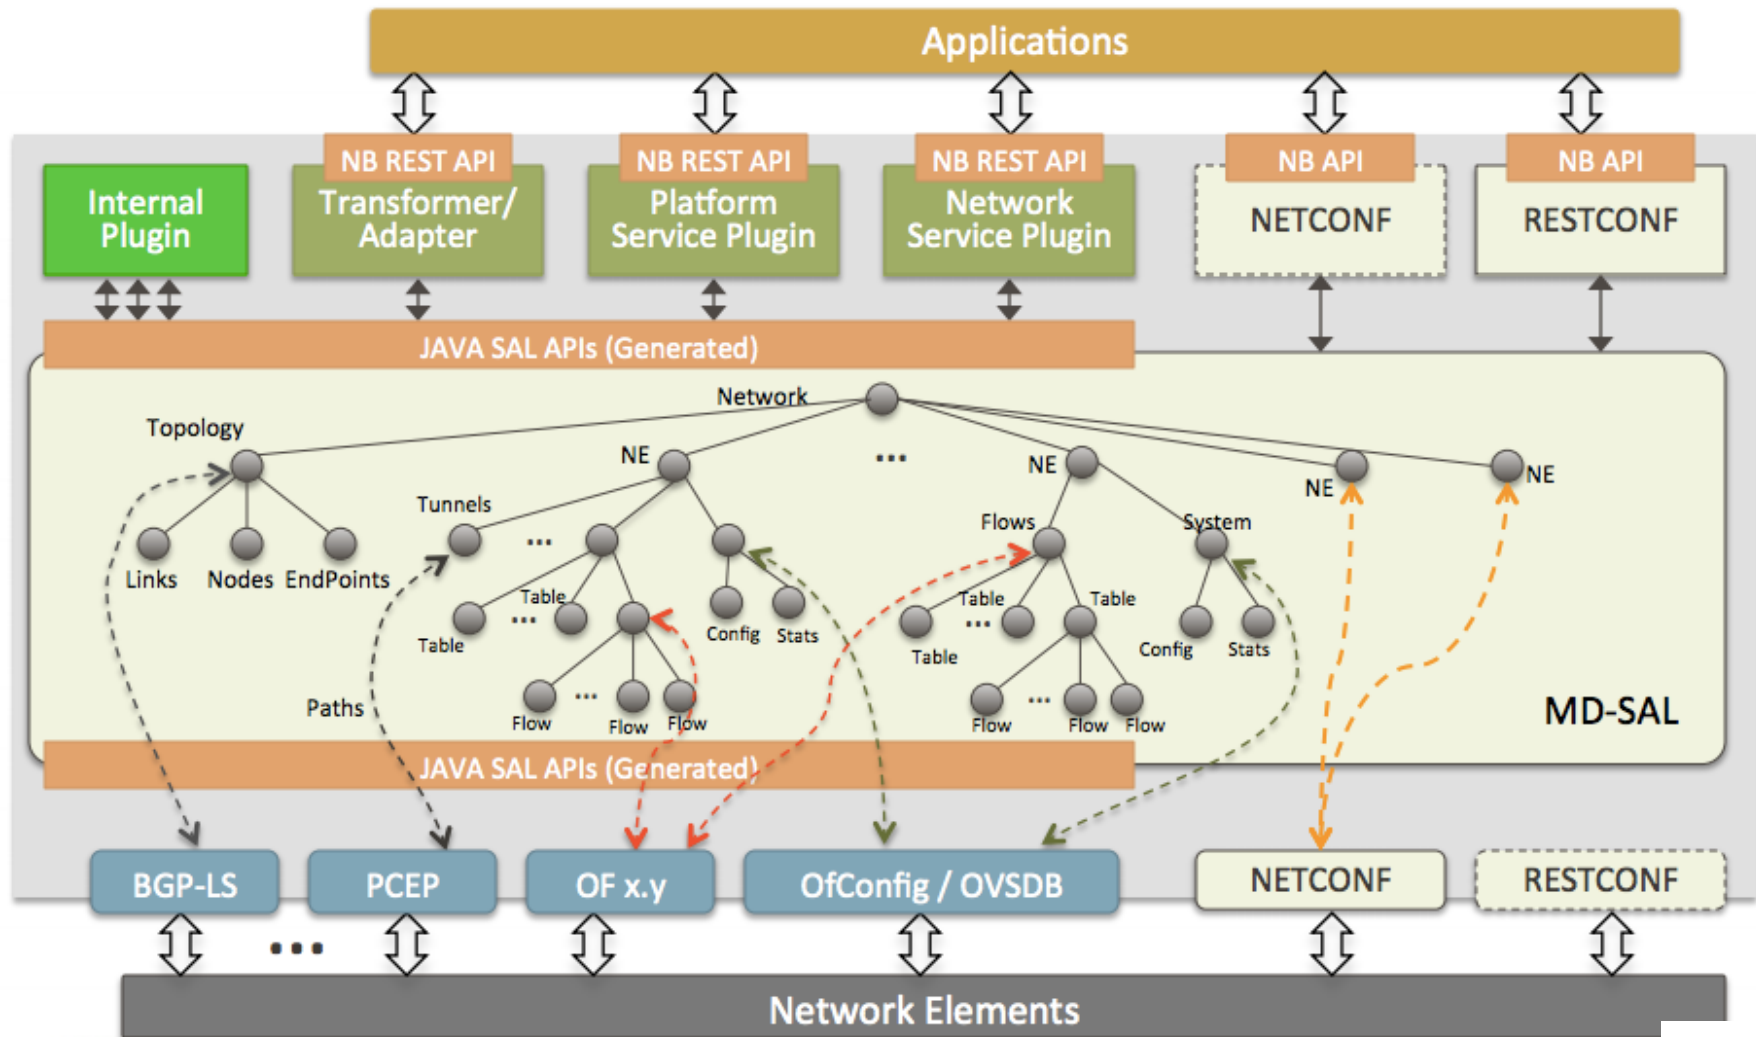
– Plugin Build Process



OpenDaylight Controller

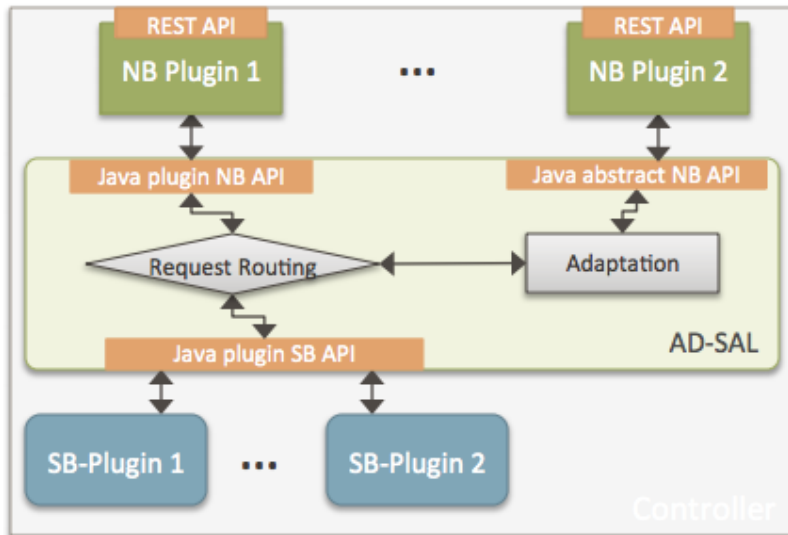


Model-Driven SAL (1/2)

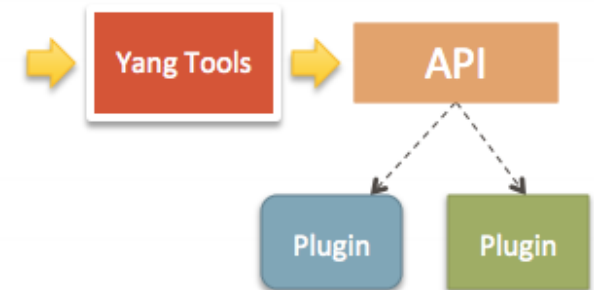
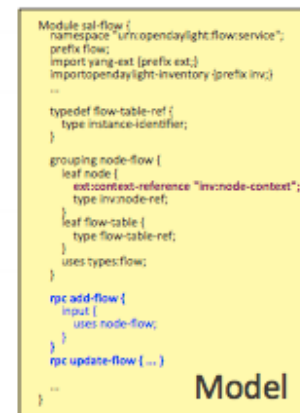


Model-Driven SAL (2/2)

- Model-Driven Service Abstraction Layer (SAL)
 - Yang tools – supporting for model-driven SAL
 - Provides tooling to build Java bindings in yang from yang models

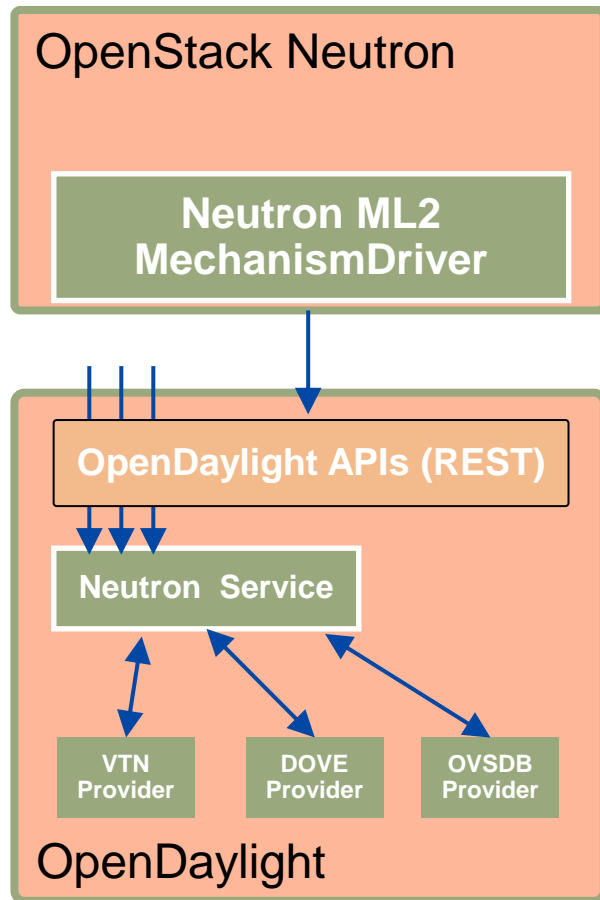


API-Driven SAL



Model-Driven SAL

OpenStack Integration



- OpenDaylight exposes a single common OpenStack Service Northbound
 - API exposed matches Neutron API precisely
 - multiple implementations of Neutron networks in OpenDaylight
 - The Modular Layer 2 (ml2) plugin is a framework allowing OpenStack Networking to simultaneously utilize the variety of layer 2 networking technologies found in complex real-world data centers.
- OpenDaylight OpenStack Neutron Plugin simply passes through
 - simplifies OpenStack plugin
 - pushes complexity to OpenDaylight

Thank you!

References:

<https://www.scs.gatech.edu/news/195201/free-online-sdn-course>

https://www.sdxcentral.com/sdn/?c_action=num_ball

<https://www.opennetworking.org/>



THE UNIVERSITY OF
SYDNEY

