

Topic 6: ONOS Controller


Presented by
Dong YUAN

School of Electrical and Computer
Engineering

dong.yuan@sydney.edu.au



THE UNIVERSITY OF
SYDNEY



Today's Popular controllers

- ONOS and OpenDayLight (ODL)
- ONOS (2014)
 - From Open Networking Foundation
 - Previously ON.LAB funded by Stanford and Berkeley
- ODL (2013)
 - From Linux Foundation
- Both ONOS and ODL are written in Java and designed for modular use with a customizable infrastructure
- Both support OpenStack and K8s
- Every ONOS partner is also an ODL member

Differences

- ONOS vs. ODL
 - Carrier-grade networks vs. Cloud provider
 - Pure SDN vs. Legacy
 - Academic initiated vs. Corporate initiated

Network operating system

- It manages the resources on behalf of users
 - It isolates and protects users from each other
 - It provides useful abstractions
 - Easy access resource
 - Shield difference of devices
 - It provides security from the external world to users
 - It supplies useful services
-
- What information should a Network OS store?
 - Topology? status of network device? route calculated? network traffic?

Introduction to ONOS

- ONOS: Open Network Operating System
 - SDN OS for service provider networks
 - Key features
 - Scalability, high availability & performance
 - Northbound & southbound abstraction
 - Modularity
 - Various usage purposes, customization and development
 - History

ON.LAB

Founded – 2012

ONOS Prototype 1 – 2013
(scalability, high availability)

ONOS Prototype 2 – 2013
(performance)

ONOS **VERSION 1** –
Open sourced on **Dec 5th, 2014**

ONOS Ecosystem

ON.LAB

- Non-profit, Carrier and vendor neutral
- Build core platform
- Provide technical shepherding, core team
- Build community



Service Providers



- Provide funding
- Provide requirements
- Develop use cases
- Drive POCs, deployments
- Bring vendors along

Vendors



- Provide funding
- Provide engineering resources
- Build products and solutions
- Provide integration, test and support services

Community



- Drive every aspect-technical, process, roadmap, deployments
- Bring in diversity
- Help ONOS evolve and thrive

Prior Work

Single Instance

NOX, POX, Beacon, Floodlight, Trema controllers

Helios, Midonet, Hyperflow, Maestro, Kandoo, ...

Distributed: ONIX

Distributed control platform for large-scale networks

ONIX: closed source; datacenter + virtualization focus

ONOS design influenced by ONIX

Community needs an open source distributed network OS

Design Goals

ONOS is a multi-module project whose modules are managed as OSGi bundles, leverages Apache Karaf as application container, and enables YANG schema language to become a programming language.

- Same as the architecture of OpenDayLight
- Code Modularity : It should be possible to introduce new functionalities as self-contained units.
- Configurability : It should be possible to load and unload various features, whether it be at startup or at runtime.
- Separation of Concern : There should be clear boundaries between subsystems to facilitate modularity.
- Protocol agnosticism : It, and its applications, should not be bound to specific protocol libraries or implementations.

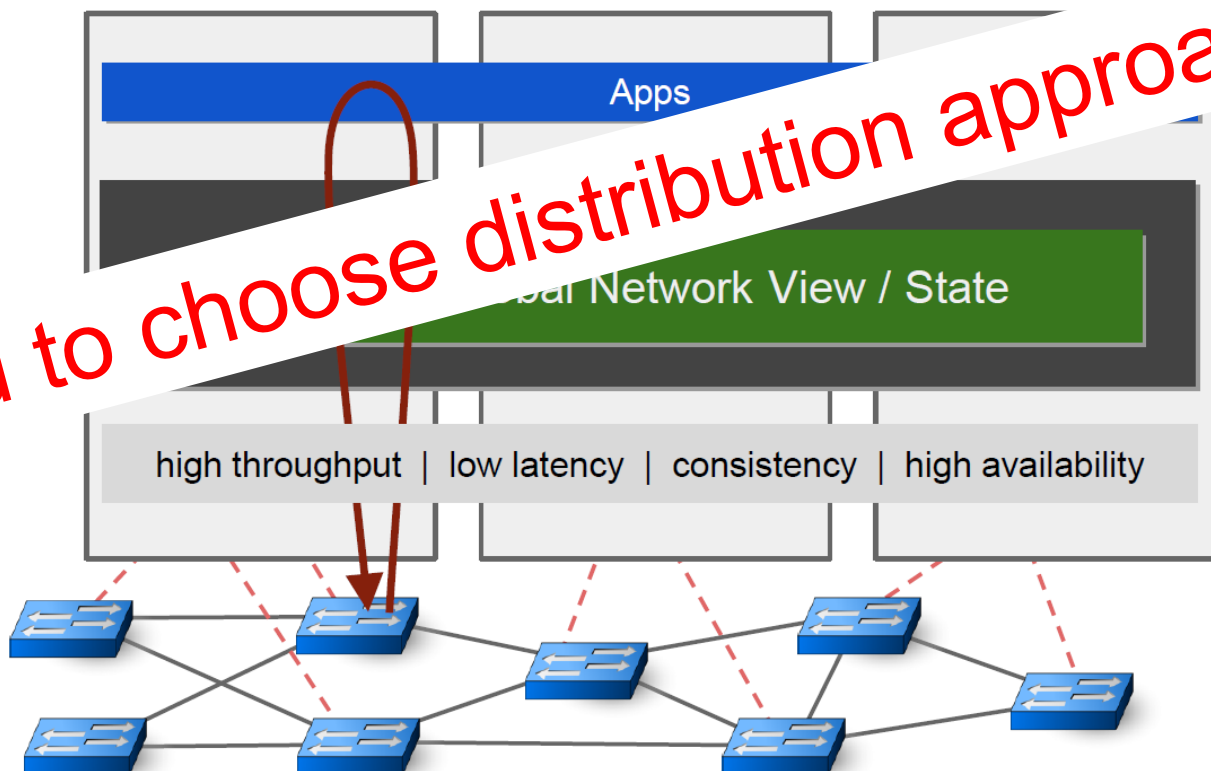
ONOS functions

- Distributed Network OS
 - Network Graph Northbound Abstraction
 - Horizontally Scalable
 - Highly Available
 - Built using open source components
- Exploring performance & reactive computation frameworks
- Expand graph abstraction for more types of network state
- Control functions: intra-domain & inter-domain routing

Key Performance Requirements

- Requirements for Supporting Service Provider Networks
 - High throughput
 - 500K ~ 1M paths setups / second, 3-6M network state operations / second
 - High volume
 - 500GB ~ 1TB of network state data

Need to choose distribution approach!



Key features of ONOS

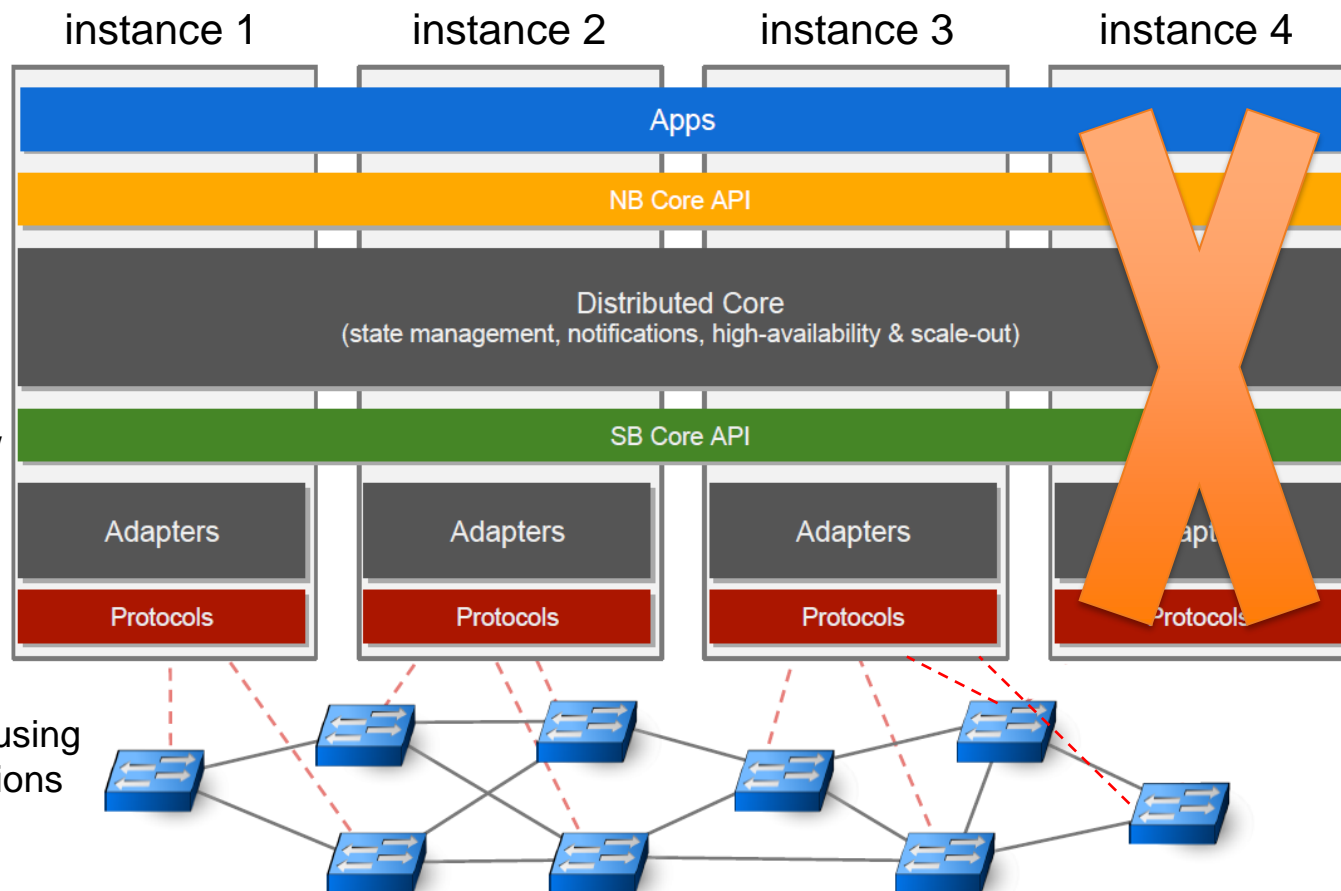
- Distributed Core
 - Scalability, availability and performance
- Northbound Abstraction/APIs
 - Ease the development of control, management and configuration services.
- Southbound abstraction/APIs
 - Enable pluggable southbound protocols for controlling both openflow and legacy devices.
- Software Modularity
 - Easy for community developers to develop, maintain, debug and upgrade ONOS

ONOS Tiers and Distributed Architecture

– Distributed Architecture

- Six tiered architecture
- Each ONOS instance is equipped with the same software stack

- Northbound Abstraction
 - Network graph
 - Application intents
- Core
 - Distributed
 - Protocol independent
- Southbound Abstraction
 - Generalized OpenFlow
 - Pluggable & extensive
- Adapters
 - Multiple southbound protocol enabling layer
- Protocols
 - Self-defined protocols using generalized SDN functions



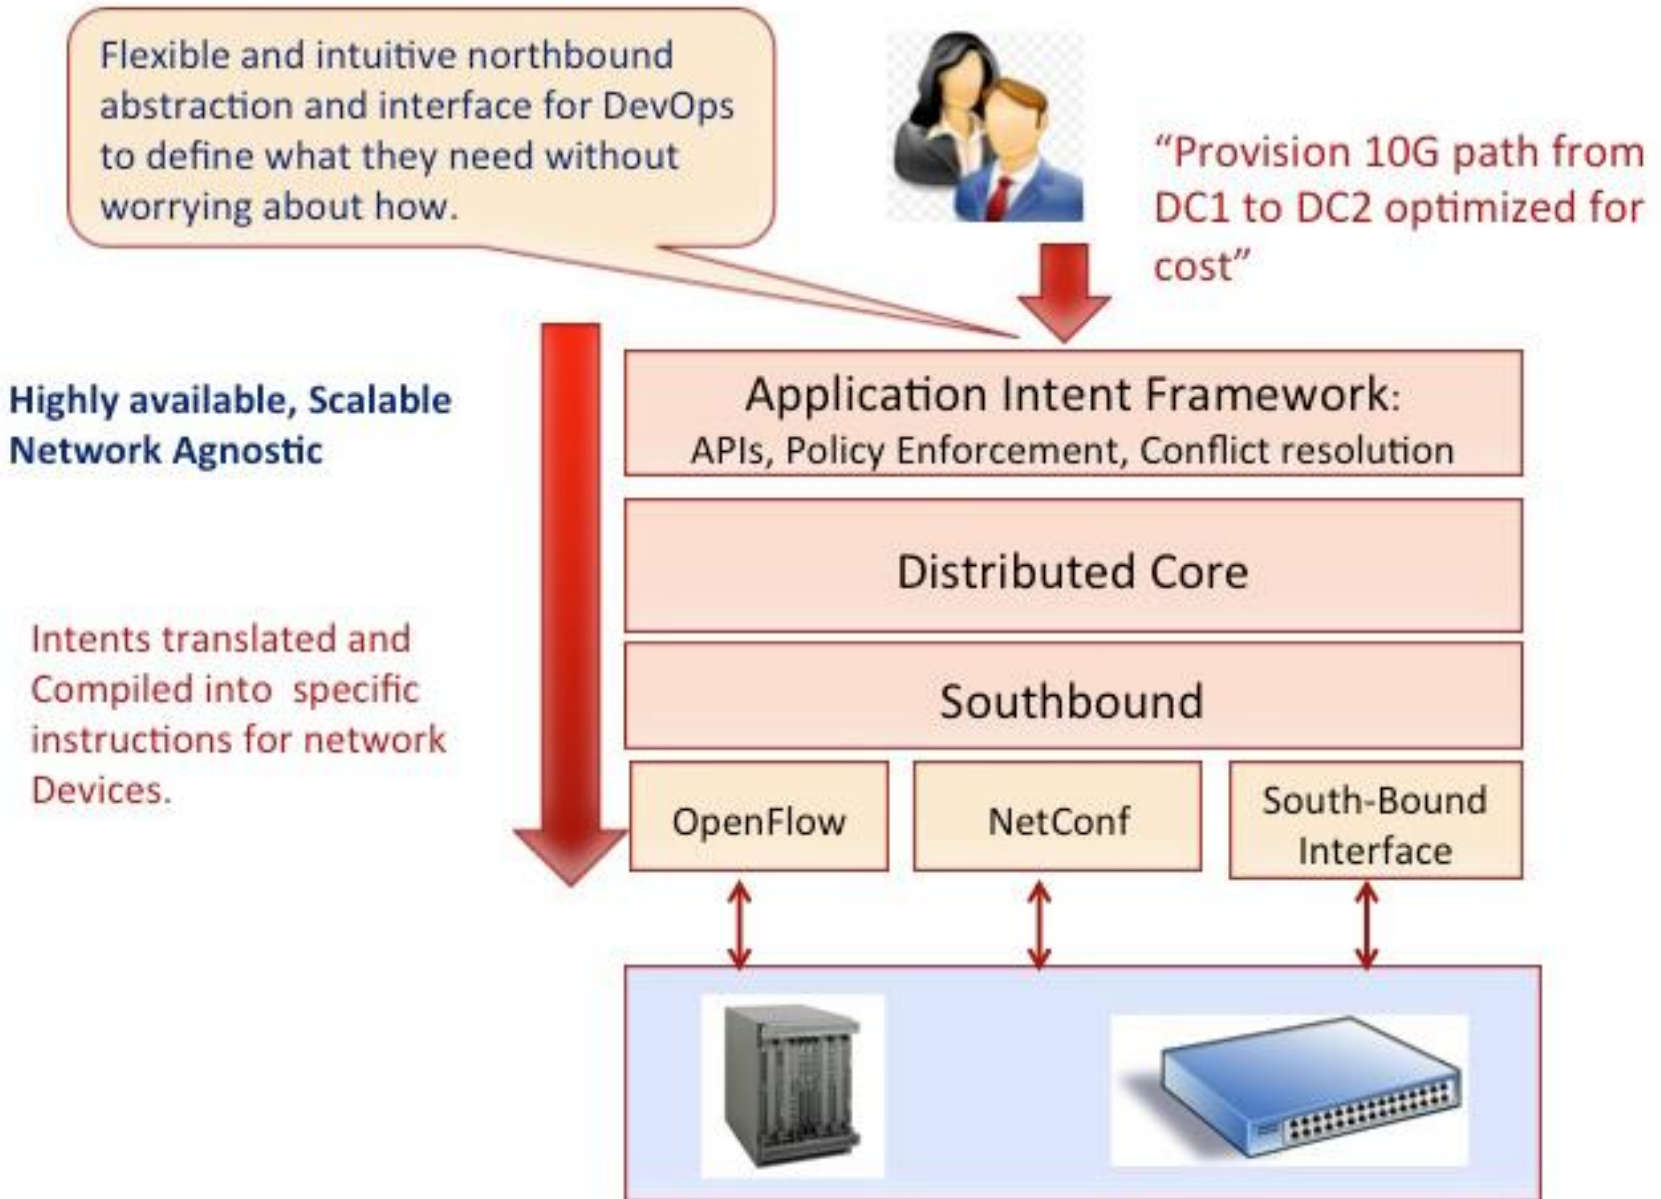
Distributed Cores

- ONOS is deployed as a service on a cluster of servers, same software runs on each server.
- The operator can add servers incrementally, as needed for control plane capacity.
- The ONOS instances work together to create what appears as a single platform.
 - Applications and devices do not have to know if they are working with a single instance or multiple instances
 - This feature makes ONOS scalable
 - The distributed Core does heavy lifting to realise this

Northbound abstraction

- Two powerful northbound abstractions
 - Intent Framework and Global Network View
- The Intent Framework allows an application to request a service from the network without having to know its details.
 - Allow operators and application developers to program the network at a high level
 - They just simply specify their intent, i.e. a script.
- The Global Network View provides a view of the network – the hosts, switches, links, etc.
 - Applications can program this network view through APIs

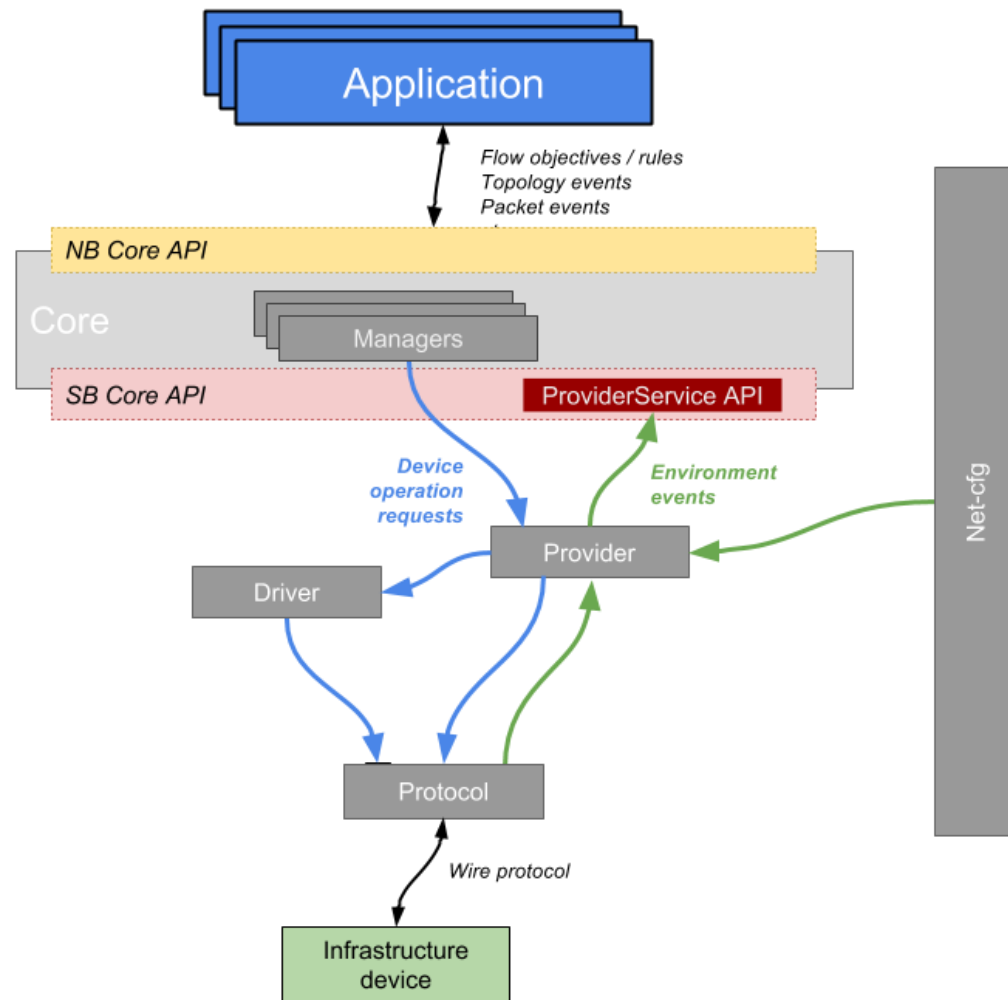
Northbound abstraction



Southbound Abstraction

- The southbound abstraction of ONOS represents each network element as an object in a generic form.
- Allow plug-ins for various southbound protocols and devices.
- Ability to manage different devices using different protocols.
- Ability to add new devices and protocols to the system.
- Ease for migration from legacy devices and protocols

Southbound Design



Provider

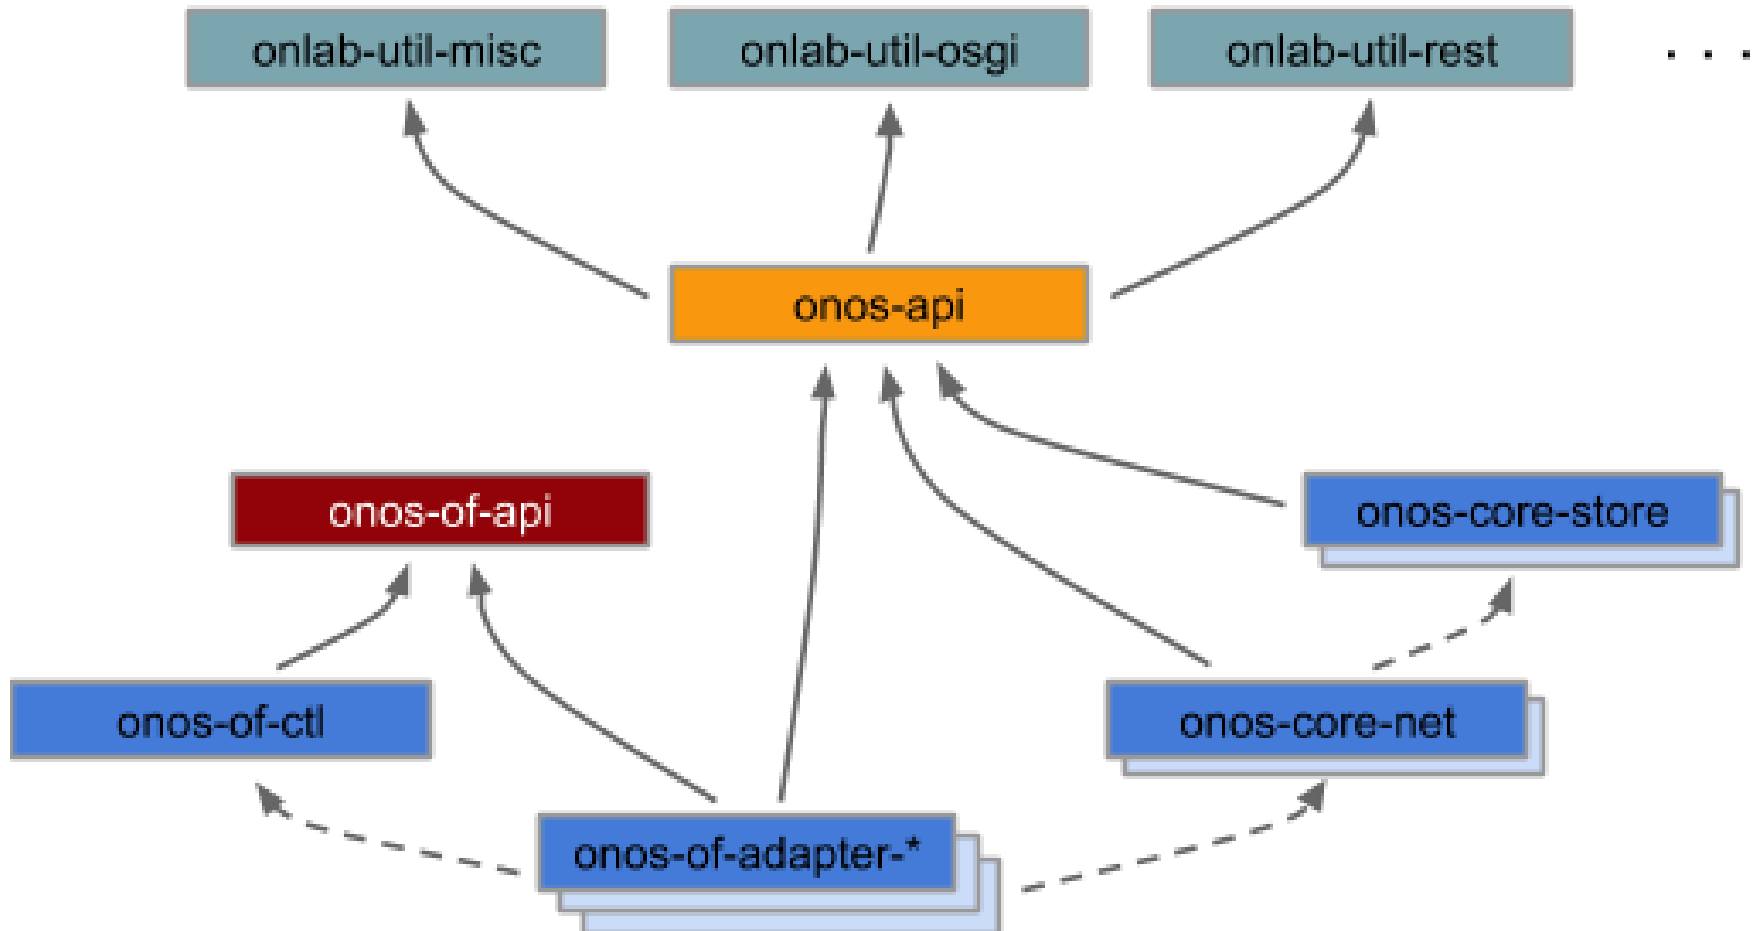
- Providers in ONOS are standalone ONOS applications based on OSGi components that can be dynamically activated and deactivated at runtime.

Protocol

- This module contains the implementation of all those features needed by ONOS to communicate with devices implementing a specific runtime control and/or management protocol. Examples of what ONOS considers protocols are OpenFlow, NETCONF, SNMP, OVSDB, etc.

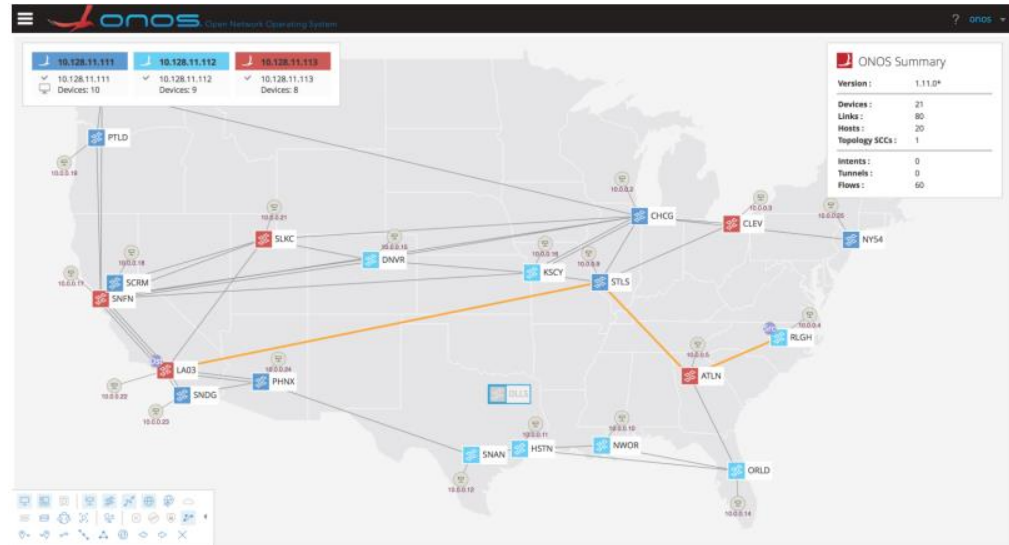
Software Modularity

- It is how the software is structured into components and how those components relate to one another.



Interacting with ONOS

- ONOS Web GUI
- (<http://...:8181/onos/ui>)
- ONOS CLI
 - An extension of Karaf's CLI and Leverage features such programmatic extensibility



```
Welcome to Open Network Operating System (ONOS)!

  _ _ _ _ _
 / _ _ _ _ \
/_ _ _ _ _ \
/_ _ _ _ _ \
/_ _ _ _ _ \
/_ _ _ _ _ \

Documentation: wiki.onosproject.org
Tutorials:    tutorials.onosproject.org
Mailing lists: lists.onosproject.org

Come help out! Find out how at: contribute.onosproject.org

Hit '<tab>' for a list of available commands
and '[cmd] --help' for help on a specific command.
Hit '<ctrl-c>' or type 'system:shutdown' or 'logout' to shutdown ONOS.

onos> |
```

Interacting with ONOS

– REST API

- Provides a way to interact with off-platform applications
- JSON, HTTP/1.1 based communication
- Swagger based REST documents


– gRPC

- Faster access than REST calls by using HTTP/2 connection multiplexing and bidirectional streaming

DELETE	/flows/application/{appld}
GET	/flows/application/{appld}
DELETE	/flows
GET	/flows
POST	/flows
DELETE	/flows/{deviceid}/{flowid}
GET	/flows/{deviceid}/{flowid}
GET	/flows/{deviceid}
POST	/flows/{deviceid}

ONOS High Level Architecture

Network Graph
Eventually consistent



Titan Graph DB

Cassandra In-Memory DHT

Distributed Registry
Strongly Consistent



Zookeeper

Instance 1


**OpenFlow
Controller+**

Instance 2

**OpenFlow
Controller+**

Instance 3

**OpenFlow
Controller+**



The software

◆ Titan Graph Database

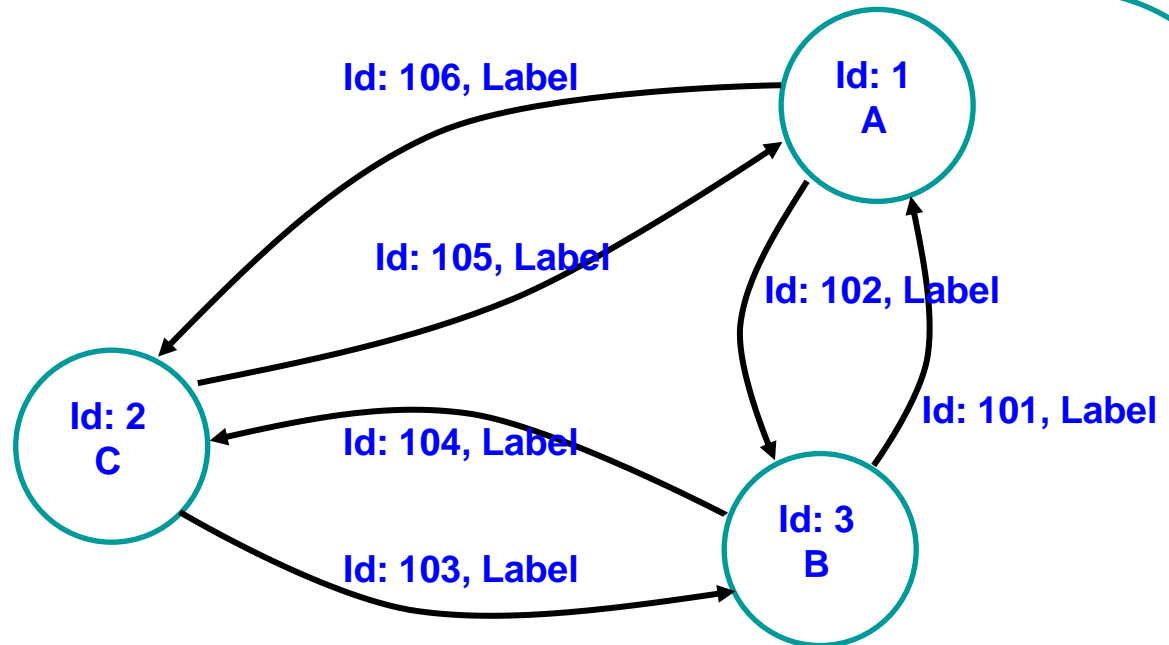

- Titan is a scalable graph database optimized for storing and querying graphs containing hundreds of billions of vertices and edges distributed across a multi-machine cluster.
- Titan is a transactional database that can support thousands of concurrent users executing complex graph traversals in real time.

◆ Apache Cassandra


- ◆ A free and open-source distributed wide column store NoSQL database management system designed to handle large amounts of data across many commodity servers, providing high availability with no single point of failure.
- ◆ Cassandra offers robust support for clusters spanning multiple datacenters, with asynchronous masterless replication allowing low latency operations for all clients.

ONOS Network Graph Abstraction


Network Graph



Cassandra In-memory DHT

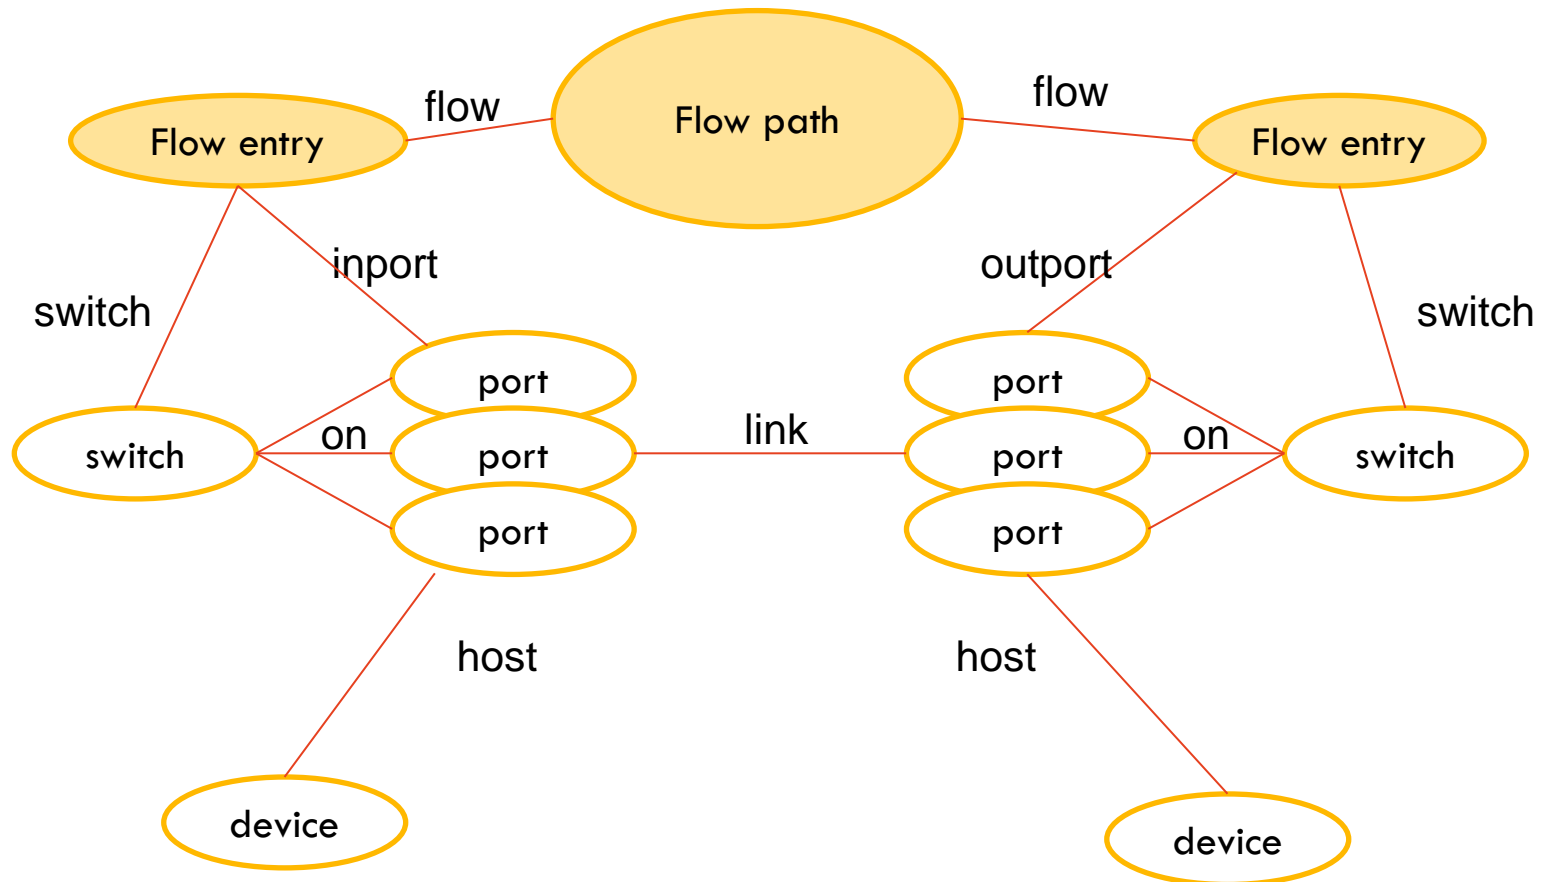


Network Graph




- **Network state is naturally represented as a graph**
- **Graph has basic network objects like switch, port, device and links**
- **Application writes to this graph & programs the data plane**

Example: Path Computation App on Network Graph



- **Application computes path by traversing the links from source to destination**
 - **Application writes each flow entry for the path**
- Thus path computation app does not need to worry about topology maintenance**

Network Graph Representation



Vertex
represented as
Cassandra row

Property (e.g. dpid)	Property (e.g. state)	Property	...	Edge	Edge
----------------------------	-----------------------------	----------	-----	------	------

Row indices for fast vertex centric queries

Edge
represented as
Cassandra
column

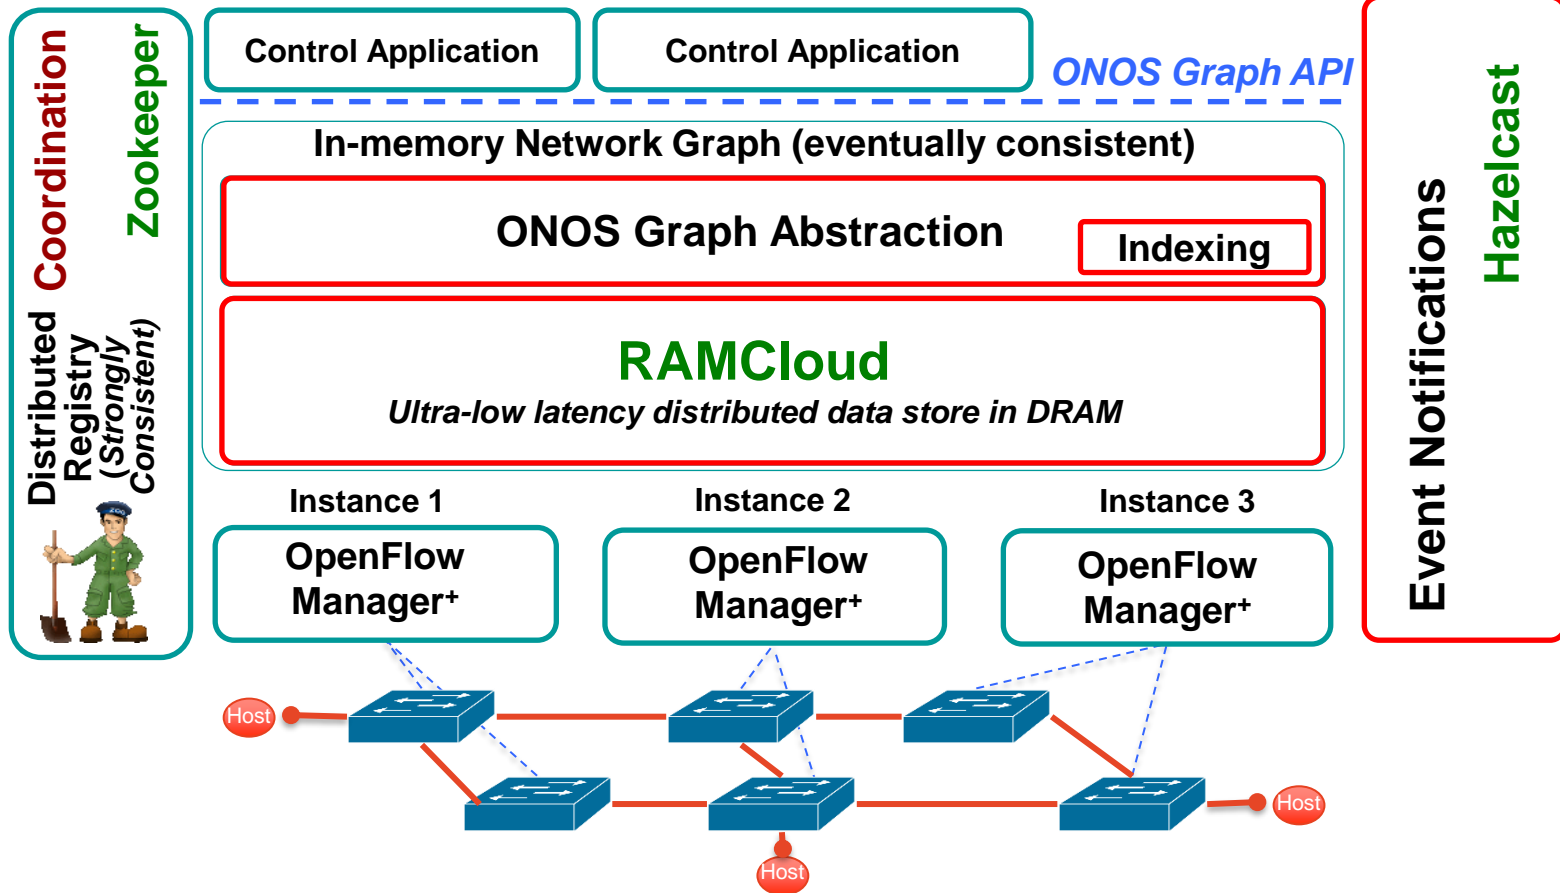
Column				Value	
Label id + direction	Primary key	Edge id	Vertex id	Signature properties	Other properties

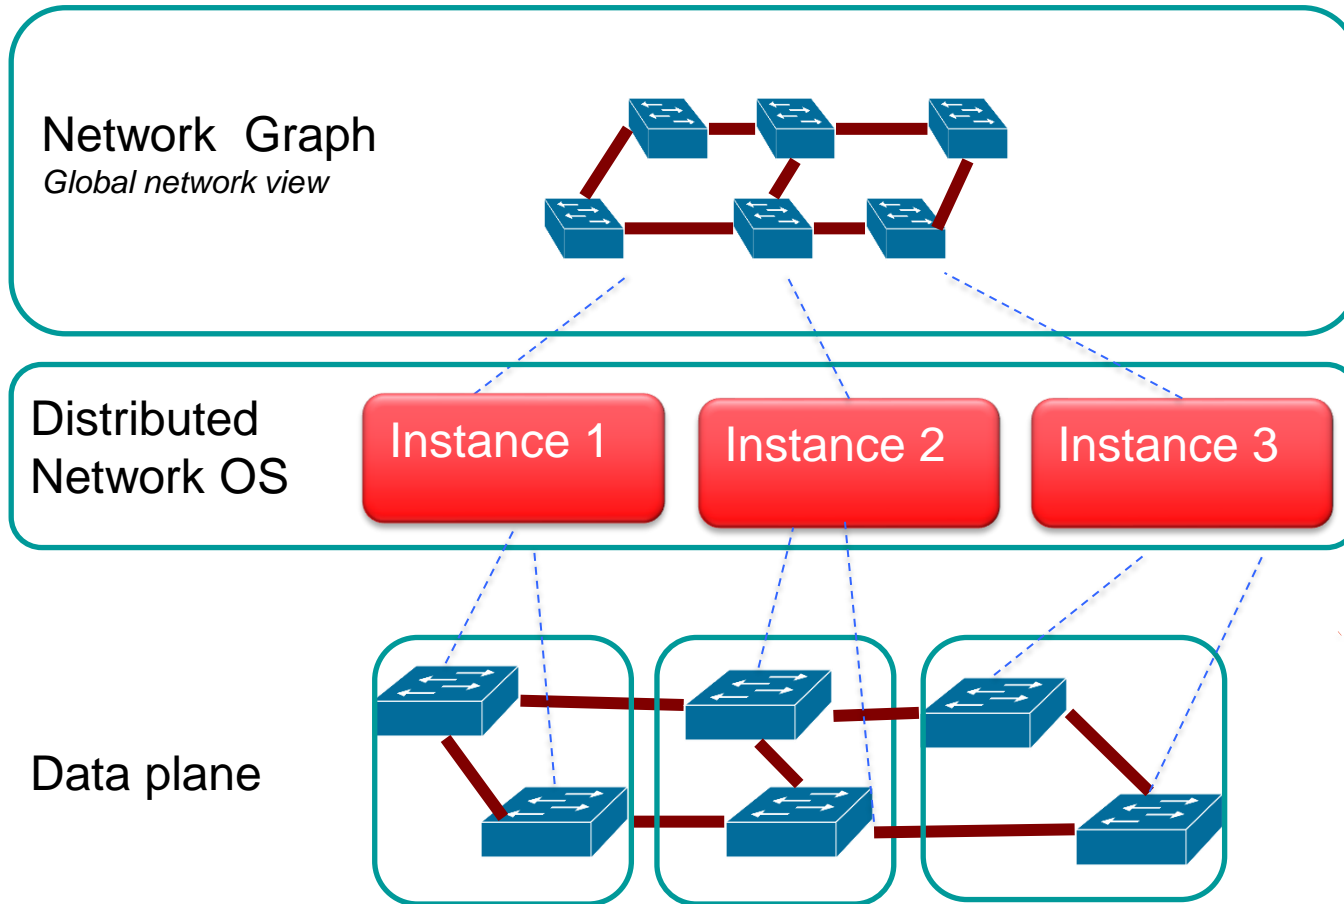
ONOS Architecture (Prototype 2)

Applications

Distributed
Network
Graph/State

Scale-out

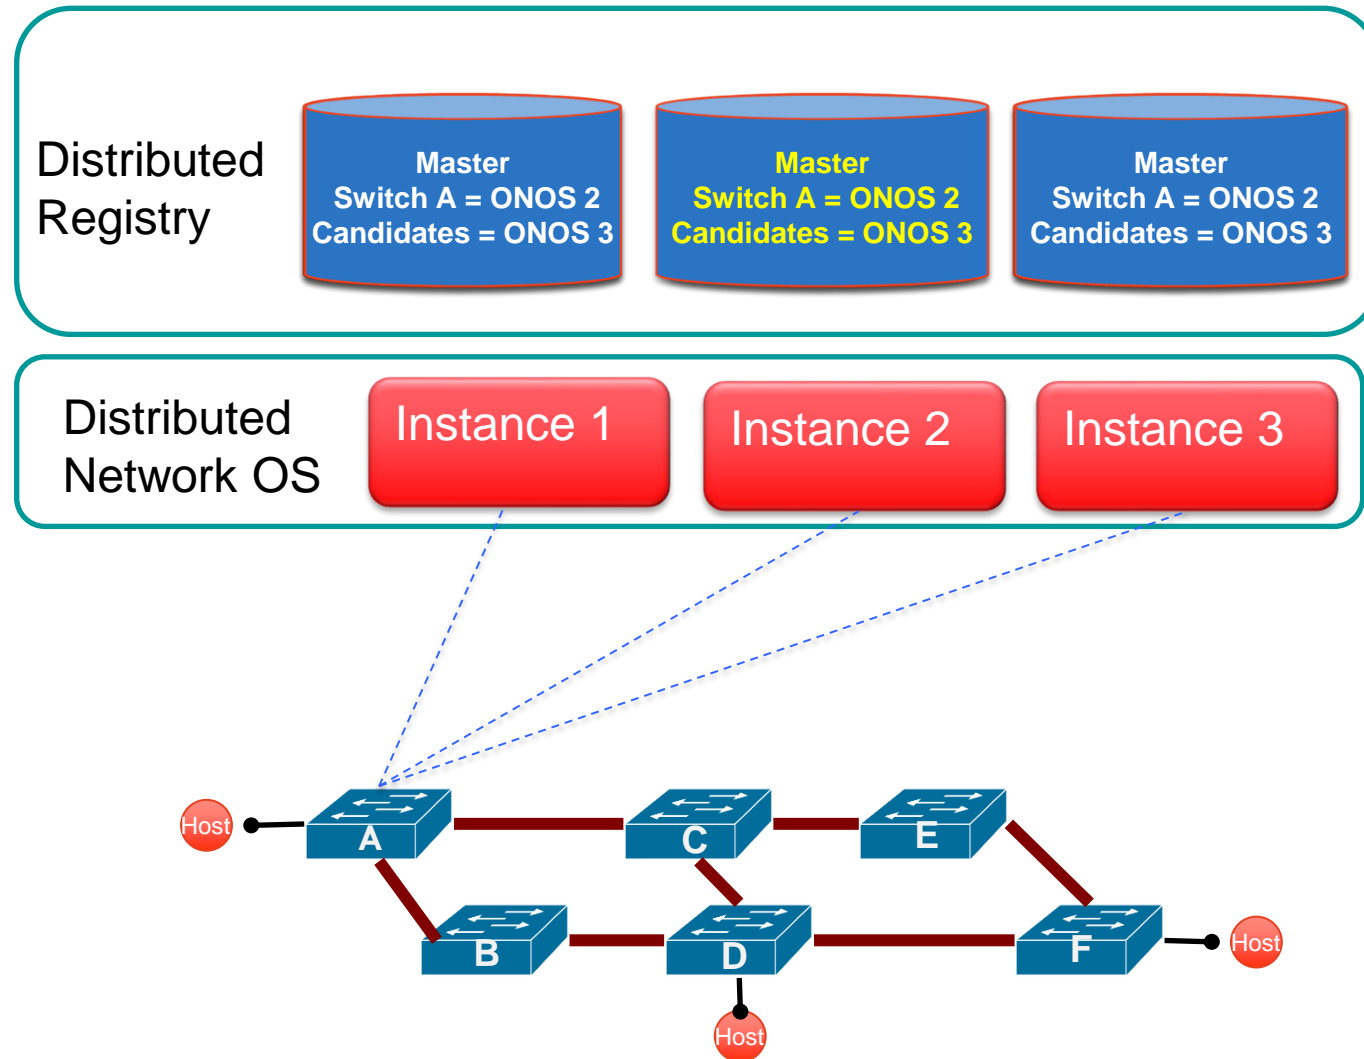




An instance is responsible for maintaining a part of network graph

Control capacity can grow with network size or application need

ONOS Control Plane Failover




Consistency Definition

- Strong Consistency: Upon an update to the network state by an instance, all subsequent reads by any instance returns the last updated value.
- Strong consistency adds complexity and latency to distributed data management.
- Eventual consistency is slight relaxation – allowing readers to be behind for a short period of time.

Strong Consistency using Registry

Network Graph



Distributed
Network OS

Instance 1

Instance 2

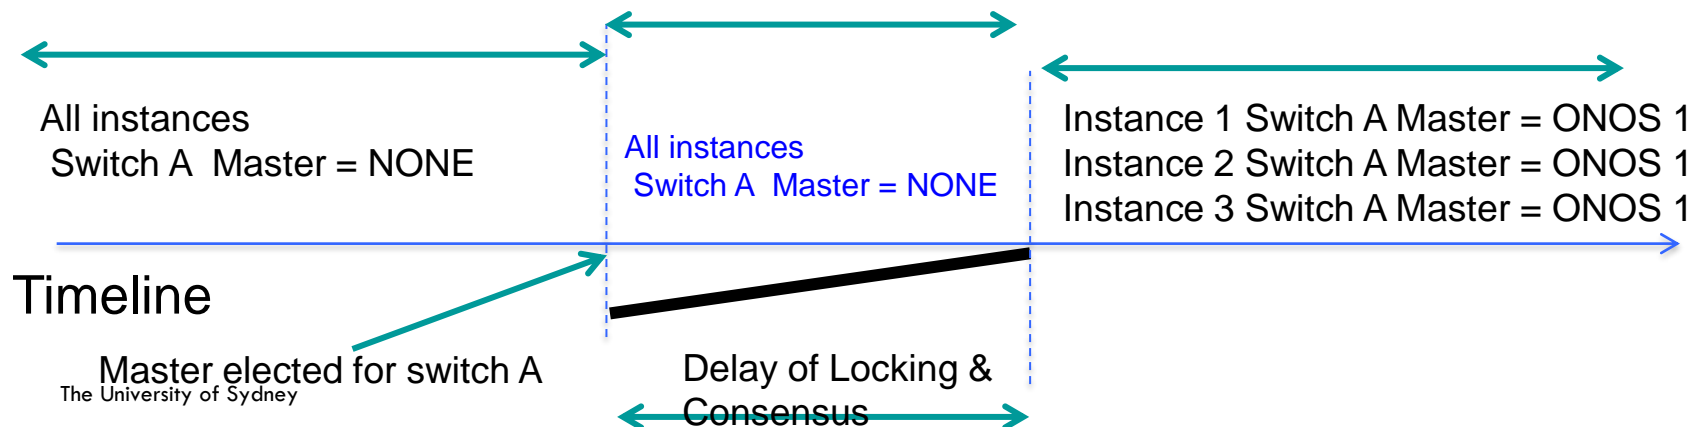
Instance 3

Registry

Switch A
Master = ONOS 1

Switch A
Master = ONOS 1

Switch A
Master = ONOS 1




Why Strong Consistency is needed for Master Election

- Weaker consistency might mean Master election on instance 1 will not be available on other instances.
- That can lead to having multiple masters for a switch.
- Multiple Masters will break the semantic of control isolation.
- Strong locking semantic is needed for Master Election

Eventual Consistency in Network Graph

Network Graph



Distributed
Network
OS

DHT

Instance 1

Instance 2

Instance 3

Switch A
State = ACTIVE

Switch A
State = ACTIVE

Switch A
STATE = ACTIVE

All instances
Switch A STATE = INACTIVE

Instance 1 Switch A = ACTIVE
Instance 2 Switch A = INACTIVE
Instance 3 Switch A = INACTIVE

All instances
Switch A STATE = ACTIVE

Timeline

Switch Connected to ONOS
The University of Sydney

Delay of Eventual Consensus

Cost of Eventual Consistency

- Short delay will mean the switch A state is not ACTIVE on some ONOS instances in previous example.
- Applications on one instance will compute flow through the switch A while other instances will not use the switch A for path computation.
- Eventual consistency becomes more visible during control plane network congestion.

Why is Eventual Consistency good enough for Network State?

- Physical network state changes asynchronously
 - Strong consistency across data and control plane is too hard
 - Control apps know how to deal with eventual consistency
- In the current distributed control plane, each router makes its own decision based on old info from other parts of the network and it works fine
- Strong Consistency is more likely to lead to inaccuracy of network state as network congestions are real.

Consistency learning



- One Consistency does not fit all
- Consequences of delays need to be well understood
- More research needs to be done on various states
using different consistency models

- Distributed Core
 - Responsible for all state management concerns
 - Organized as a collection of “STORES”
 - E.g., topology, links, link resources and etc.
 - State management choices (ACID vs. BASE)
 - ACID (A^{Atomicity}, C^{Consistency}, I^{Isolation}, D^{Durability})
 - BASE (B^{Basically} A^{Available}, S^{Soft state}, E^{Eventually consistency})
- State and Properties

State	Properties
Network Topology	Eventually consistent, low latency access
Flow Rules, Flow Stats	Eventually consistent, shardable, soft state
Switch – Controller Mapping Distributed Locks	Strongly consistent, slow changing
Application Intents Resource Allocations	Strongly consistent, durable

Controller Placement in SDN

- How many controller should we use and where to place them?
 - Optimisation objectives: Average latency, resilience, QoS, throughput, etc.
 - Constrains: Controller Locations, Processing Latency, Number of Controllers, Switch load, Traffic Profile, etc.
 - Graph Model: $G(V, E, C)$, . Let $\phi : V \rightarrow C$ return the controller assigned to a switch, i.e., switch v is assigned to controller $\phi(v)$.



Thank you!

References:

<https://www.scs.gatech.edu/news/195201/free-online-sdn-course>

https://www.sdxcentral.com/sdn/?c_action=num_ball

<https://www.opennetworking.org/>



THE UNIVERSITY OF
SYDNEY

