

Unit Review

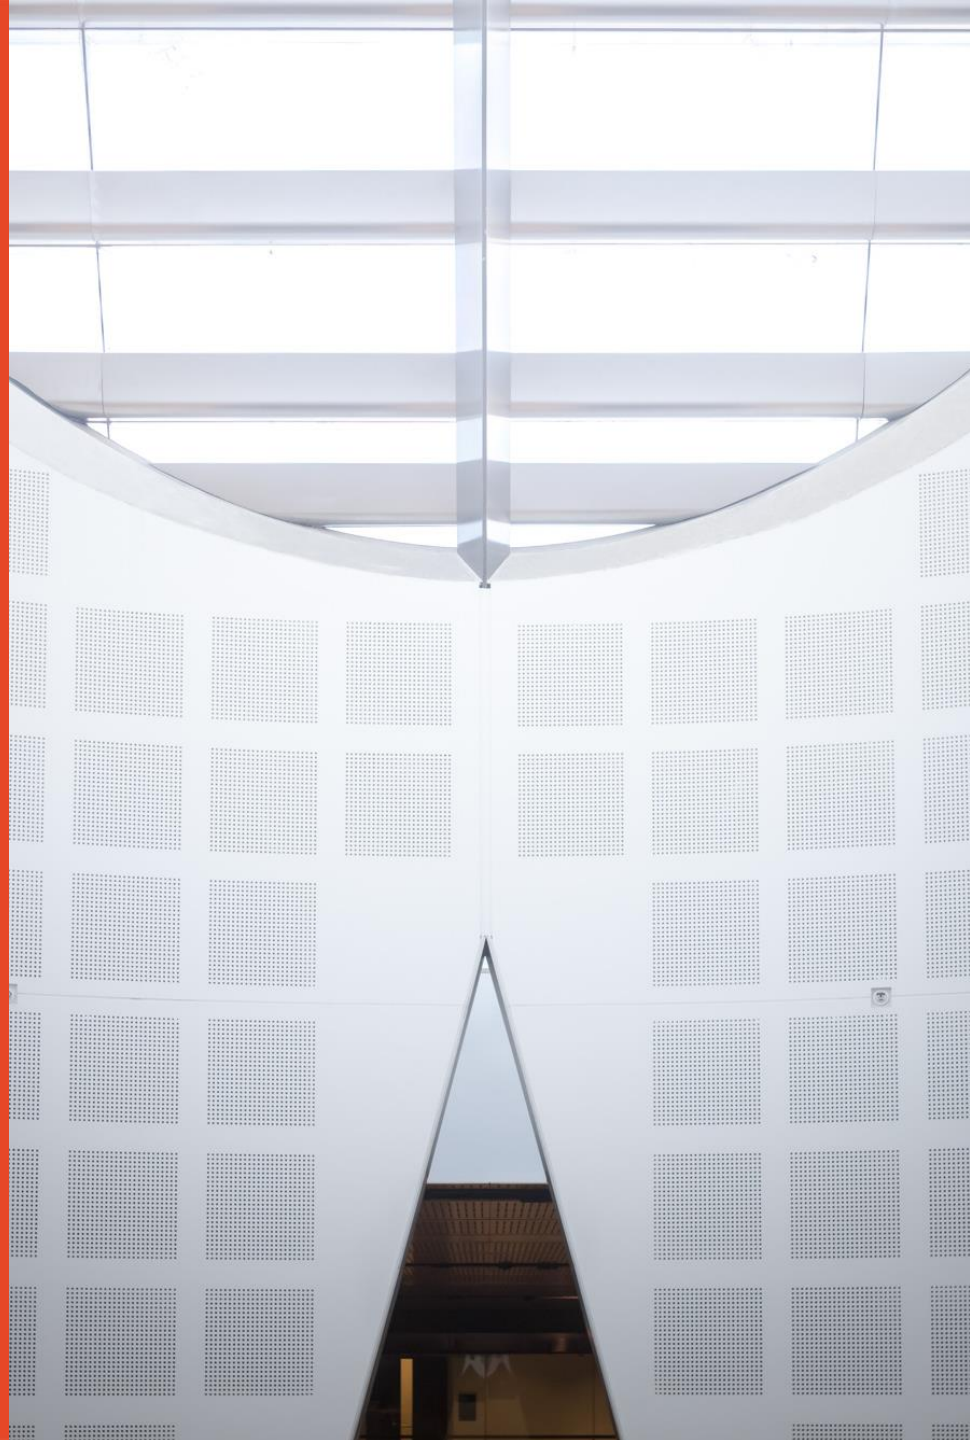
Presented by
Dong YUAN

School of Electrical and Computer
Engineering

dong.yuan@sydney.edu.au



THE UNIVERSITY OF
SYDNEY



Topics covered in this unit

- Introduction to SDN
- Networking Technologies Basics
- OpenFlow and Mininet
- Smart Switches
- Controllers Design
- ONOS Controller
- Programmable Data Planes
- Virtualisation
- Network Function Virtualisation
- SDN applications (CORD & Datacenter)
- SDN applications (SDX & SDWan)

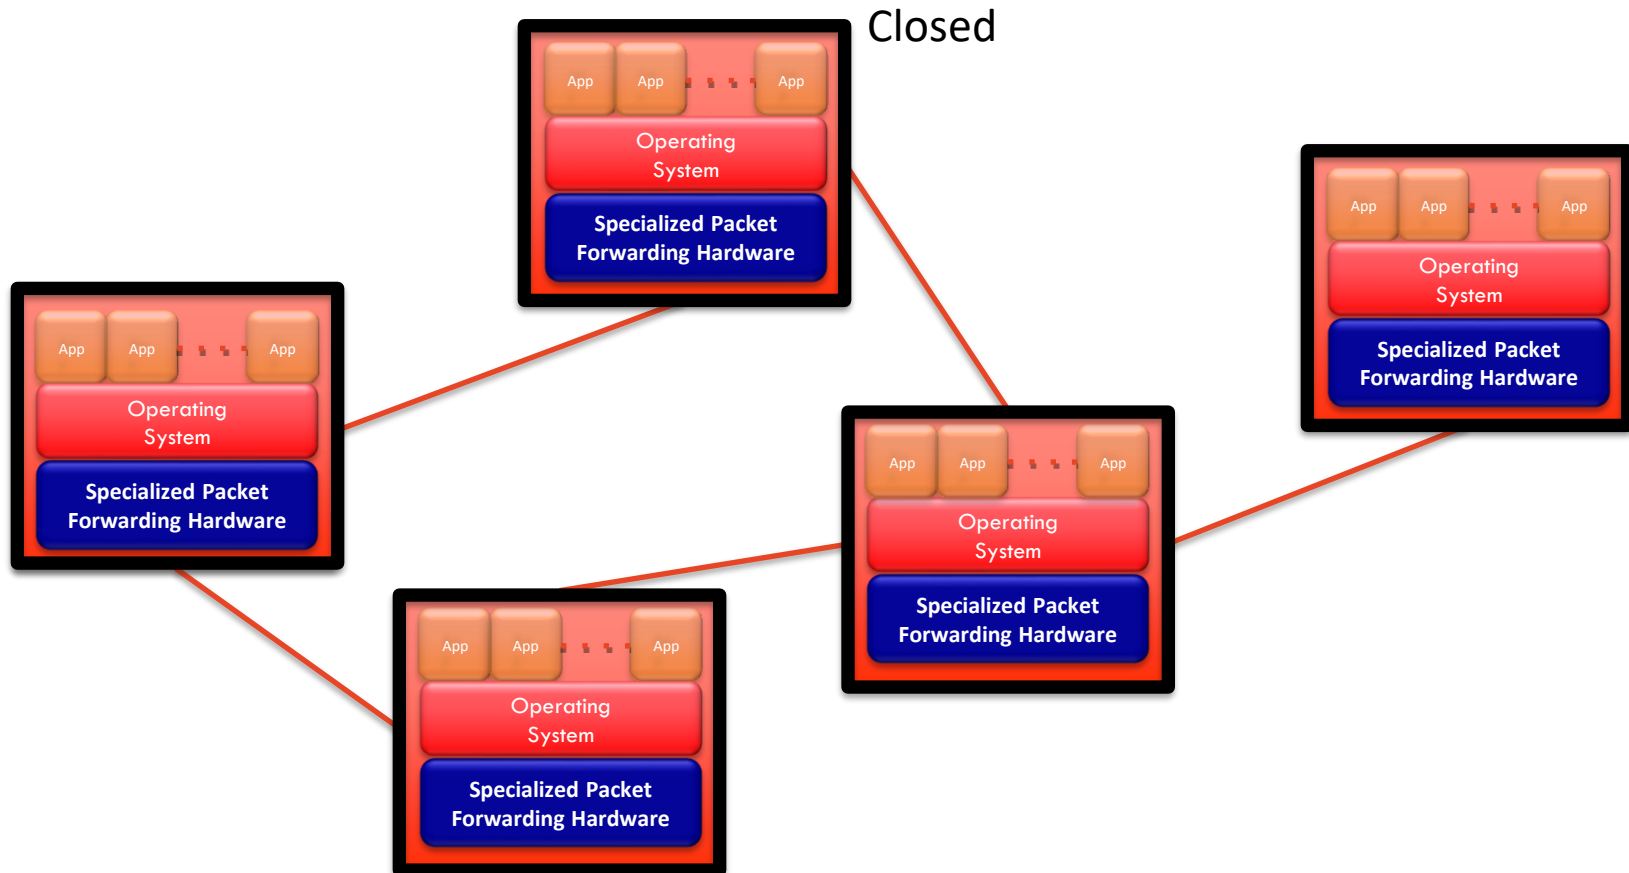
New Internet Architecture

- Cannot make changes to the fundamental technologies of Internet, e.g., IP, routing, etc.
 - High cost to adopt new tech
 - Risks of malfunctions
 - Commercial concerns: need to see the benefit
- Redesign the Internet Architecture
 - **Stanford's Clean Slate** Initiative
 - 100s of project funded world-wide
 - WWW Conference and Journal

What is SDN?

- A new architecture that makes networks more programmable than in the past
- Key principles:
 - Centralised control
 - **Open interfaces**
 - **Flow-based routing**
- **OpenFlow is the most widely adopted protocol.**
 - It is not SDN, but is one important SDN technology

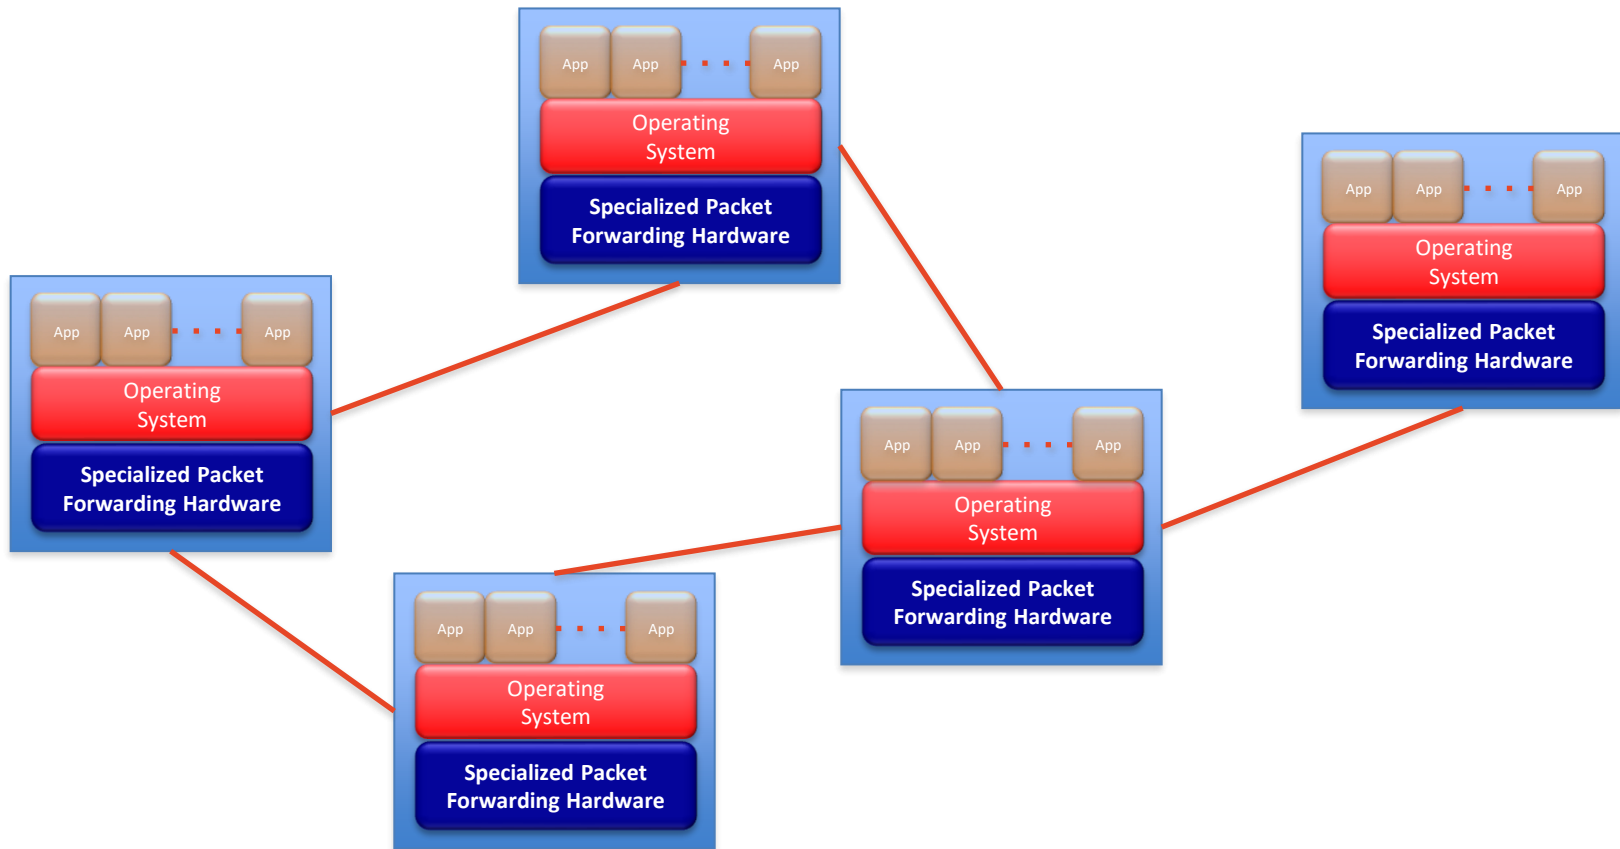
Idea: An OS for Networks



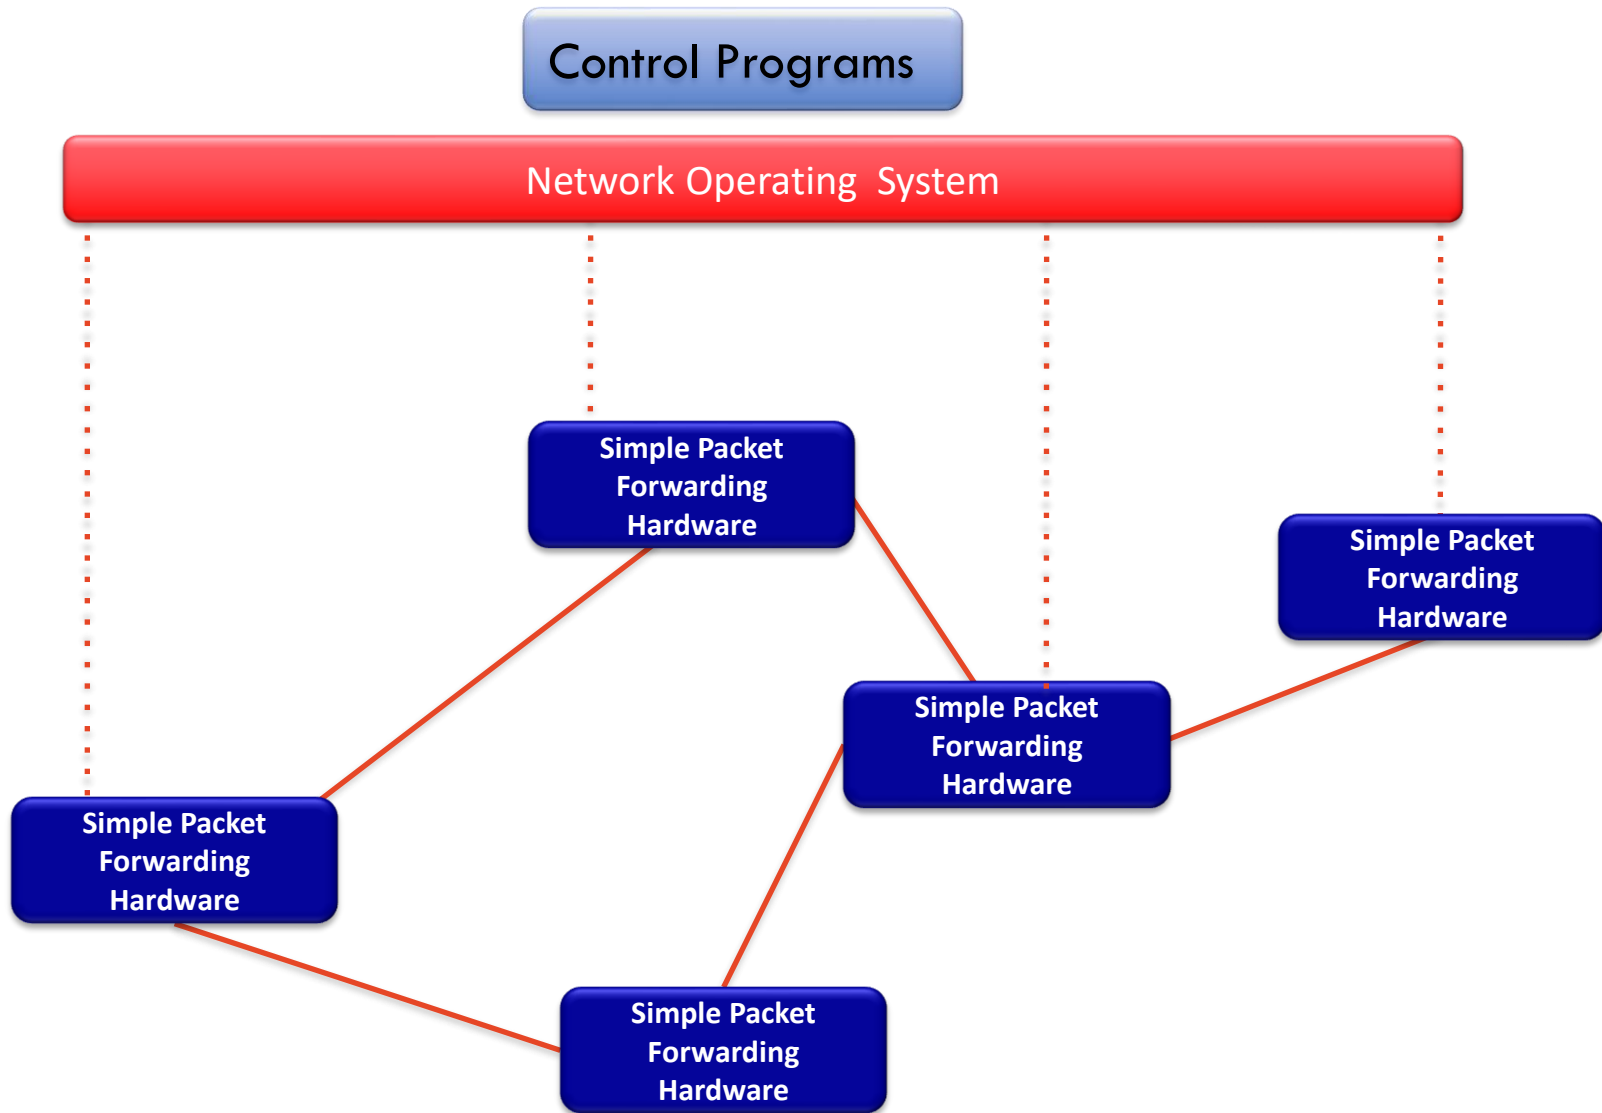
Idea: An OS for Networks

Control Programs



Network Operating System



Idea: An OS for Networks



Key difference



- *Separation of data plane and control plane*

SDN Basic Concept

- Separate Control plane and Data plane entities.
 - Network intelligence and state are logically centralized.
 - The underlying network infrastructure is abstracted from the applications.
- Execute or run Control plane software on general purpose hardware.
 - Decouple from specific networking hardware.
 - Use commodity servers and switches.
- Have programmable data planes.
 - Maintain, control and program data plane state from a central entity.
- An architecture to control not just a networking device but an entire network.

Mininet

- Provides a simple and inexpensive network testbed for developing OpenFlow Applications
 - OpenFlow: protocol to send/receive forwarding rules from controller to switches
 - Some key ideas of OpenFlow: centralization and flow based control
- Enables complex topology testing (without need to wire up a physical network)
- Multiple concurrent developers can work independently on the same topology
- Usable out of the box without programming
- Includes a topology-aware Command Line Interface (CLI) for running or debugging networks
- Supports system-level regression tests (verifies that the previously developed and tested network still performs the same way after changes)

OpenFlow

- OpenFlow-protocol to send/receive forwarding rules from controller to switches
- Control logic is represented as a controller
- Switches perform forwarding
- Packet Arrival: Once packet arrives, the header fields are matched with flow entries in a table, if any entry matches, the counters indicated in that entry are updated and indicated actions are performed.

Flow table

- Match criteria: source, destination, protocol, etc
 - Defined and added dynamically (so called SDN)
- Actions:
 - flow to which port
 - sending packet to controller
 - Drop a packet,
 - Modify packet header, etc.
- Priority: for match criteria
 - Time out period:
 - Hard TO, Idle TO

Controller

- In openflow, controller communicates with switch over a secure channel
 - Openflow protocol defines message format
 - Purpose of control channel: update flow table
 - Logic is executed at controller
 - Communication with external controllers
- Flow Table: Packet switching
 - All packets compare to flow table for match
 - Packet header fields matched against one of the N tables
 - Actions depend on match being found
 - Forward, Drop, modify, enqueue
 - If no match, traffic is send to controller

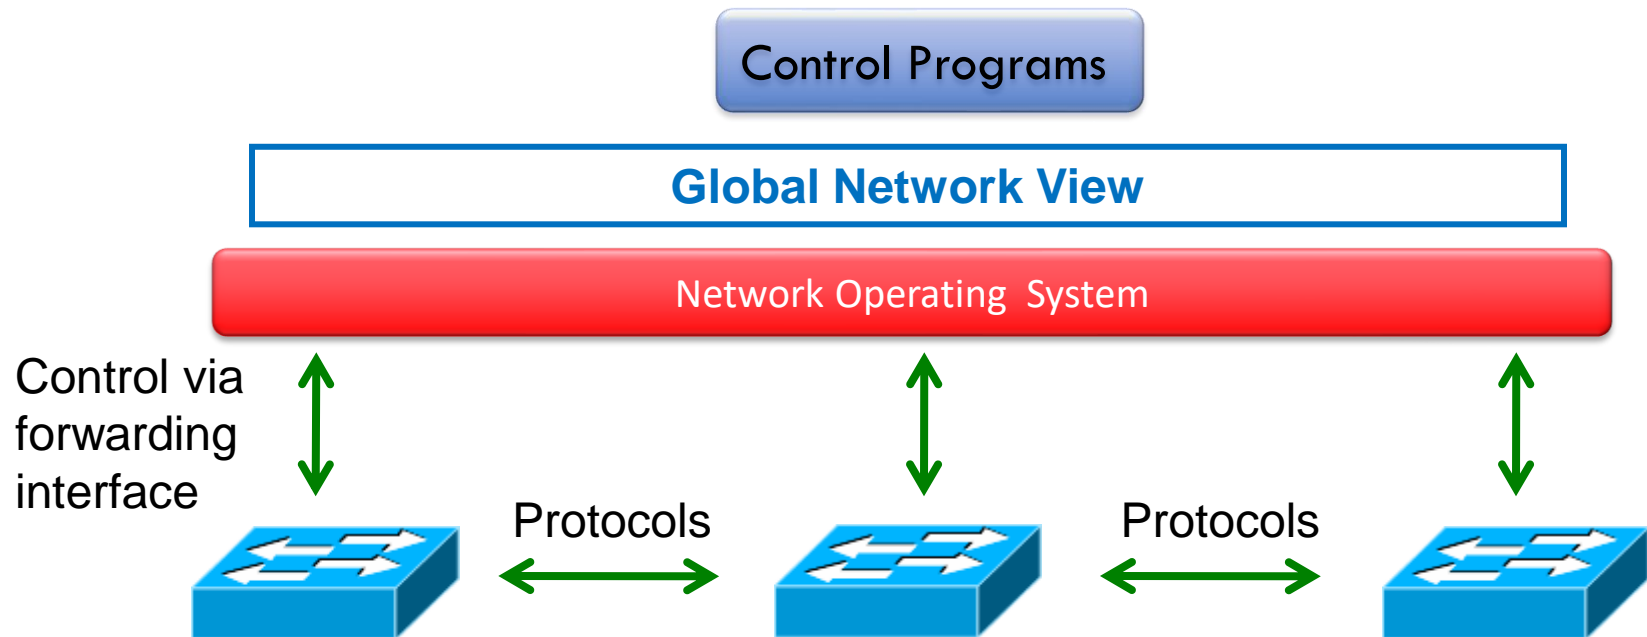
Openflow switches

– Whitebox Switch

- White box switches refers to the ability to use ‘generic,’ off-the-shelf switching (or white box switching) and routing hardware, in the forwarding plane of a software-defined network (SDN).
- They represent the foundational element of the commodity networking ecosystem
- A common operating system for white box switches is Linux-based because of the many open and free Linux tools available that help administrators customize the devices to their needs.

Idea: An OS for Networks

Software-Defined Networking (SDN)



Network OS

- Maintain an up-to-date view of the network state (e.g., topology, etc.)
 - Next page
- Configure network elements
 - A.k.a “south-bound” interface or API
- Provide a graph abstraction to the applications on the top
 - A.k.a “north-bound” interface or API

Network View (graph abstraction)

- In the network view:
 - Switch-level topology
 - Location of host, middlebox and other network elements
 - Location of users
 - Namespace: bindings between names and addresses
- Not in the network view:
 - Some states of network, e.g., traffic.
- Network Information Base (NIB)
 - Graph and abstraction

Challenges in building an NOS

- Scalability
- Reliability
- Good performance (fast, etc.)
- Generality and Simplicity of the “north-bound” API.
 - Easily to be used by applications

Today's Popular controllers

- ONOS and OpenDayLight (ODL)
- ONOS (2014)
 - From Open Networking Foundation
 - Previously ON.LAB funded by Stanford and Berkeley
- ODL (2013)
 - From Linux Foundation
- Both ONOS and ODL are written in Java and designed for modular use with a customizable infrastructure
- Both support OpenStack
- Every ONOS partner is also an ODL member

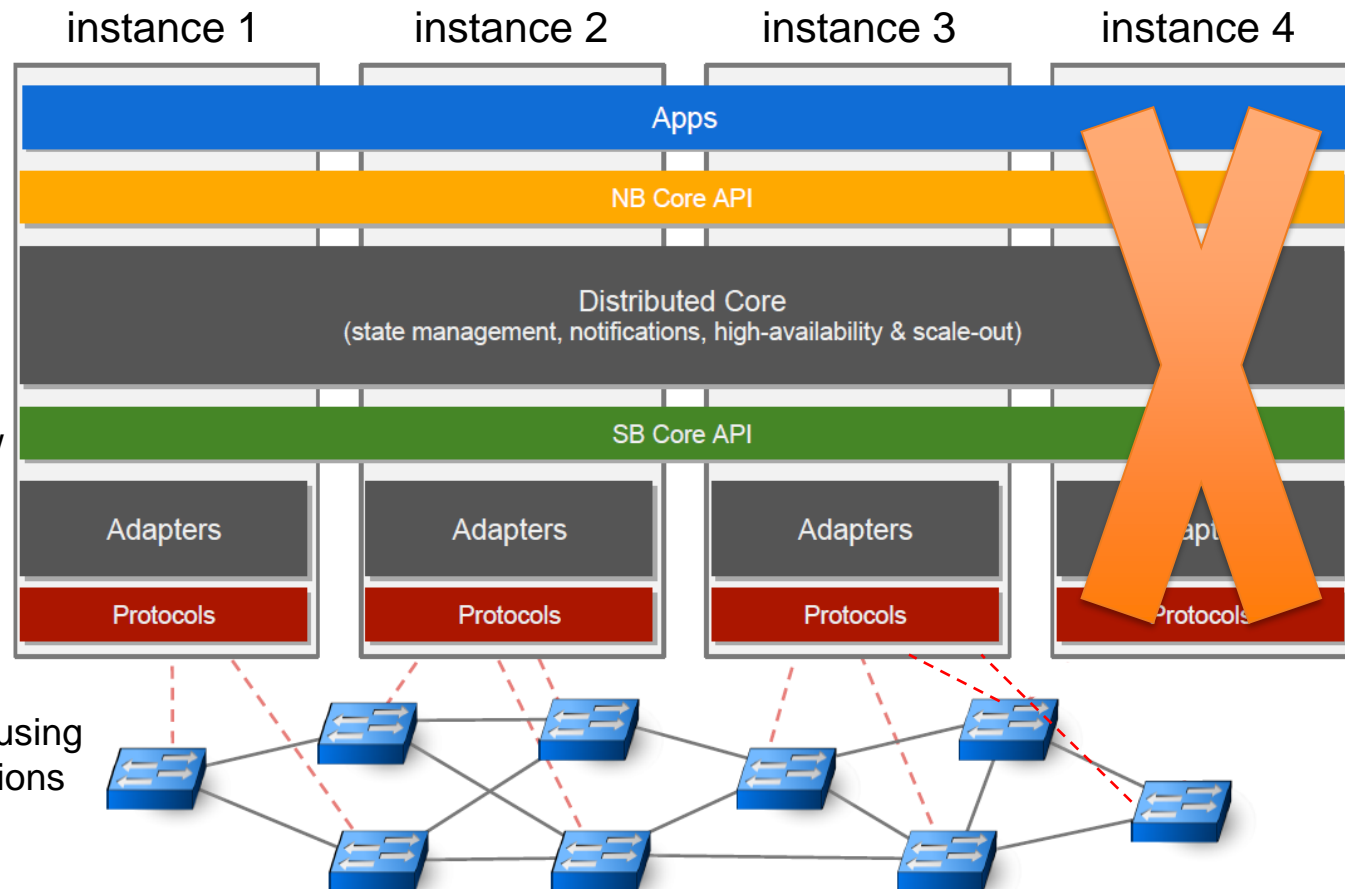
Differences

- ONOS vs. ODL
 - Carrier-grade networks vs. Cloud provider
 - Pure SDN vs. Legacy
 - Academic initiated vs. Corporate initiated

ONOS Tiers and Distributed Architecture

- Distributed Architecture
 - Six tiered architecture
 - Each ONOS instance is equipped with the same software stack

- Northbound Abstraction
 - Network graph
 - Application intents
- Core
 - Distributed
 - Protocol independent
- Southbound Abstraction
 - Generalized OpenFlow
 - Pluggable & extensive
- Adapters
 - Multiple southbound protocol enabling layer
- Protocols
 - Self-defined protocols using generalized SDN functions



Consistency Definition

- Strong Consistency: Upon an update to the network state by an instance, all subsequent reads by any instance returns the last updated value.
- Strong consistency adds complexity and latency to distributed data management.
- Eventual consistency is slight relaxation – allowing readers to be behind for a short period of time.

- Distributed Core
 - Responsible for all state management concerns
 - Organized as a collection of “STORES”
 - E.g., topology, links, link resources and etc.
 - State management choices (ACID vs. BASE)
 - ACID (A^{Atomicity}, C^{Consistency}, I^{Isolation}, D^{Durability})
 - BASE (B^{Basically} A^{Available}, S^{Soft state}, E^{Eventually consistency})
- State and Properties

State	Properties
Network Topology	Eventually consistent, low latency access
Flow Rules, Flow Stats	Eventually consistent, shardable, soft state
Switch – Controller Mapping Distributed Locks	Strongly consistent, slow changing
Application Intents Resource Allocations	Strongly consistent, durable

Overview of data plane

- Wide range of functions
 - Forwarding
 - Access control
 - Mapping header fields
 - Traffic monitoring
 - Buffering and marking
 - Shaping and scheduling
 - Deep packet inspection
- Data plane design goals
 - Flexible
 - Extensible
 - Clean interfaces

Motivation

- SDN protocols require data-plane changes
- Performance requirement
 - Protocols must forward packets at acceptable speeds.
- Support different protocols
 - Run in parallel with existing protocols
- Requirement of SDN data plane – a platform that
 - Forwards packets at high speed
 - Run multiple protocols in parallel

Existing Approaches

- Develop Custom Software
 - Flexible, easy to program
 - Slow forwarding speeds
- Develop Custom Hardware
 - Excellent performance
 - Long development cycles, rigid
- Develop programmable hardware
 - Flexible and fast
 - Programming is difficult

Network Assembly Language

- Openflow's design was motivated by the underlying device layout
 - Controller is limited in supporting new functions not supported by Openflow
- New Chipsets are adding data plane functions.
- New languages are specifying data plane at a high level – people like to use
- What's in between?
- <http://netasm.cs.princeton.edu/>

Need for network assembly


- A low-level programming language for programmable network devices
- Provides a one to one correspondence with the underlying hardware
- Uses well-defined constructs to define low-level packet operations
- Enables writing highly optimized network programs

P4: Programming protocol-independent packet processors

- P4 is a high-level language for programming protocol-independent packet processors.
- P4 works in conjunction with SDN control protocols like OpenFlow.
 - OpenFlow explicitly specifies protocol headers on which it operates. This set has grown from 12 to 41 fields in a few years, increasing the complexity of the specification.
- P4 propose how OpenFlow should evolve in the future.


What is virtualization?

Virtualization is a broad term (virtual memory, storage, network, etc)
Virtualization basically allows one computer to do the job of multiple computers, by sharing the resources of a single hardware across multiple environments



'Nonvirtualized' system

A single OS controls all hardware platform resources



Virtualized system

It makes it possible to run multiple Virtual Containers on a single physical platform

What is network virtualization?

Network Virtualization is the process of combining hardware and software network resources and network functionality into a single, software-based administrative entity, a virtual network. - In computing

- Two categories :
 - External network virtualization
 - Combining many networks, or parts of networks, into a virtual unit.
 - Internal network virtualization
 - Providing network-like functionality to the software containers on a single system.
 - This was before SDN and NFV


NFV

- Combining forwarding devices and middleboxes into a common control framework
- NFV enables network operators to implement network policies without worrying about:
 - Placement: Where to place the functions (middleboxes) in the network.
 - Steering: How to route traffic through these functions.
- Placement and steering are two difficult problems for the traditional network with middleboxes.

NFV Infrastructure as a Service (NFVlaaS)

NFV Infrastructure :

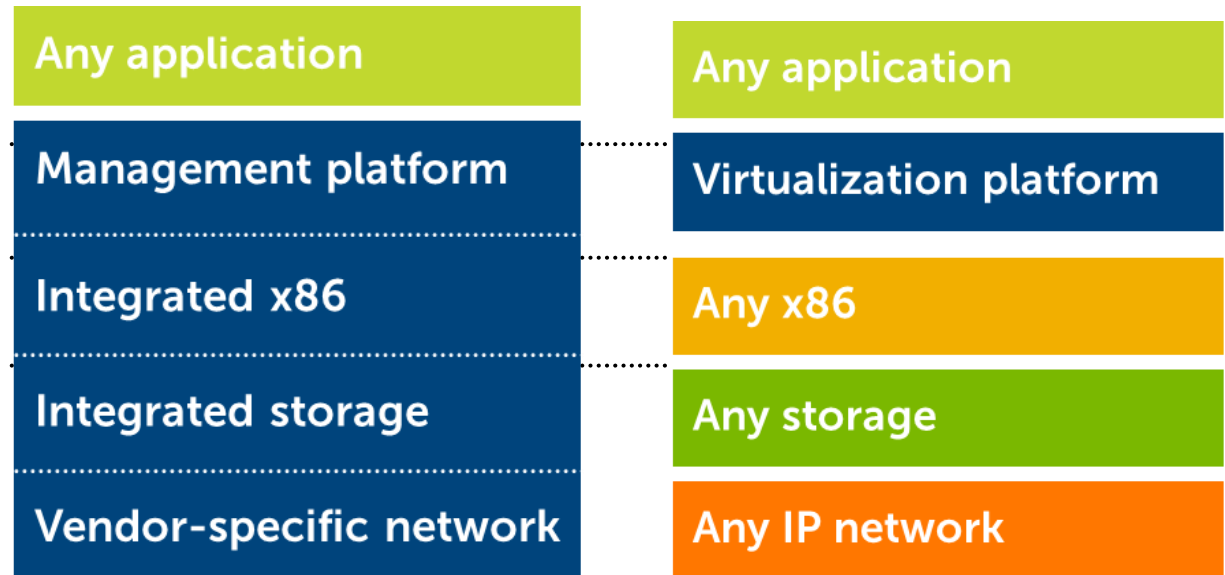
- provide the capability or functionality of providing an environment in which Virtualized network functions (VNF) can execute
- **NFVlaaS** provides compute capabilities comparable to an **IaaS cloud computing service** as a run time execution environment **as well as support the dynamic network connectivity services** that may be considered as comparable to **NaaS**



NFV vs SDN


- **NFV: re-definition of network equipment architecture**
- NFV was born to meet Service Provider (SP) needs:
 - Lower CAPEX by reducing/eliminating proprietary hardware
 - Consolidate multiple network functions onto industry standard platforms
- **SDN: re-definition of network architecture**
- SDN comes from the IT world:
 - Separate the data and control layers, while centralizing the control
 - Deliver the ability to program network behavior using well-defined interfaces

SDDC delivers needed agility and efficiency



Benefit	Hardware-defined (HDDC)	Software-defined (SDDC)
Innovation	Slow Long hardware/ASIC cycles	Fast Rapid software innovation
Flexibility	No Lock-in	Yes Choice of infrastructure
Ease of insertion/ deployment	Low Requires forklift upgrade	High Non-disruptive

Enabling the Future Ready Enterprise



Network Slicing

- One of the enabling technologies for SDN in Data Center
- Traditional Network Device Control

Control
Plane

- Computes forwarding rules
 - “128.8.128/16 --> port 6”
- Pushes rules down to data plane



Rules

Control/Data
Protocol

Exceptions

Data
Plane

- Enforces forwarding rules
- Exceptions pushed back to control plane



How to facilitate network slicing

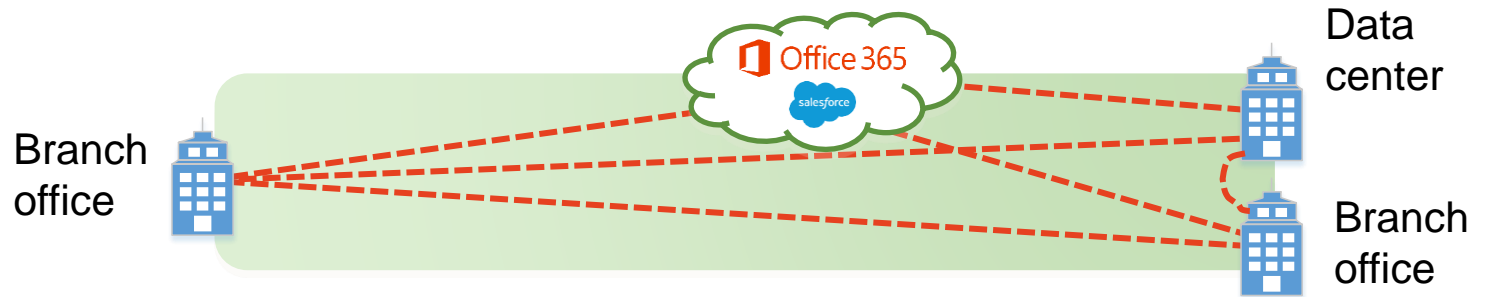
- Data plane unmodified
 - No performance penalty
- Control Policy: Specify resource limits for each slice
 - Link bandwidth
 - Maximum number of forwarding rules
 - Topology
 - Fraction of switch/router CPU

SDN for Interdomain Routing - SDX


- Forwarding on **multiple header fields**
(not just destination IP prefixes)
- Ability to **control entire networks** with a single software program (not just immediate neighbors)
- **Direct control** over data-plane forwarding (not indirect control via control-plane arcana)

SD-WAN

Overlay



Infrastructure (“Underlay”)



B4: Google's Software Defined WAN

- Google's private WAN connecting its data centers
 - Elastic bandwidth demands
 - Can tolerate periodic failures with temporary BW reduction
 - Small number of sites
 - Allows special optimization
 - Complete control of end application
 - Application priorities and control bursts
 - Cost Sensitivity
 - Unsustainable cost projection with traditional approach (2-3x cost of a fully utilized WAN).

Good Luck!

Now this is not the end.

It is not even the beginning of the end.

But it is, **perhaps**, the end of the beginning.

-Winston Churchill

USS feedback

- I need at least 67% response!
- <http://sydney.edu.au/itl/surveys/complete/>
- Your participation is appreciated!
- Lottery – iPad, gift card, etc.